

Network Applications: P2P, Streaming, and CDN

September 24, 2024

Min Suk Kang

Associate Professor

School of Computing/Graduate School of Information Security



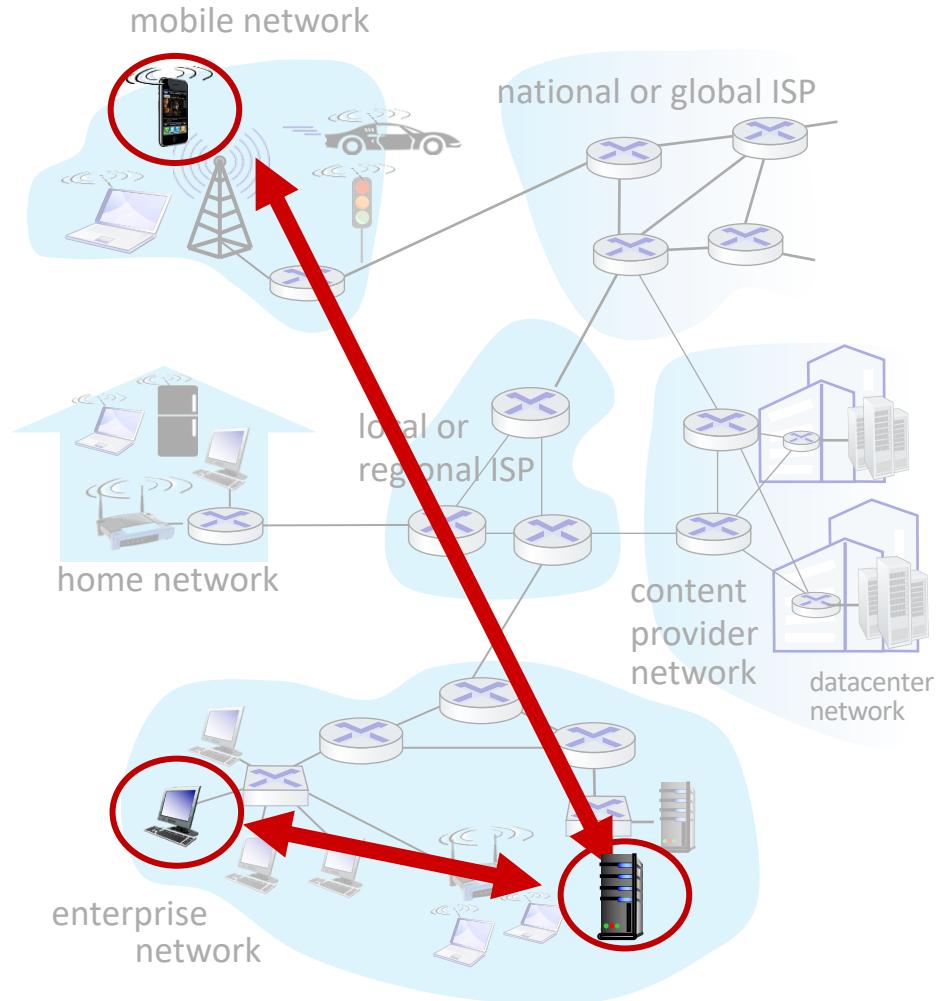
Client-server paradigm

server:

- always-on host
- permanent IP address
- often in data centers, for scaling

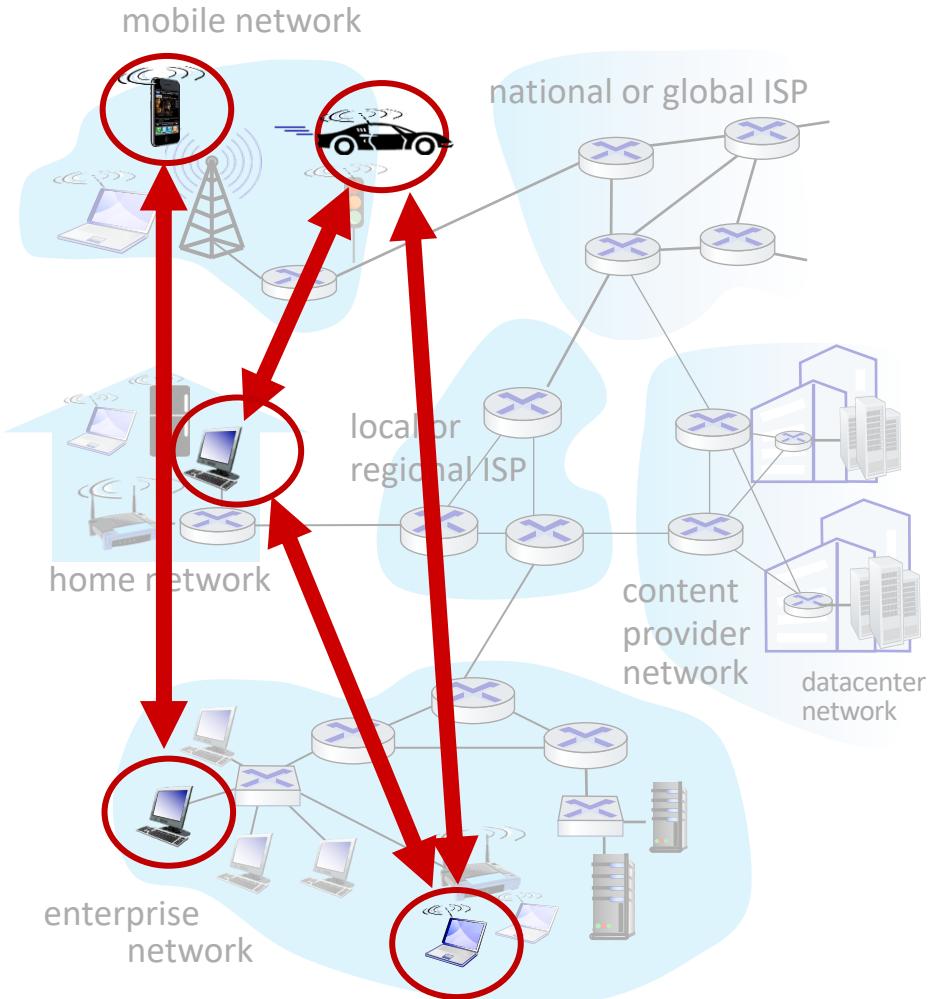
clients:

- contact, communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do *not* communicate directly with each other
- examples: HTTP, IMAP, FTP



Peer-peer architecture

- no always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
 - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers may be intermittently connected and change IP addresses
 - complex management
- example: P2P file sharing



Real-world P2P applications?

- BitTorrent
- Decentralized webs
- Blockchain
- ...

Defining blockchain *without* any jargons

- “Agreeing on something among a group of parties who care about it”

Defining blockchain *with* some jargons

- “Agreeing on something among a group of parties who care about it”

consensus algorithms

- safety
- liveness
- scalability
- finality

distributed ledger

- immutable
- total order
- fair ordering

decentralized peers

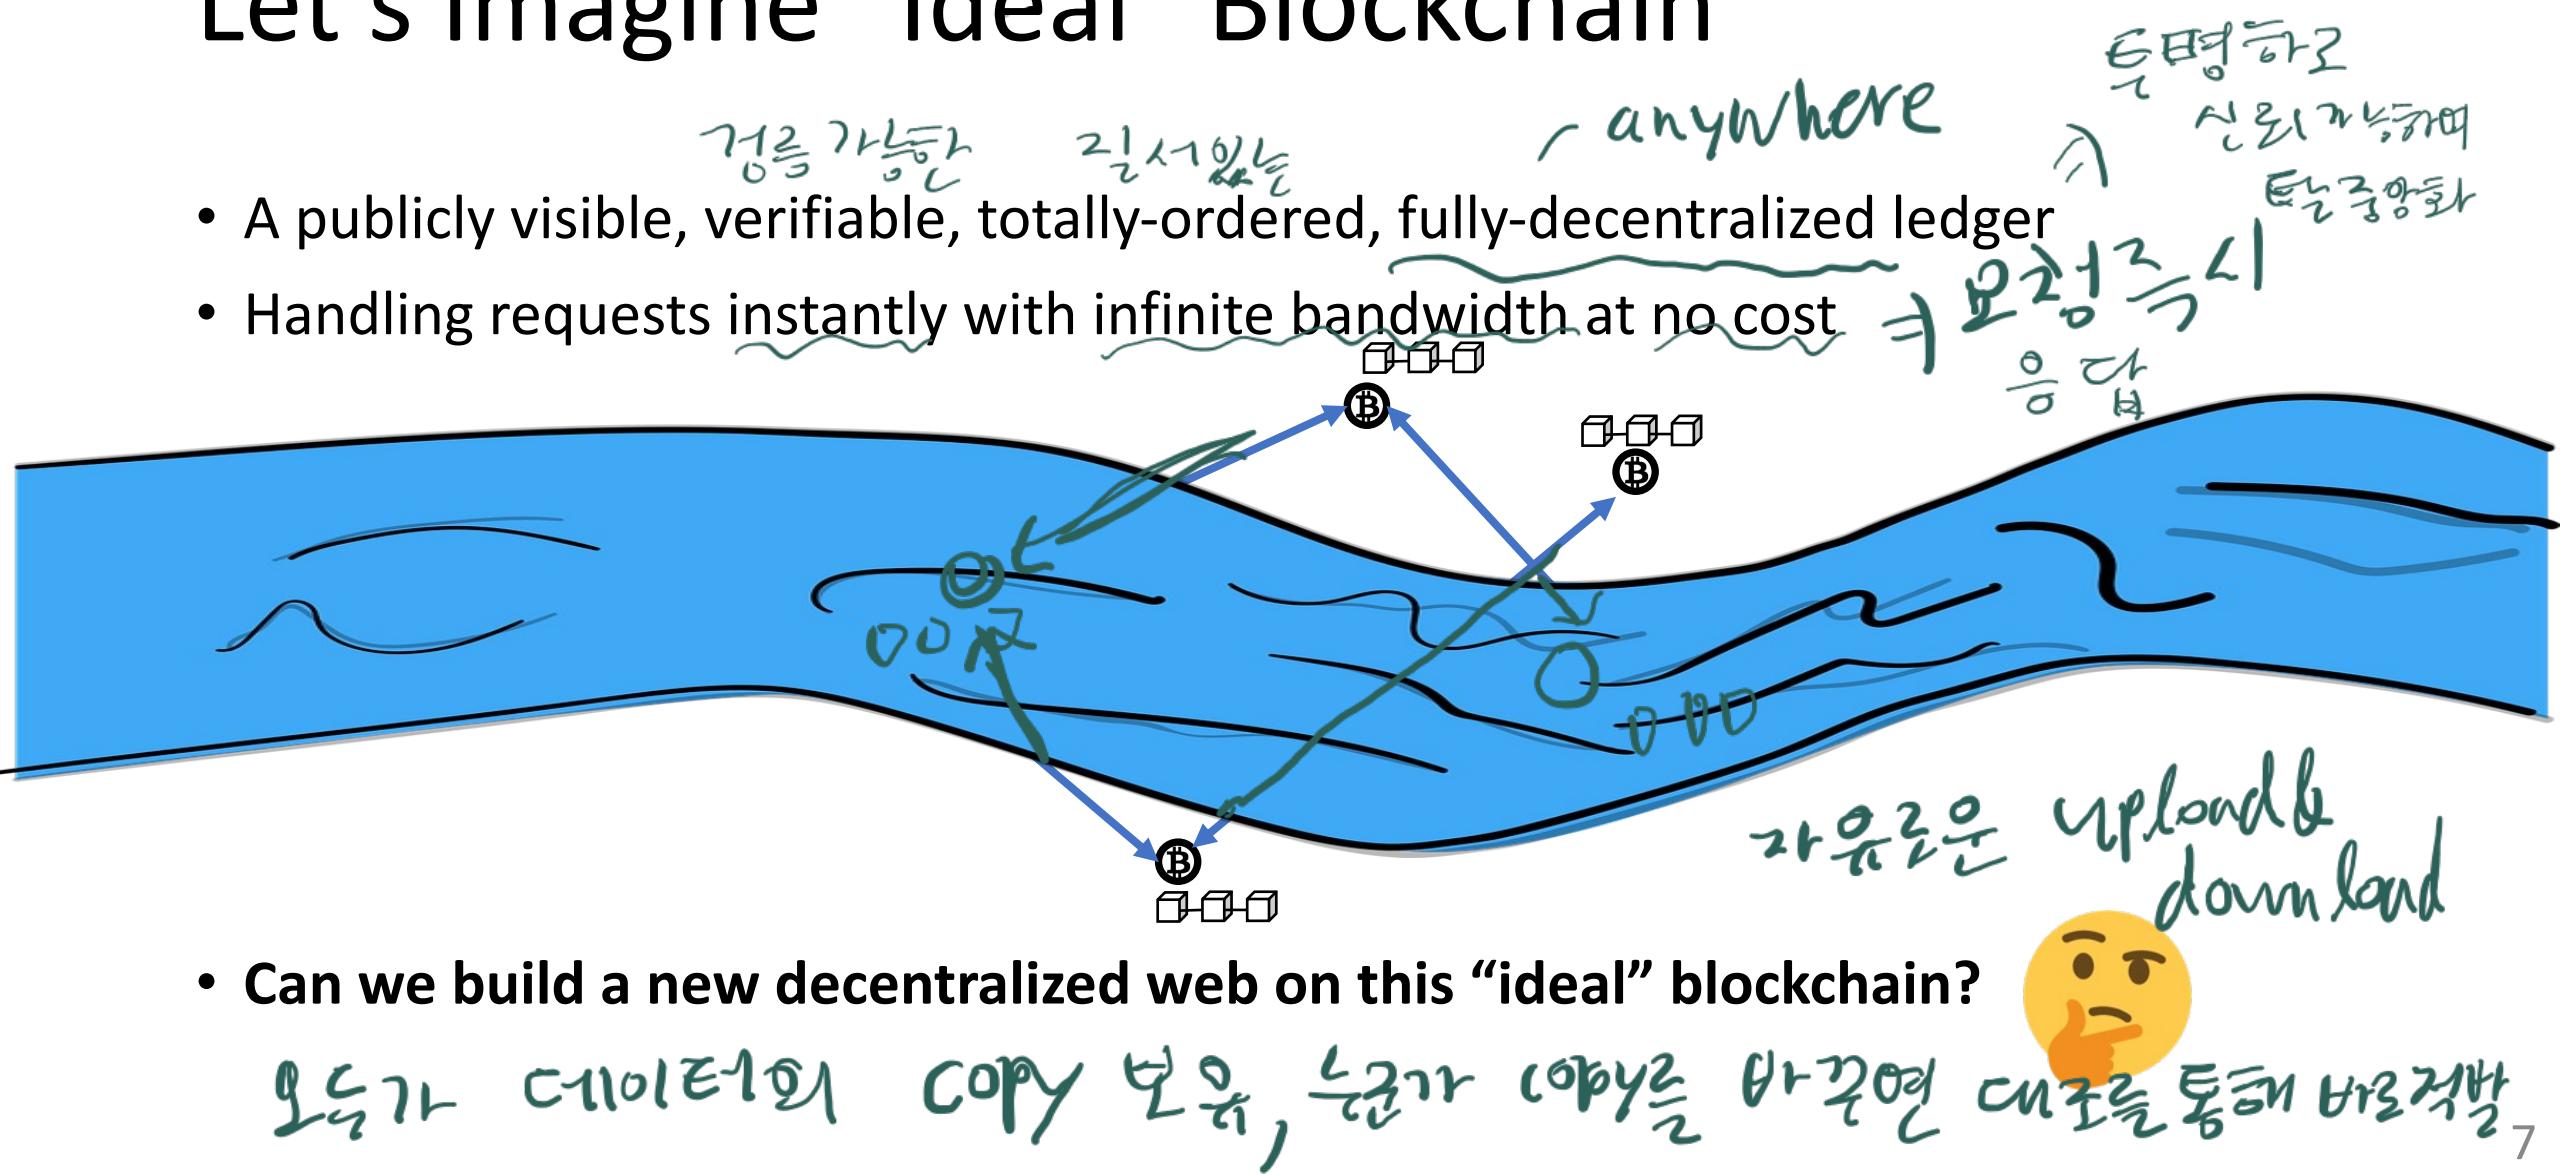
- permissionless
- permissioned
- Byzantine peers

Sybil protection

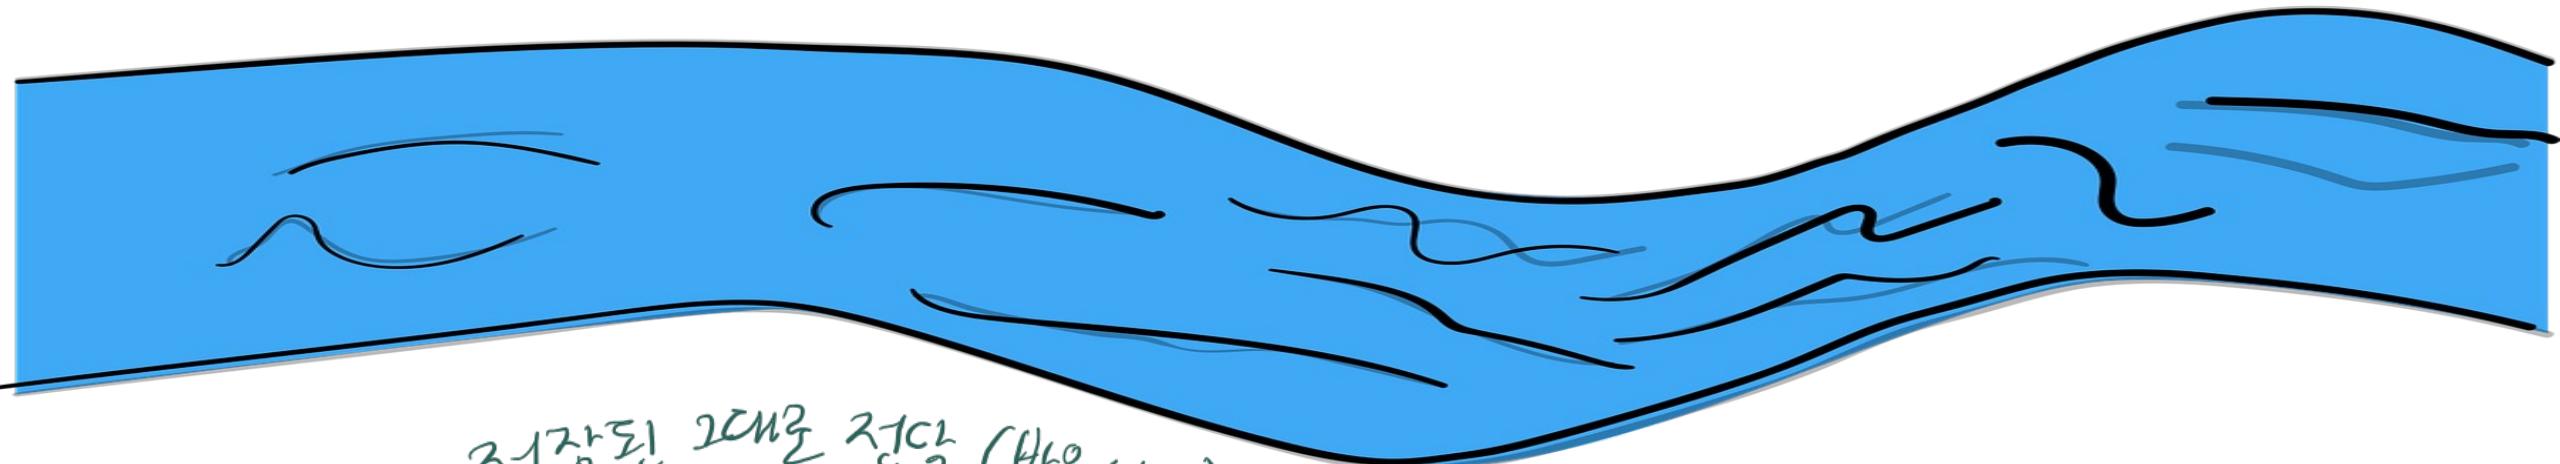
- proof of work
- proof of stake

Let's Imagine “Ideal” Blockchain

- A publicly visible, verifiable, totally-ordered, fully-decentralized ledger
- Handling requests instantly with infinite bandwidth at no cost



Building a New Web



기장된 구조로 전달 (비용성X)

- Static web
- Microblogging 단편적 정보 실시간 전달
- Facebook
 - Social graph, chats
- ...

비유적
web service 가능
peer 간 통신을 통해

- How would this new web change the ecosystem?

Back to Reality: Existing Blockchains

- Reality check:

- 7 transactions per second, 1-hour for finalizing a transaction, burning too much fossil fuel?

한당 개화 트랜잭션

트랜잭션 완료를 위해 1시간 → 시간 ↑

정보 위치 및 변조여부 확인



- Getting better (very quickly)

- >100k transactions per second
 - <1sec finalization
 - near zero cost operation

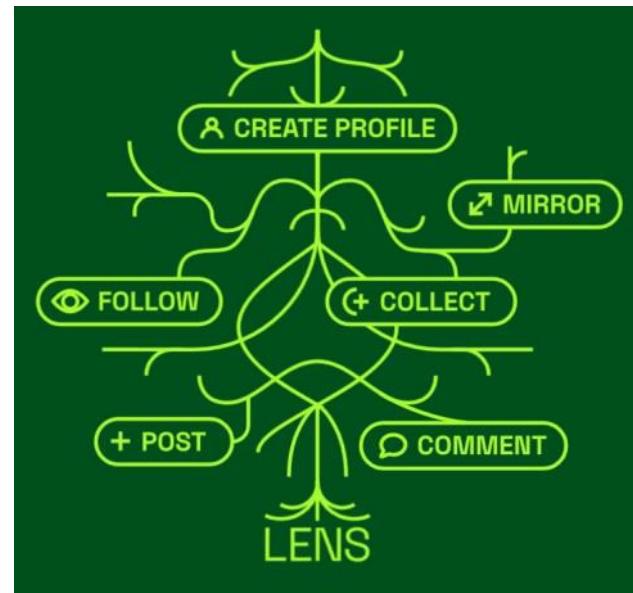
- Technologies behind them

- new distributed system techniques (e.g., DAG, proof of stake)
 - new cryptography (e.g., zero-knowledge proof)
 - ... 암호학

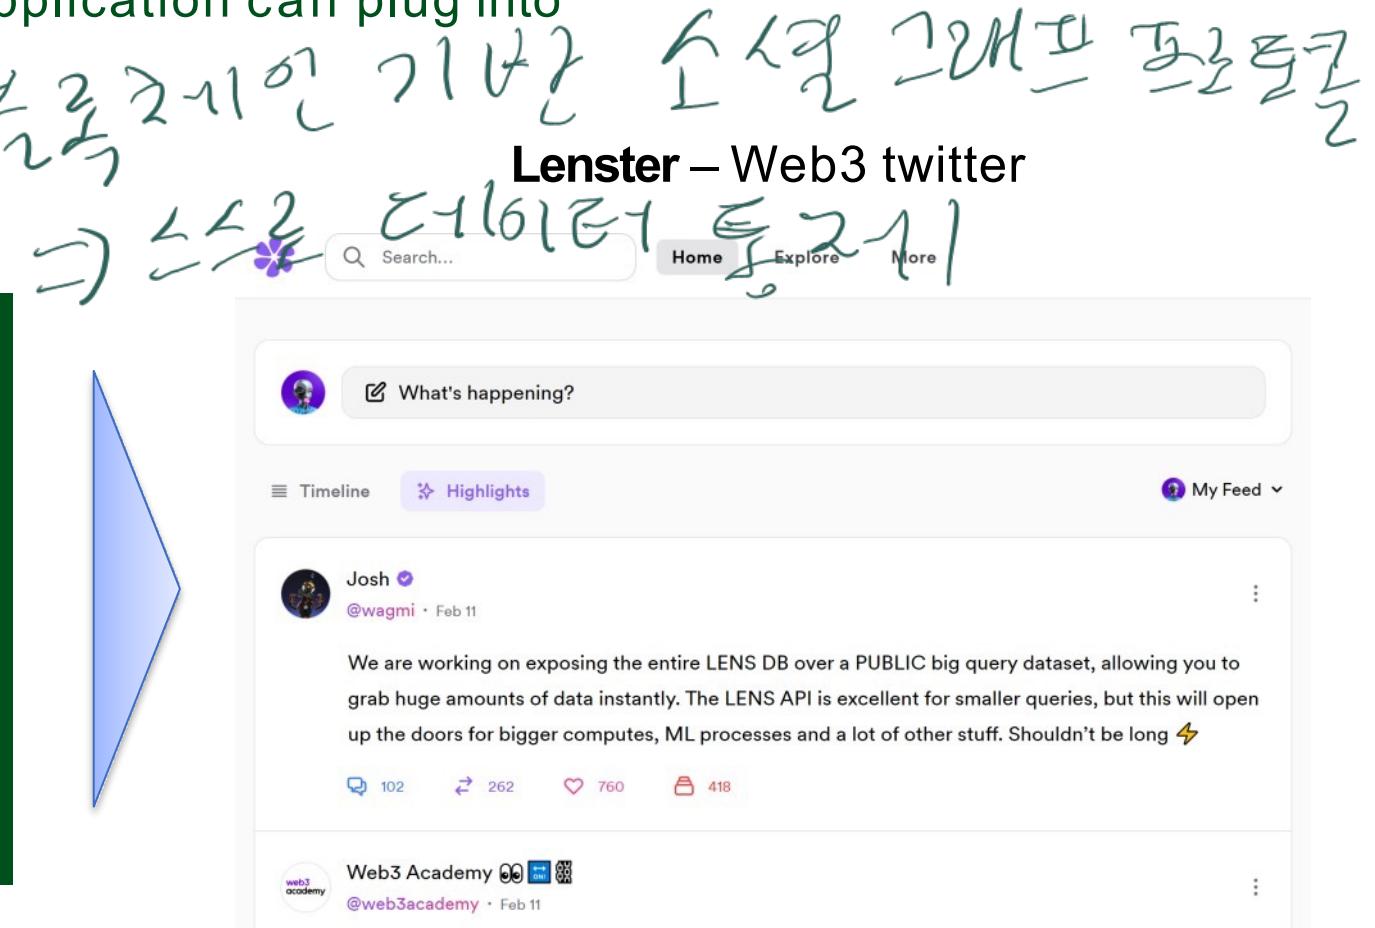
Use Case: Lens Protocol

A user-owned, open social graph
that any application can plug into

social graph
data elements



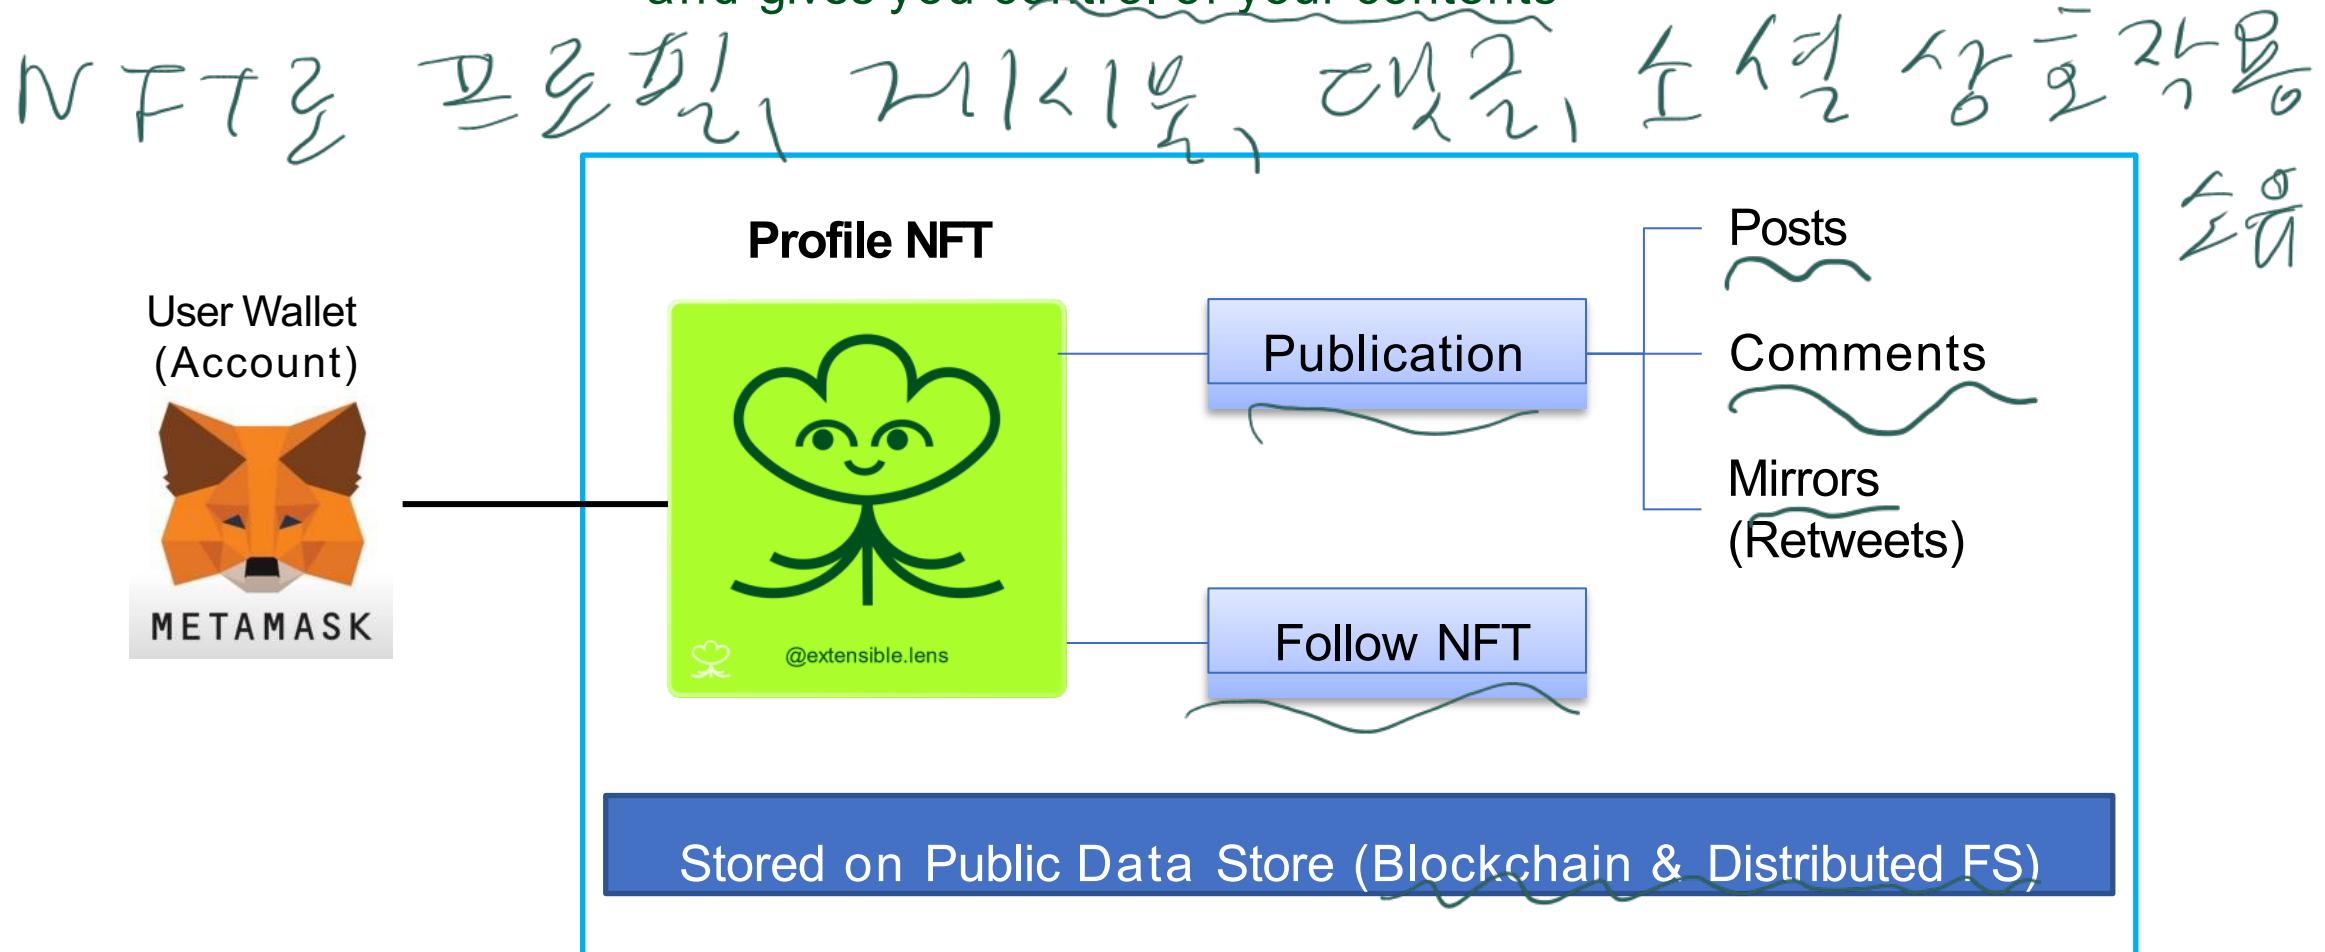
<https://www.lens.xyz/>



<https://lenster.xyz/>

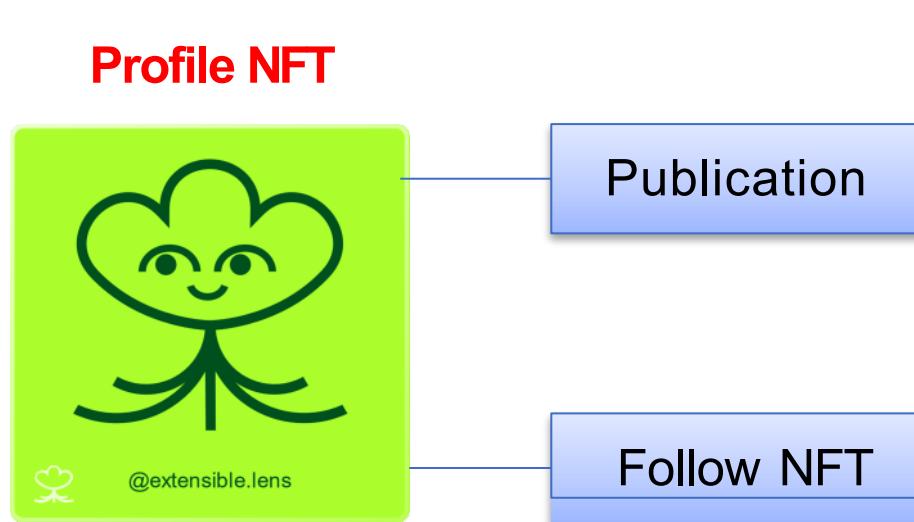
1. Data Ownership to Users

Profile NFT is a key element to owns your data and gives you control of your contents



2. Data as Digital Assets

Follow NFT, Collect, Profile NFT as a digital asset



Buy & Sell
a Profile NFT
with all content

근친친구 수영장

Collect
Paid content

Posts
Comments
Mirrors
(Retweets)

Profile NFT market on OpenSea



<https://opensea.io/collection/lens-protocol-profiles>

3. Communities as data governance actors

Governance for Lens Protocol

Community Multisig

by the Lens community

- Setting up Governance and Emergency Admin Addresses
 - Setting Treasury Addresses and Fees
 - Whitelisting Assets
 - Moving the Lens Protocol System into a Publishing Paused or fully Paused state
 - Whitelisting addresses to create profiles
 - Whitelisting Follow, Collect, and Reference Modules
 - Upgrading the Lens Protocol Hub Contract
- Lens Protocol 허가 및 개선
Lens Protocol 허가 및 개선

Governance for social communities

내장

Built-in governance of Follow NFTs

to create a social DAO

(Decentralized Autonomous Organization)

단점화된

자율적

조직

bitcoin.lens

DAO



Follow
NFTs

extensible.lens
DAO



Follow
NFTs



Follow
NFTs

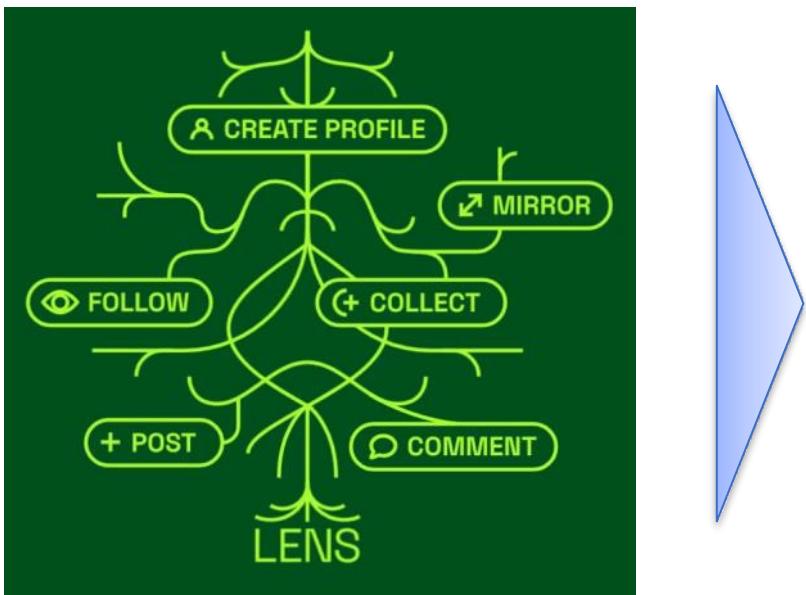


4. Protocols as data utilization tools

Lens Protocol makes user's social graph data a **social graph protocol**

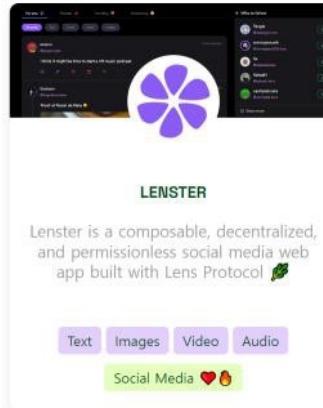
Applications utilize Lens Protocol to access user's social graph

↳ protocol을
백엔드로
app 사용하기



Lensverse
Hundreds of applications built on top of Lens Protocol

Web3 Twitter

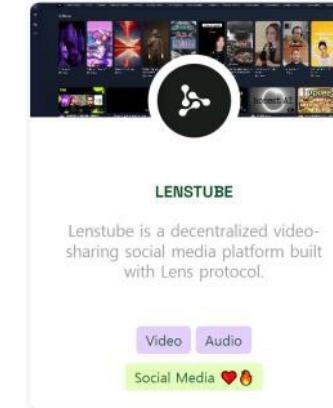


LENSTER

Lenster is a composable, decentralized, and permissionless social media web app built with Lens Protocol.

Text Images Video Audio
Social Media ❤️

Web3 Youtube

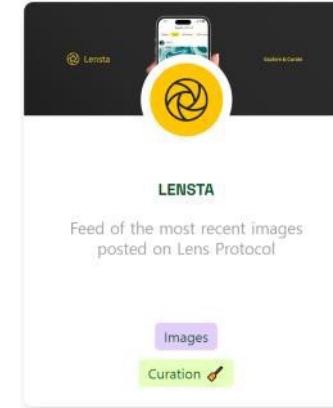


LENSTUBE

Lenstube is a decentralized video-sharing social media platform built with Lens protocol.

Video Audio
Social Media ❤️

Web3 Instagram



LENSTA

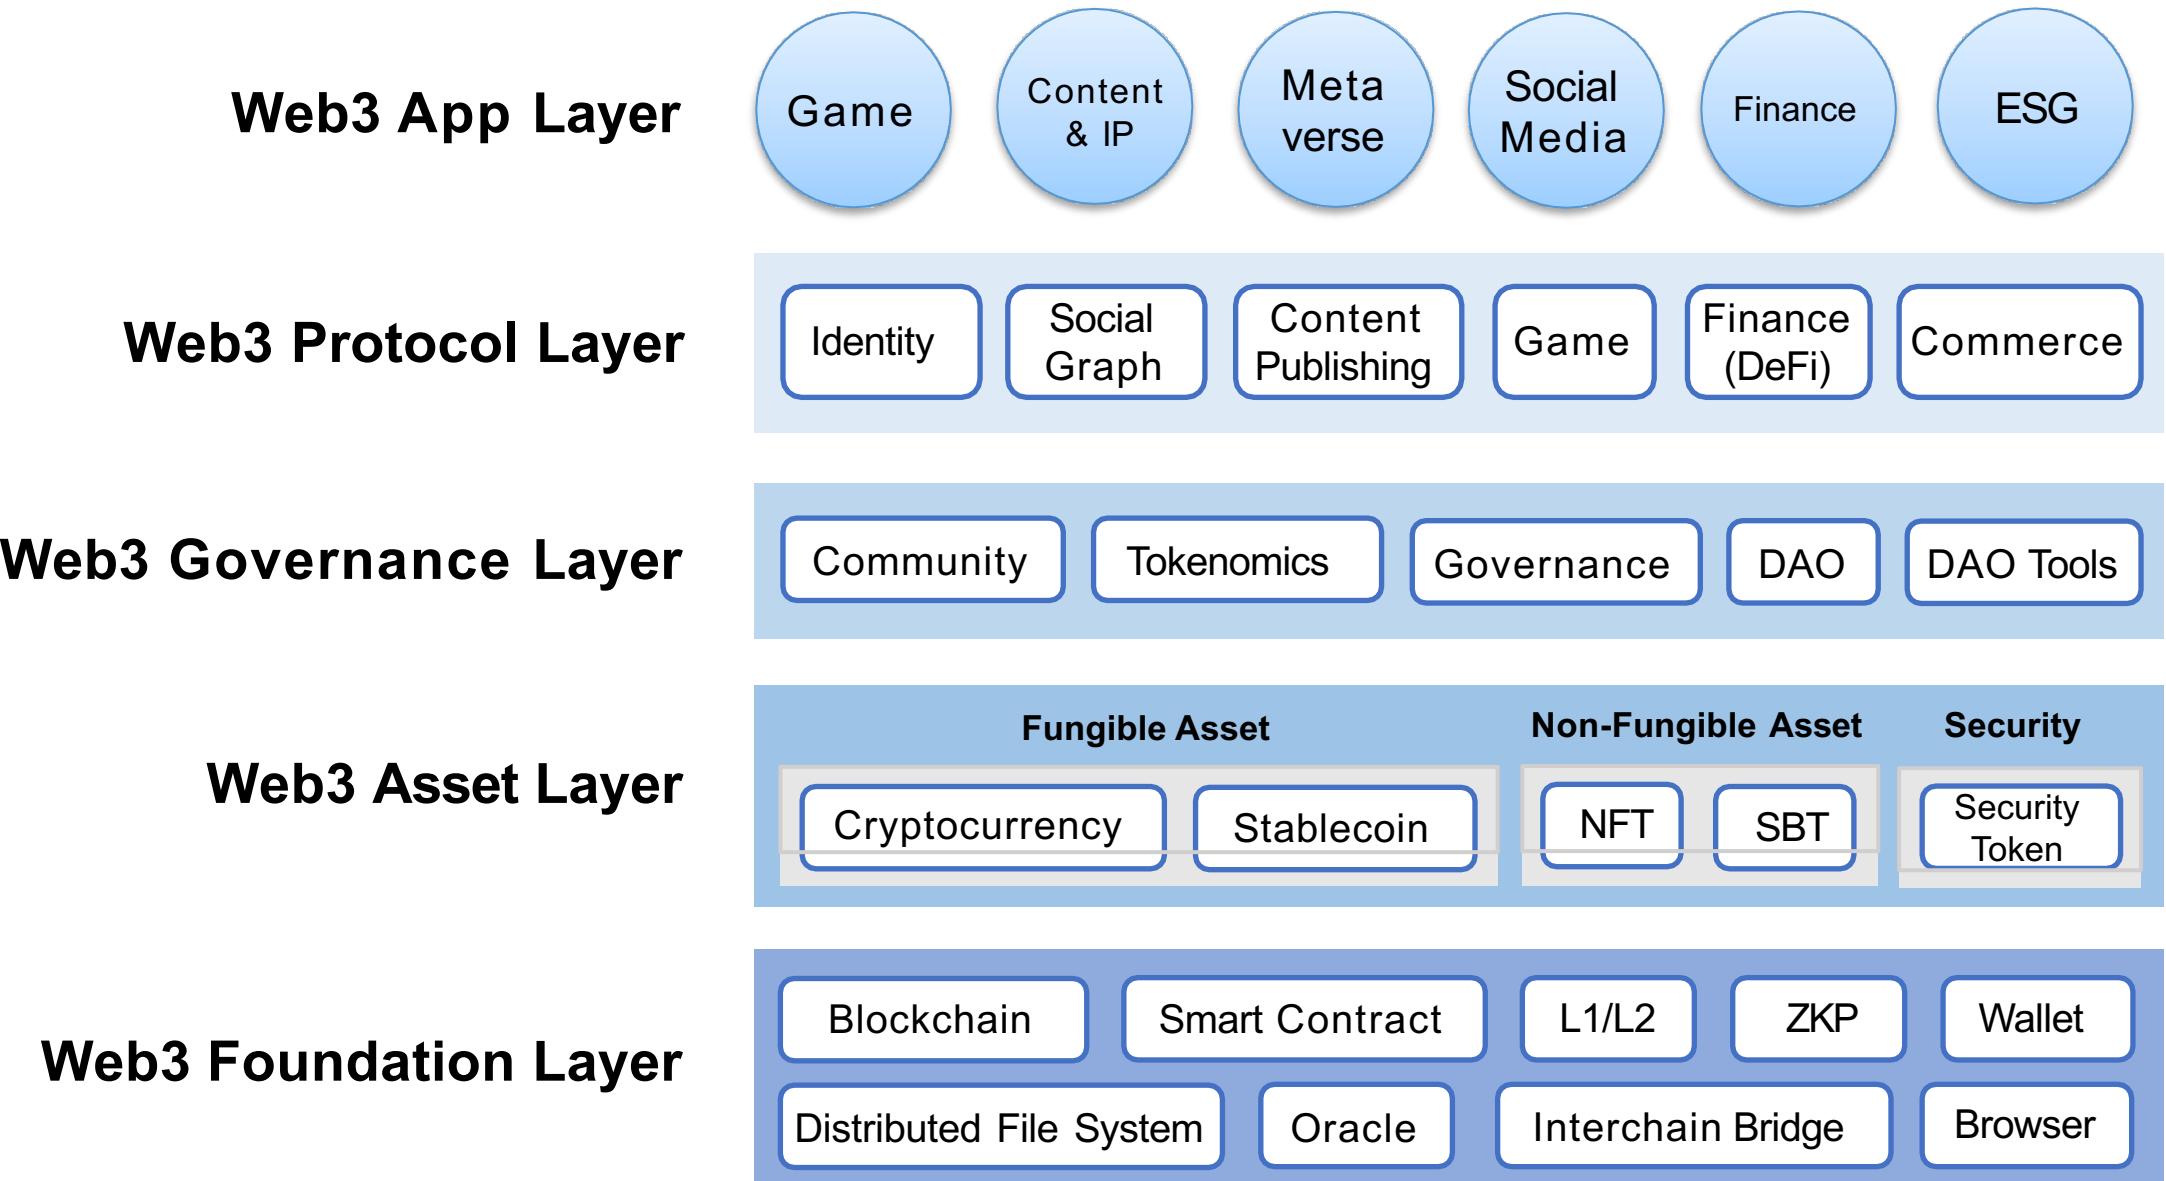
Feed of the most recent images posted on Lens Protocol

Images
Curation 📚

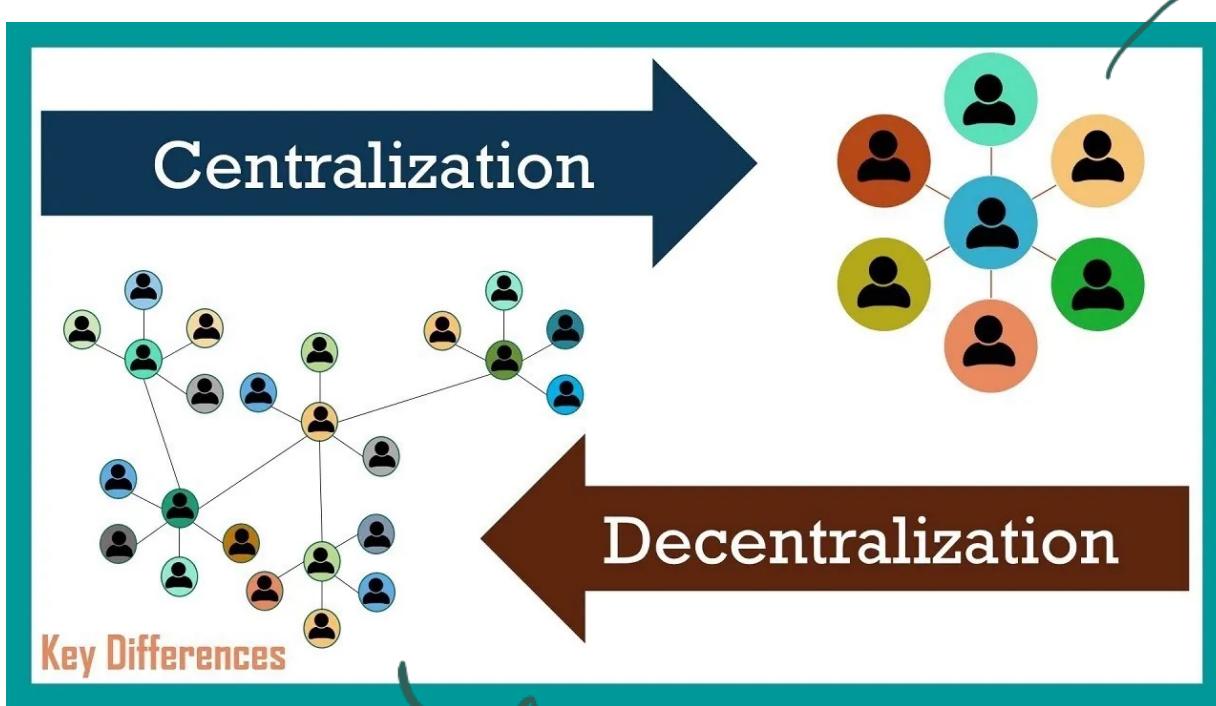
<https://www.lens.xyz/apps>



Web3 Stack (in the view of data)



Takeaways



router
protocol



속도↑,

중앙 통제↓

- Can you imagine the next Internet or Web?

SNS → 탈중앙화,
경매 or 조작X
(데이터가 빠르게 송수)

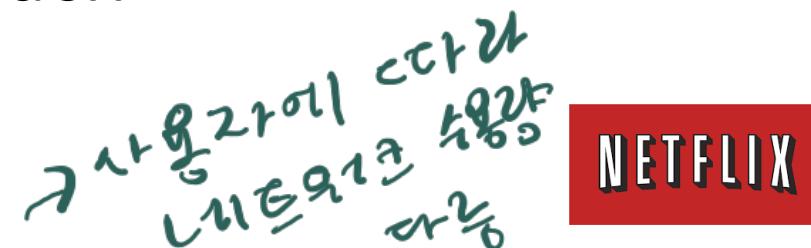
Application layer: overview

- Principles of network applications
- Web and HTTP
- E-mail, SMTP, IMAP
- The Domain Name System DNS
- P2P applications
- video streaming and content distribution networks
- socket programming with UDP and TCP



Video Streaming and CDNs: context

- stream video traffic: major consumer of Internet bandwidth
 - Netflix, YouTube, Amazon Prime: 80% of residential ISP traffic (2020)
 - *challenge*: scale - how to reach ~1B users?
 - *challenge*: heterogeneity
 - different users have different capabilities (e.g., wired versus mobile; bandwidth rich versus bandwidth poor)
 - *solution*: distributed, application-level infrastructure



Multimedia: video

- video: sequence of images displayed at constant rate
 - e.g., 24 images/sec
- digital image: array of pixels
 - each pixel represented by bits
- coding: use redundancy *within* and *between* images to decrease # bits used to encode image
 - spatial (within image) 공간적
 - temporal (from one image to next) 시간적

일정 속도로
이미지의
연속 display



frame *i*

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (purple) and number of repeated values (N)



frame *i+1*

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i

Multimedia: video

- CBR: (constant bit rate): video encoding rate fixed

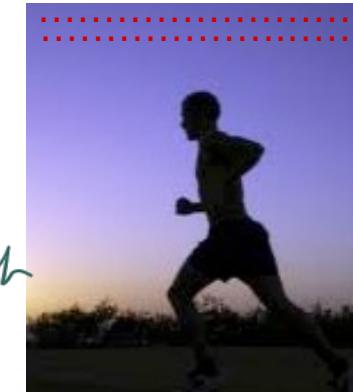
- VBR: (variable bit rate): video encoding rate changes as amount of spatial, temporal coding changes → *coding 비율이 변화하는 경우
encoding 속도도 변화*

- examples:

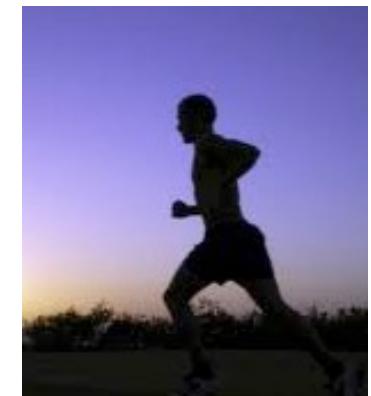
- MPEG 1 (CD-ROM) 1.5 Mbps ~ CBR
- MPEG2 (DVD) 3-6 Mbps ~ VBR
- MPEG4 (often used in Internet, 64Kbps – 12 Mbps)

VBR

spatial coding example: instead of sending N values of same color (all purple), send only two values: color value (*purple*) and number of repeated values (N)



frame i

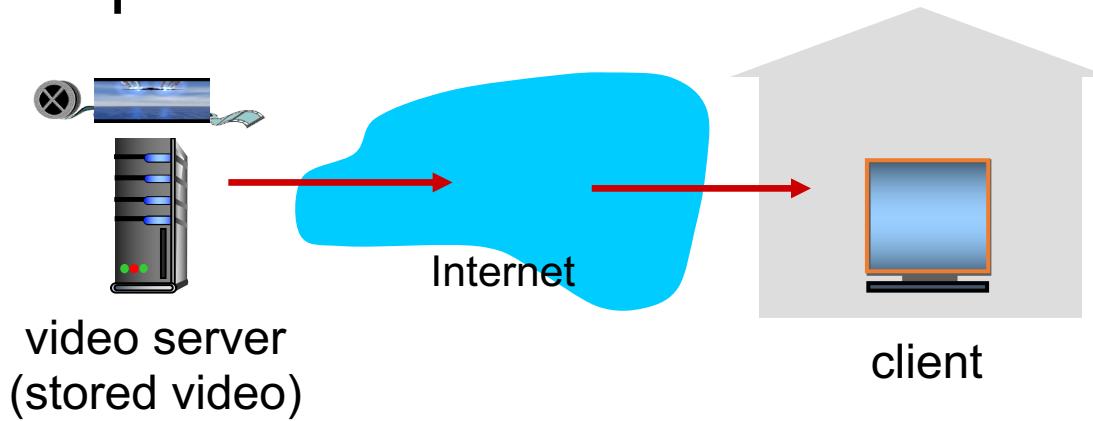


frame $i+1$

temporal coding example: instead of sending complete frame at $i+1$, send only differences from frame i

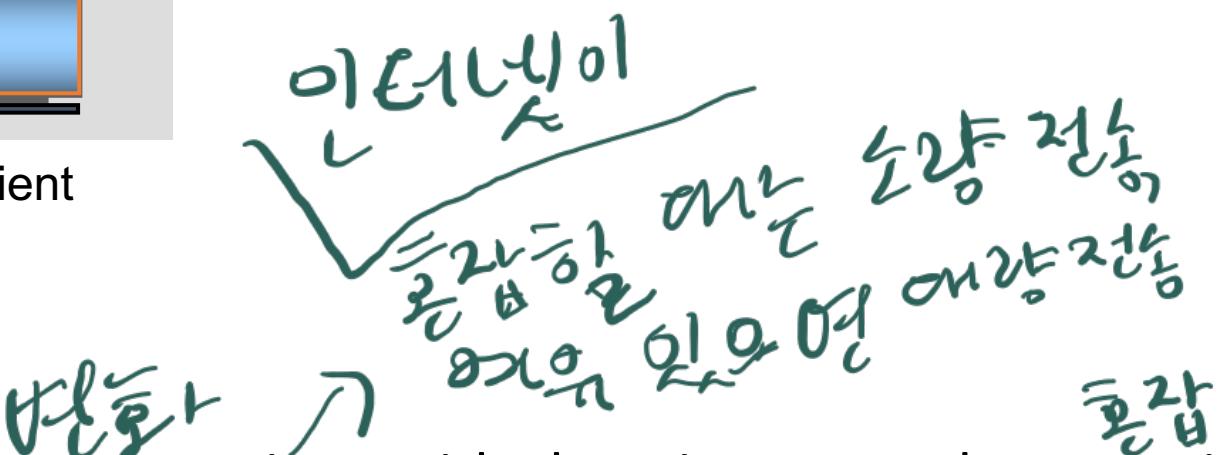
Streaming stored video

simple scenario:

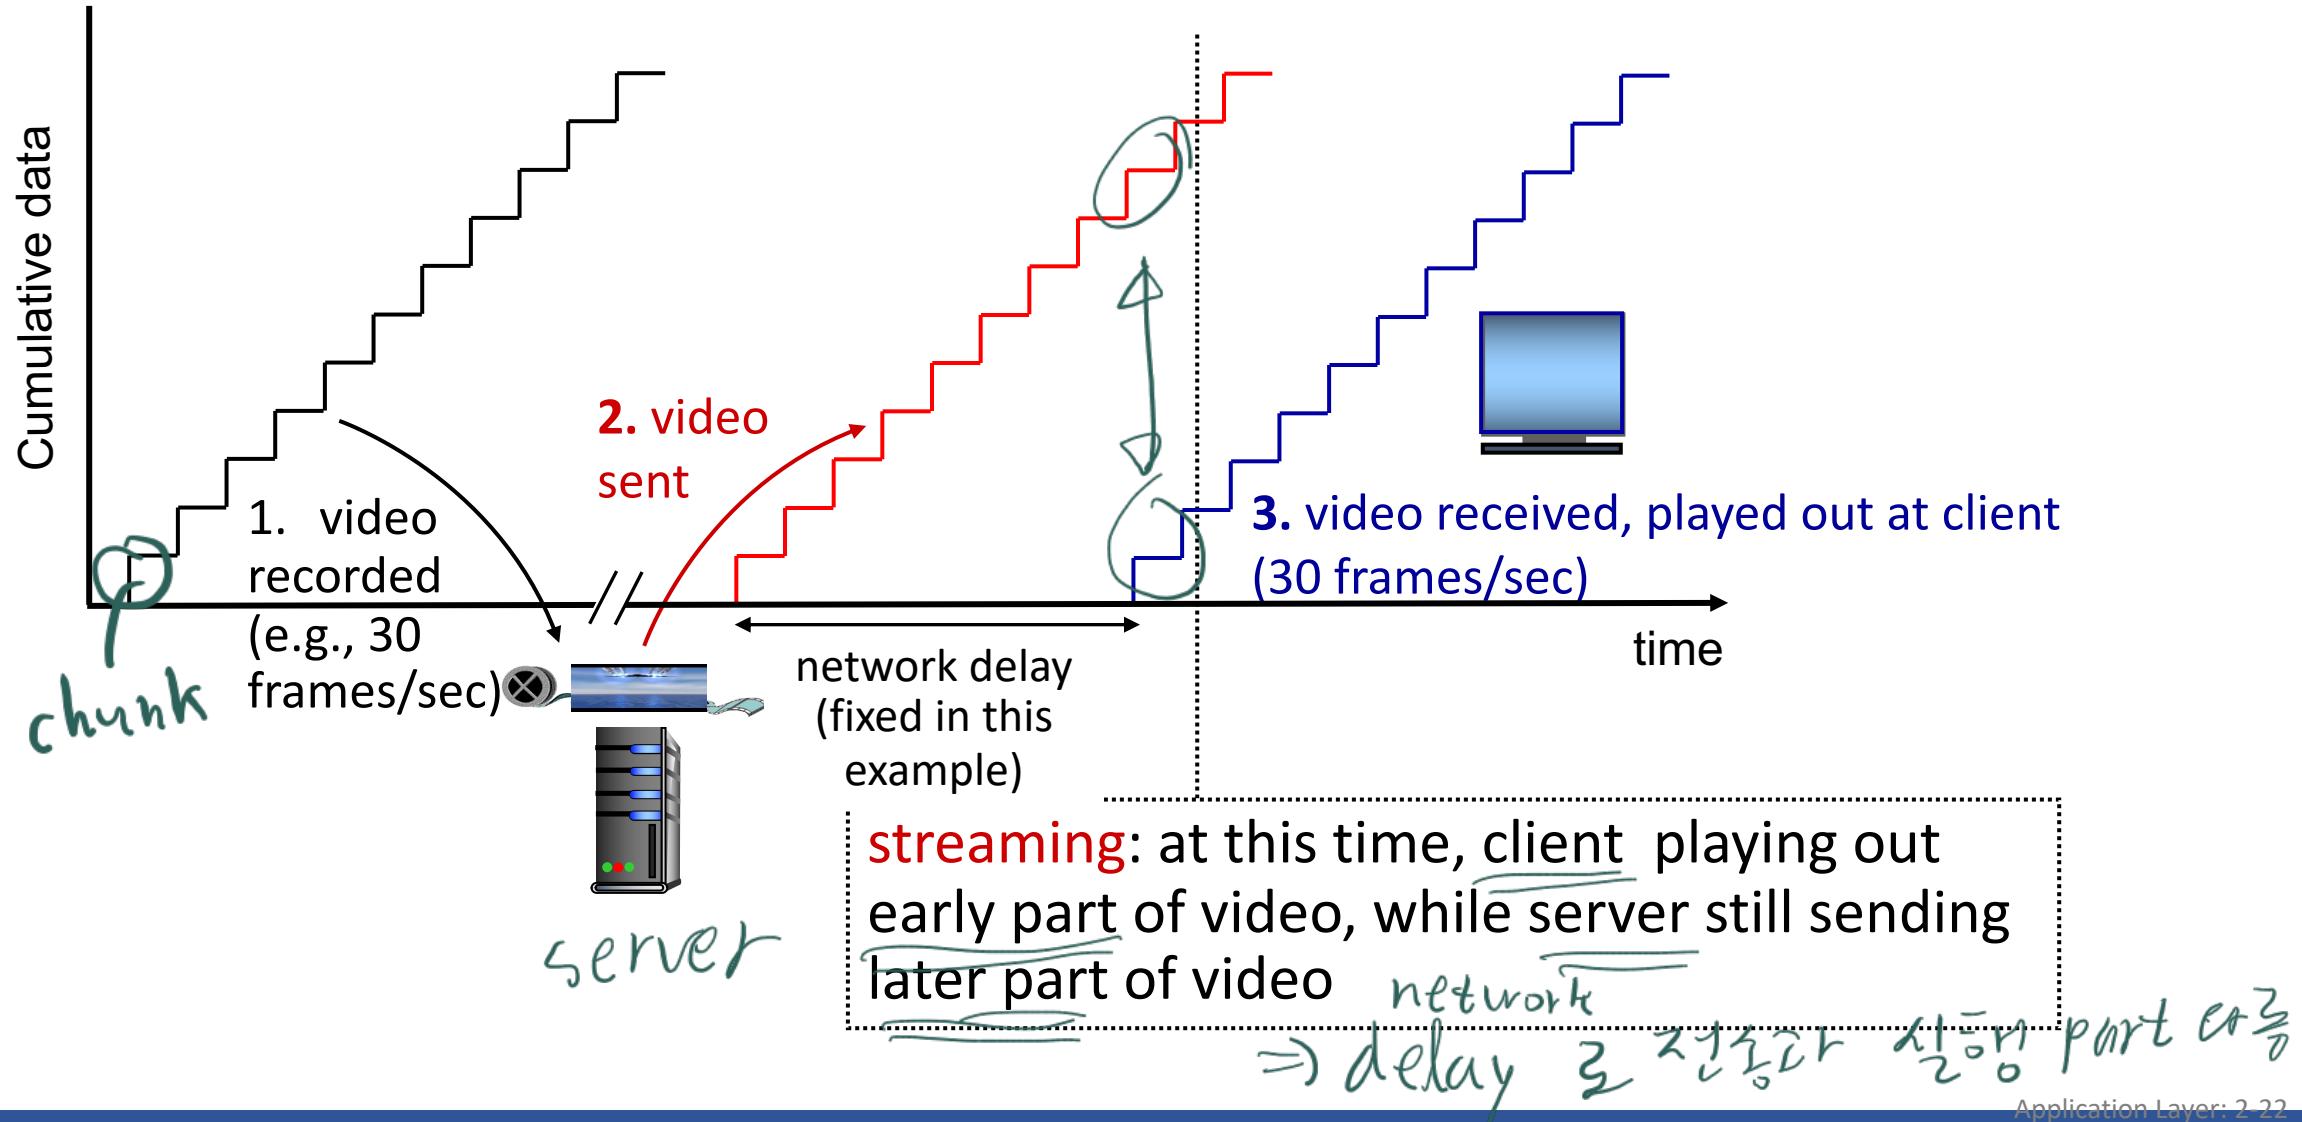


Main challenges:

- server-to-client bandwidth will *vary* over time, with changing network congestion levels (in house, access network, network core, video server)
 - packet loss, delay due to congestion will delay playout, or result in poor video quality



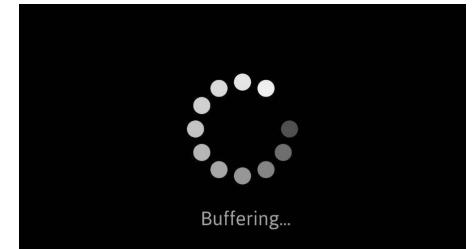
Streaming stored video



Streaming stored video: challenges

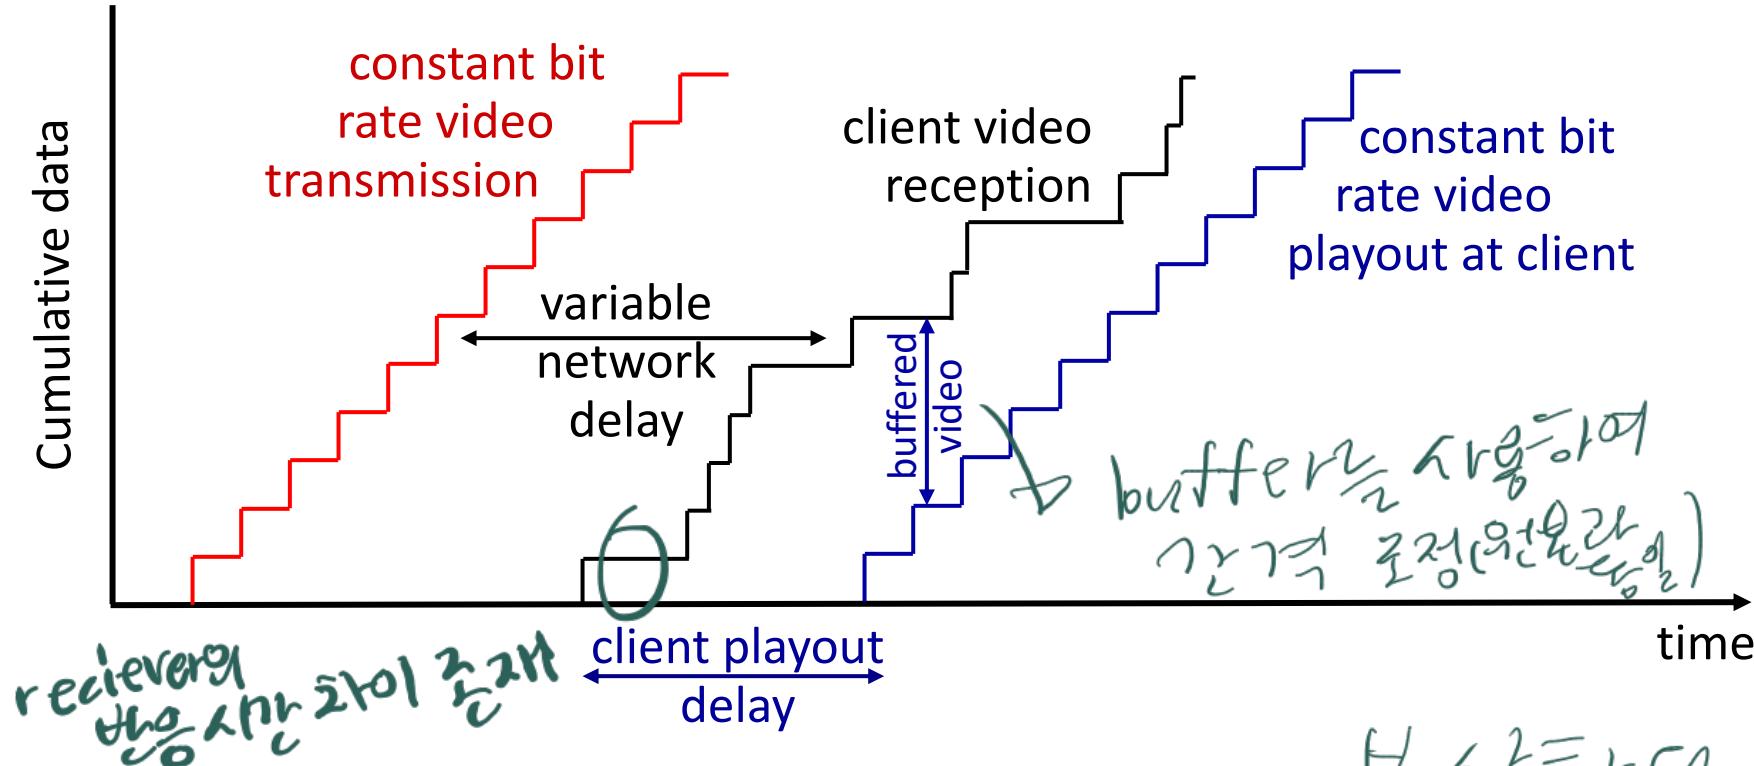
2/10

- **continuous playout constraint:** during client video playout, playout timing must match original timing *(client video playout = original)*
 - ... but **network delays are variable** (jitter), so will need **client-side buffer** to match continuous playout constraint
 \rightarrow delay 정도가 변하므로 이를 동일하게 위해 buffer 사용



- other challenges:
 - client interactivity: pause, fast-forward, rewind, jump through video \rightarrow 동작의 상호작용
 - video packets may be lost, retransmitted
 \rightsquigarrow 재전송

Streaming stored video: playout buffering



- **client-side buffering and playout delay:** compensate for network-added delay, delay jitter → ~~타이밍 불안정성~~

Streaming multimedia: DASH

Dynamic, Adaptive
Streaming over HTTP

server:

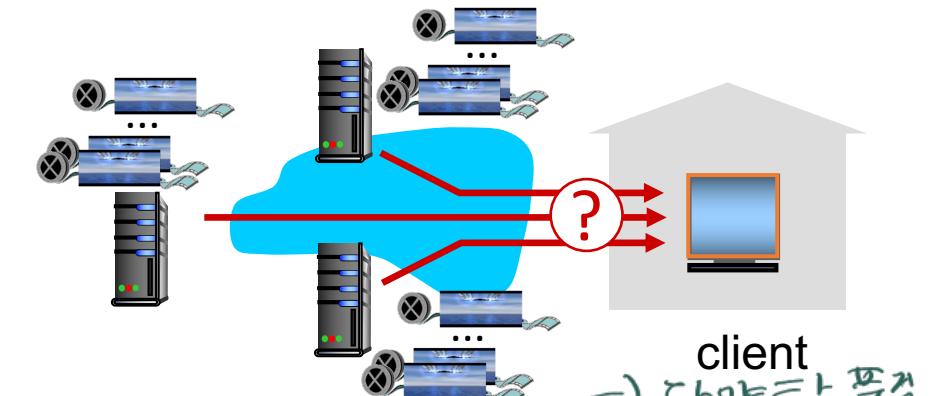
- divides video file into multiple chunks
- each chunk encoded at multiple different rates
- different rate encodings stored in different files
- files replicated in various CDN nodes
- manifest file*: provides URLs for different chunks

목록

file을 다수의 chunk로 만들고

client:

- periodically estimates server-to-client bandwidth
- consulting manifest, requests one chunk at a time
 - chooses maximum coding rate sustainable given current bandwidth
 - can choose different coding rates at different points in time (depending on available bandwidth at time), and from different servers



client

다양한 풍질
수준 제공

encoding

을 위한 고정

등급에 따라 chunk 요청

풍질 지정

고정

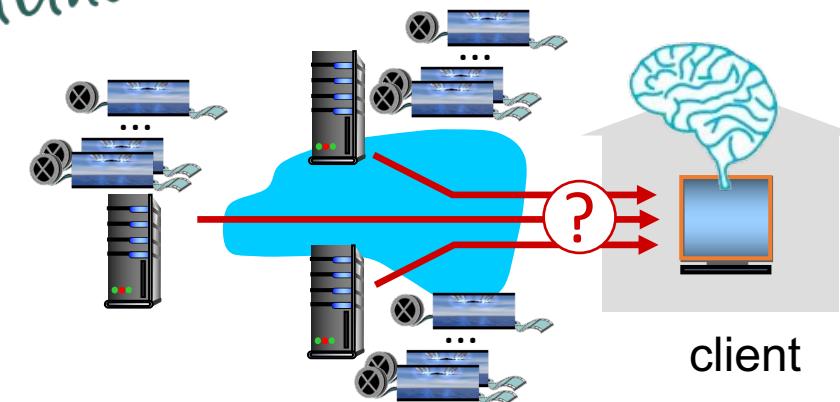
지정

Streaming multimedia: DASH

- “intelligence” at client: client determines

- when to request chunk (so that buffer starvation, or overflow does not occur)
- what encoding rate to request (higher quality when more bandwidth available) → 자동으로 풍질 수준 결정 (해당 타이밍의 네트워크 조건에 따라)
- where to request chunk (can request from URL server that is “close” to client or has high available bandwidth) → 가깝거나 (속도 빠를 가능성↑), 대역폭↑ (속도↑)인 URL

온투보의 경우, 3~4초마다 chunk 요청



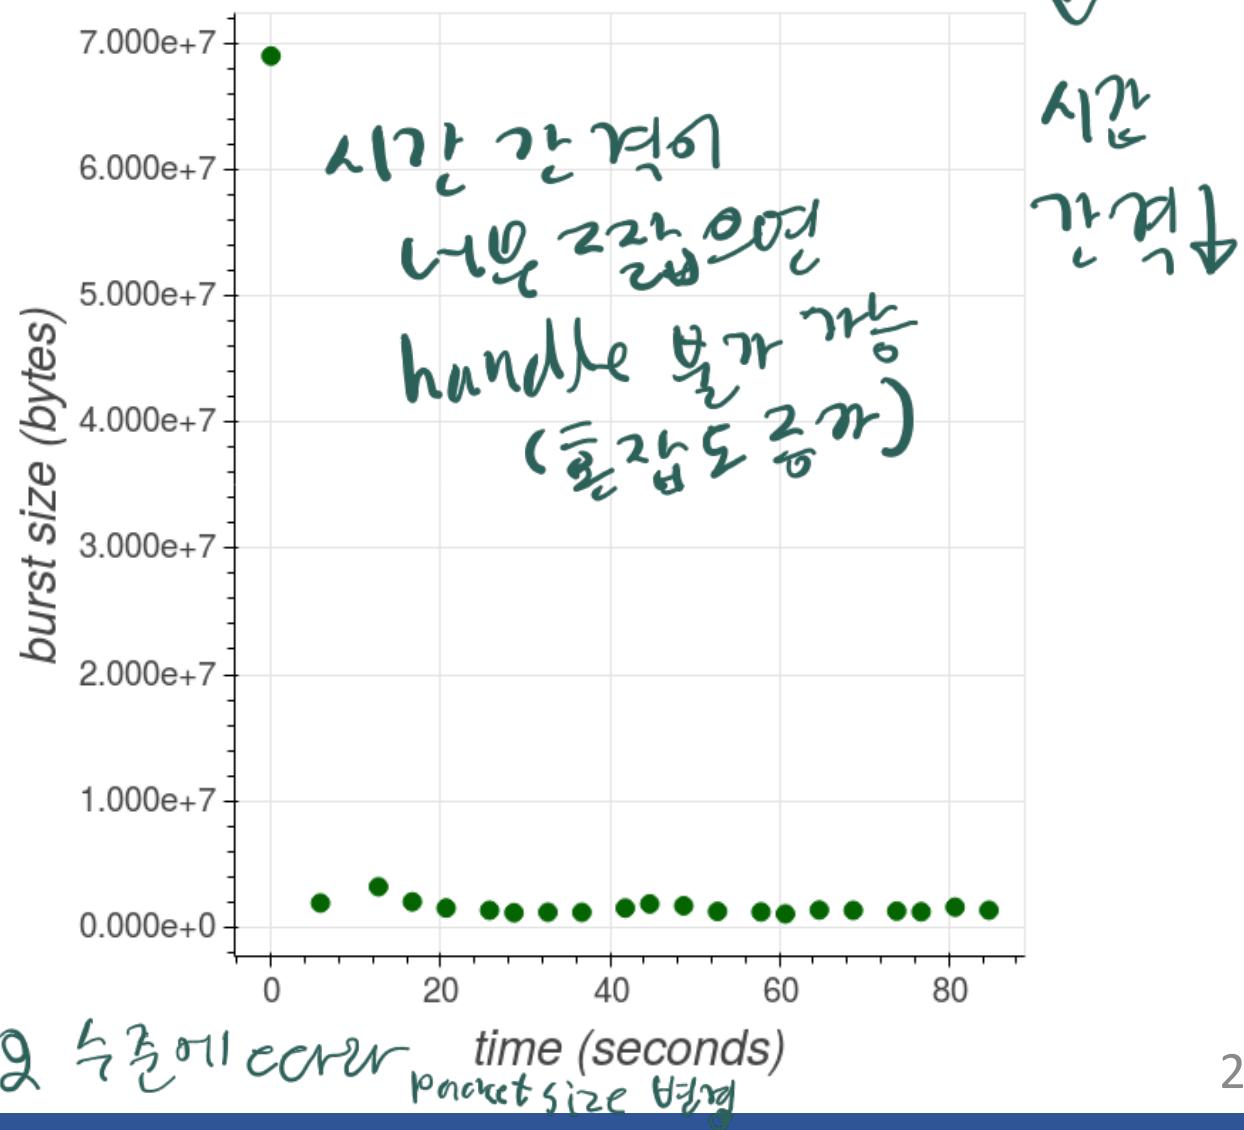
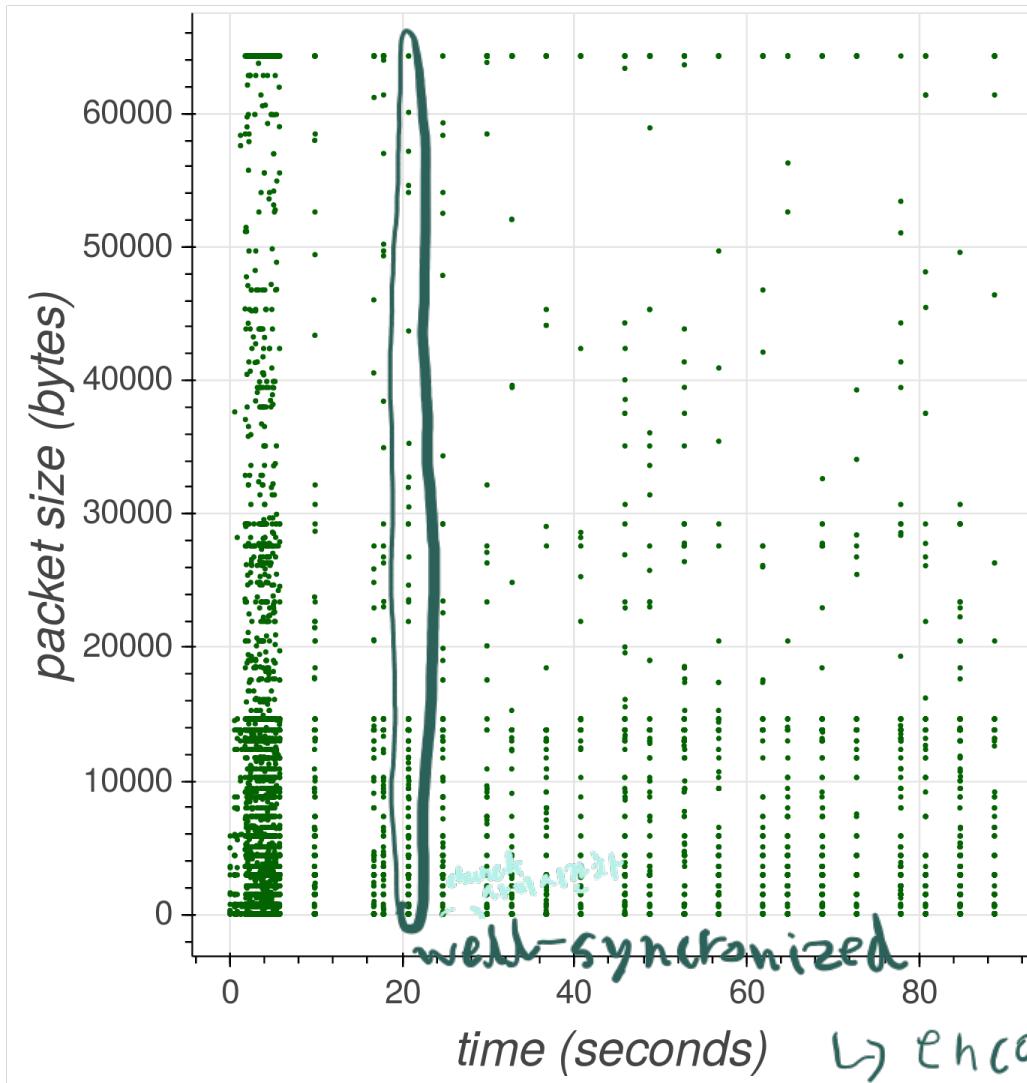
Streaming video = encoding + DASH + playout buffering

내용을 일련의 작은 크기의 HTTP 기반 파일 세그먼트들로 분리시킴으로써 동작하며, 각 세그먼트는 영화나 스포츠 이벤트 생방송 등 잠재적으로 수 시간에 걸친 내용물의 재생 시간의 짧은 간격(interval)을 포함하고 있다. 이 콘텐츠는 다양한 비트레이트로 이용이 가능하다.

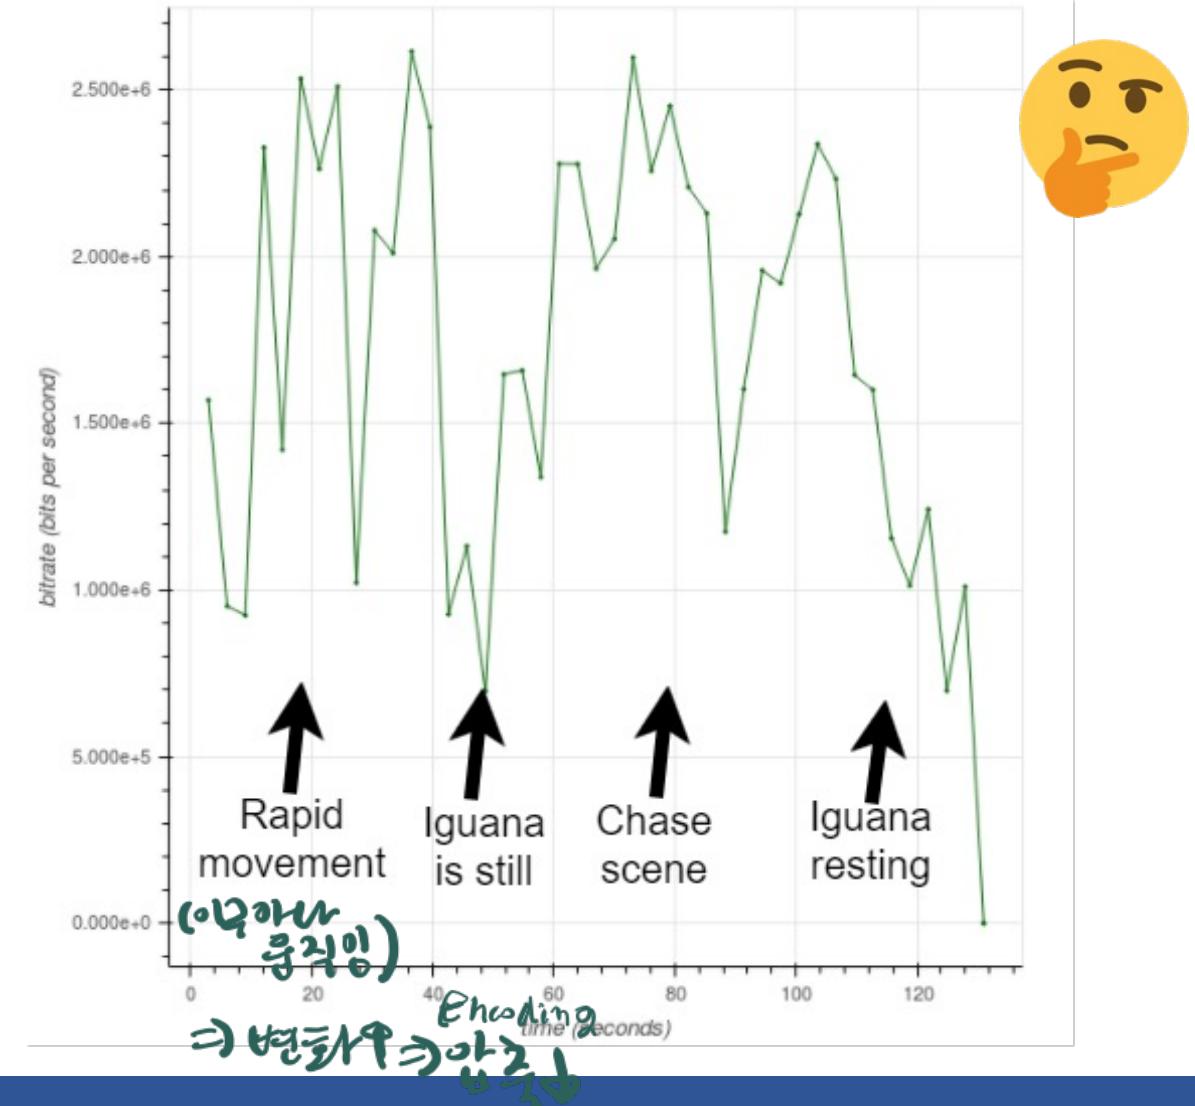
콘텐츠가 MPEG-DASH 클라이언트에 의해 재생되면 클라이언트는 비트레이트 적응(ABR) 알고리즘을 사용하여 재생 시 멈춤이나 재버퍼링을 일으키지 않고 다운로드할 수 있도록 가능한 최고 비트레이트의 세그먼트를 자동으로 선별한다.

Privacy concerns in DASH

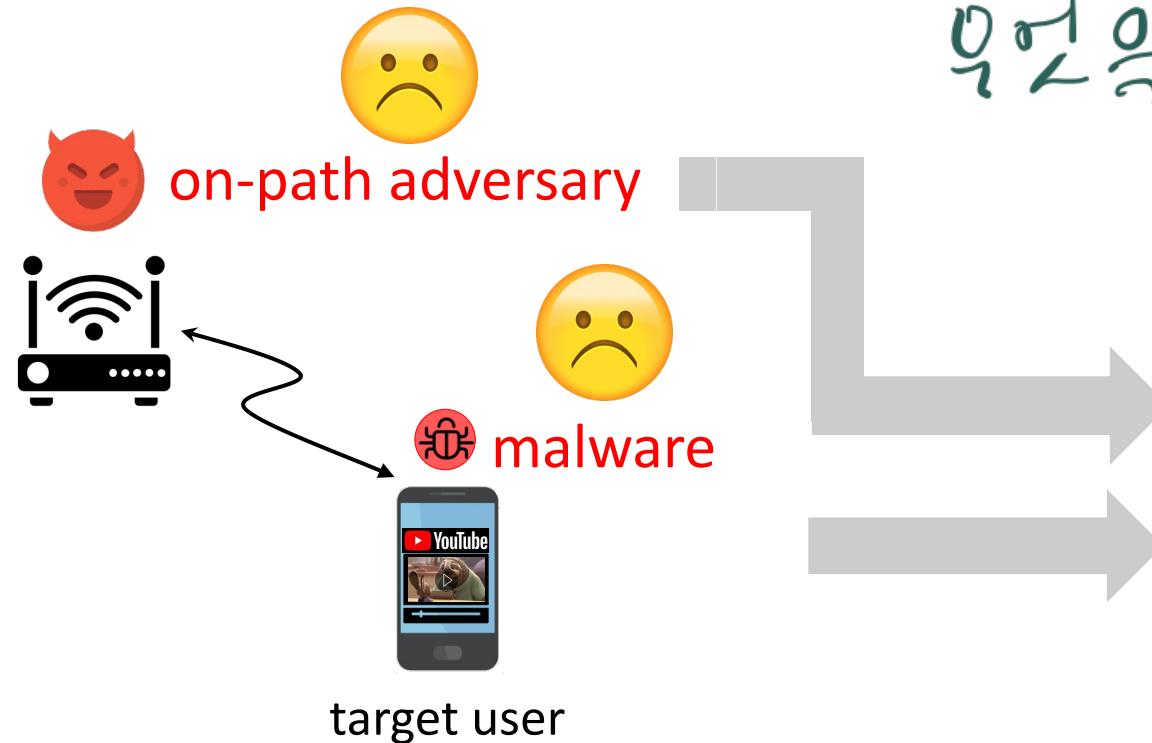
packet loss, delay



Privacy concerns in DASH

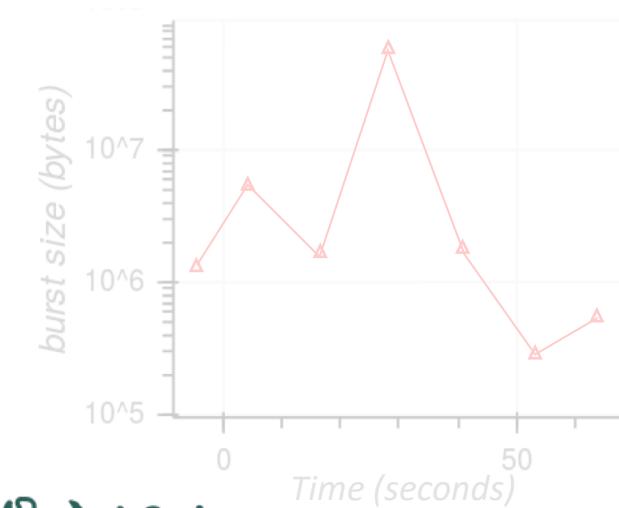


Infer Video Titles from *Encrypted Traffic*



wifi-traffic 220222
욕구를 보고 있는지 파악

traffic burst pattern
leaks video title
[USENIX Security 2017,
CODASPY 2017]



보는 방향을 210221 unique pattern 찾기

Moba: A Novel Video Inference Attack

(‘What are you watching?’ in Korean)

- **Identify YouTube videos** watched by the target user

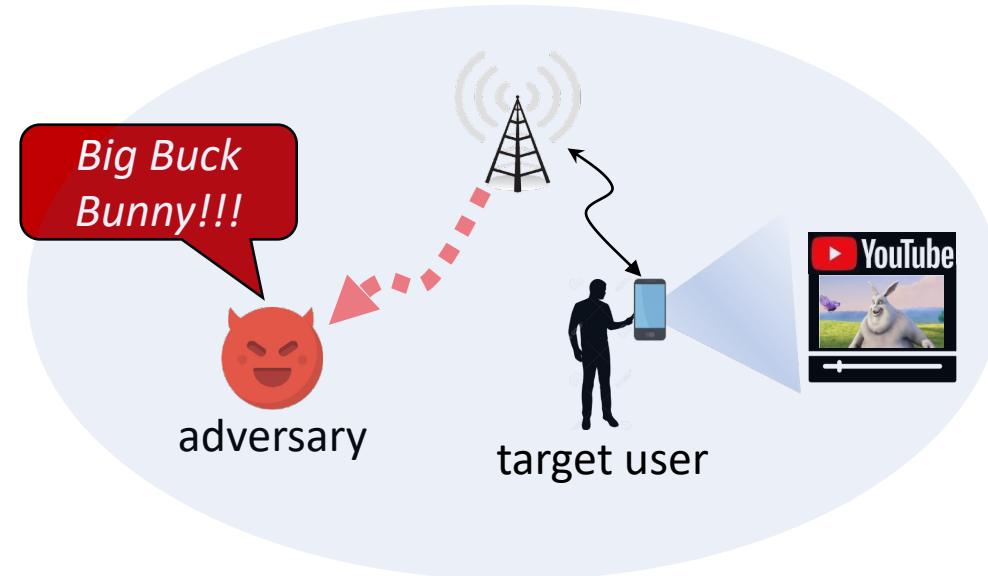
보내는 동영상 목록

- **Indirectly** estimate the downlink traffic burst pattern

burst patter 주파수

- No requirement for an **on-path adversary** or **malware installation**

터미널의 malware 설치 X

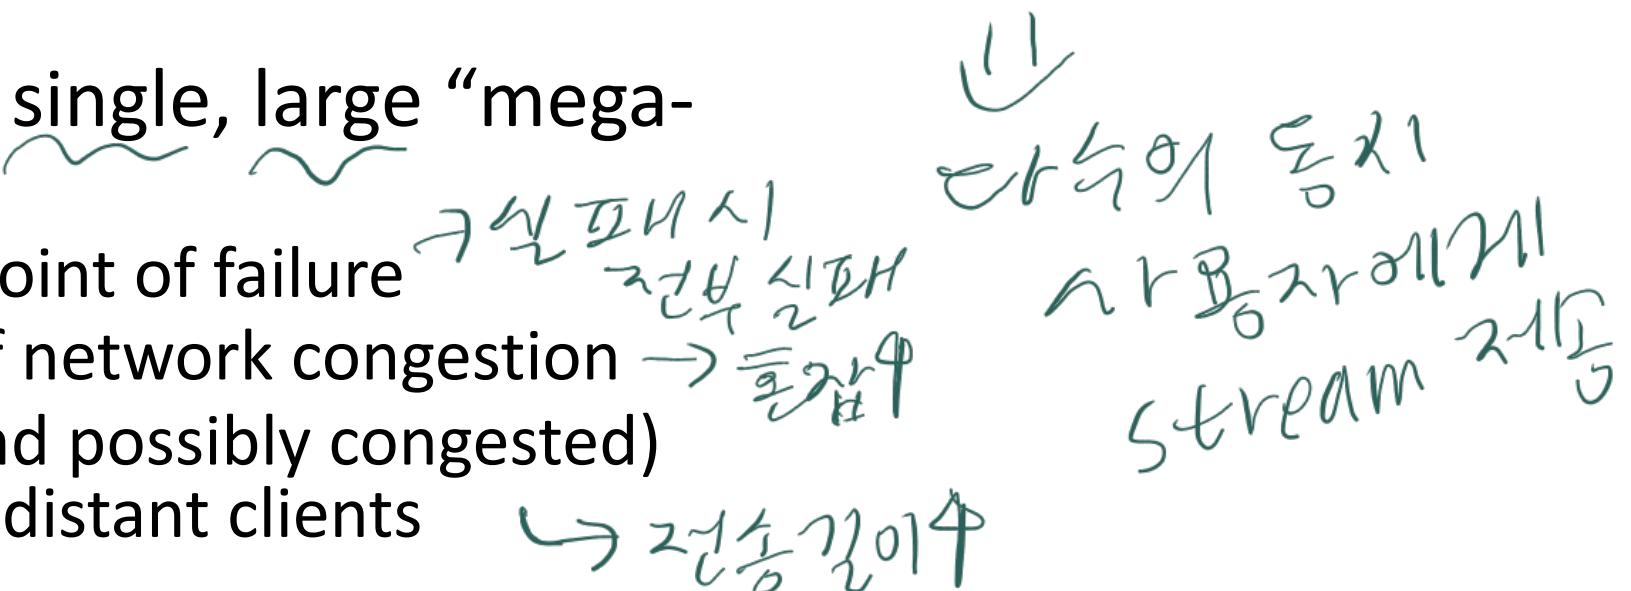


Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- **option 1:** single, large “mega-server”

- single point of failure
- point of network congestion
- long (and possibly congested) path to distant clients



....quite simply: this solution **doesn't scale**

한국어로

Content distribution networks (CDNs)

challenge: how to stream content (selected from millions of videos) to hundreds of thousands of *simultaneous* users?

- **option 2:** store/serve multiple copies of videos at multiple geographically distributed sites (**CDN**) → *인터넷으로 혼란된 복수의 캐시를*
 - **enter deep:** push CDN servers deep into many access networks
 - close to users → *액세스 네트워크에 있는 복수의 캐시를*
 - Akamai: 240,000 servers deployed (*240,000 개의 서버를*)
in > 120 countries (2015) *120 개 이상의 국가에서*
 - **bring home:** smaller number (10's) of larger clusters in POPs near access nets
 - used by Limelight → *액세스 네트워크에 있는 클러스터의 복수의 캐시를*
클리리스터(대량의 콘텐츠를) *EM21*



Akamai

복수의
캐시를



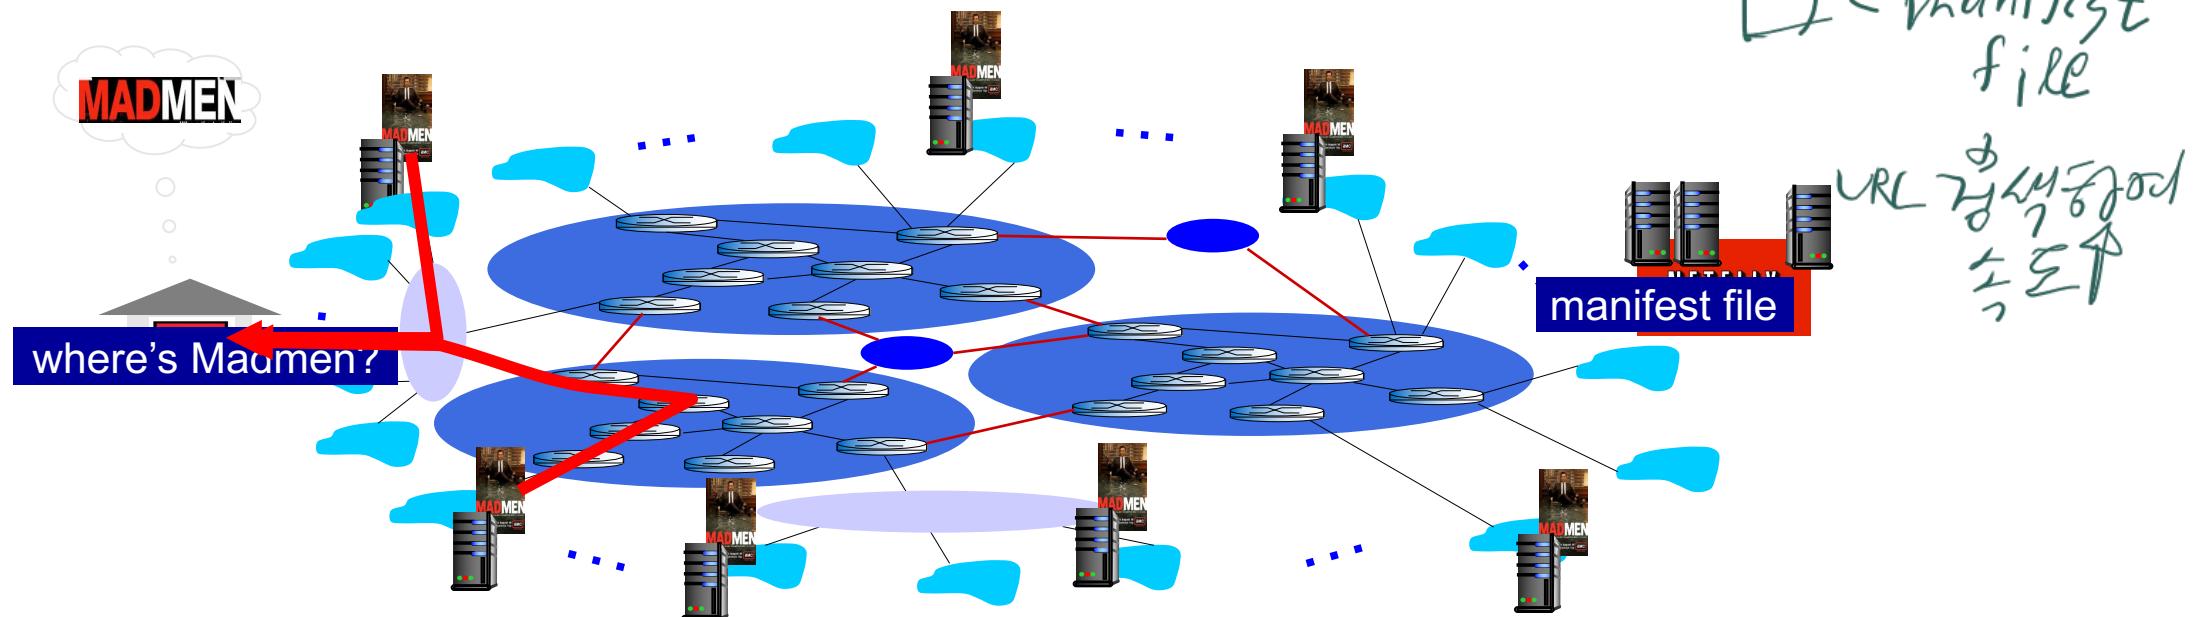
& 혼란된

EM21

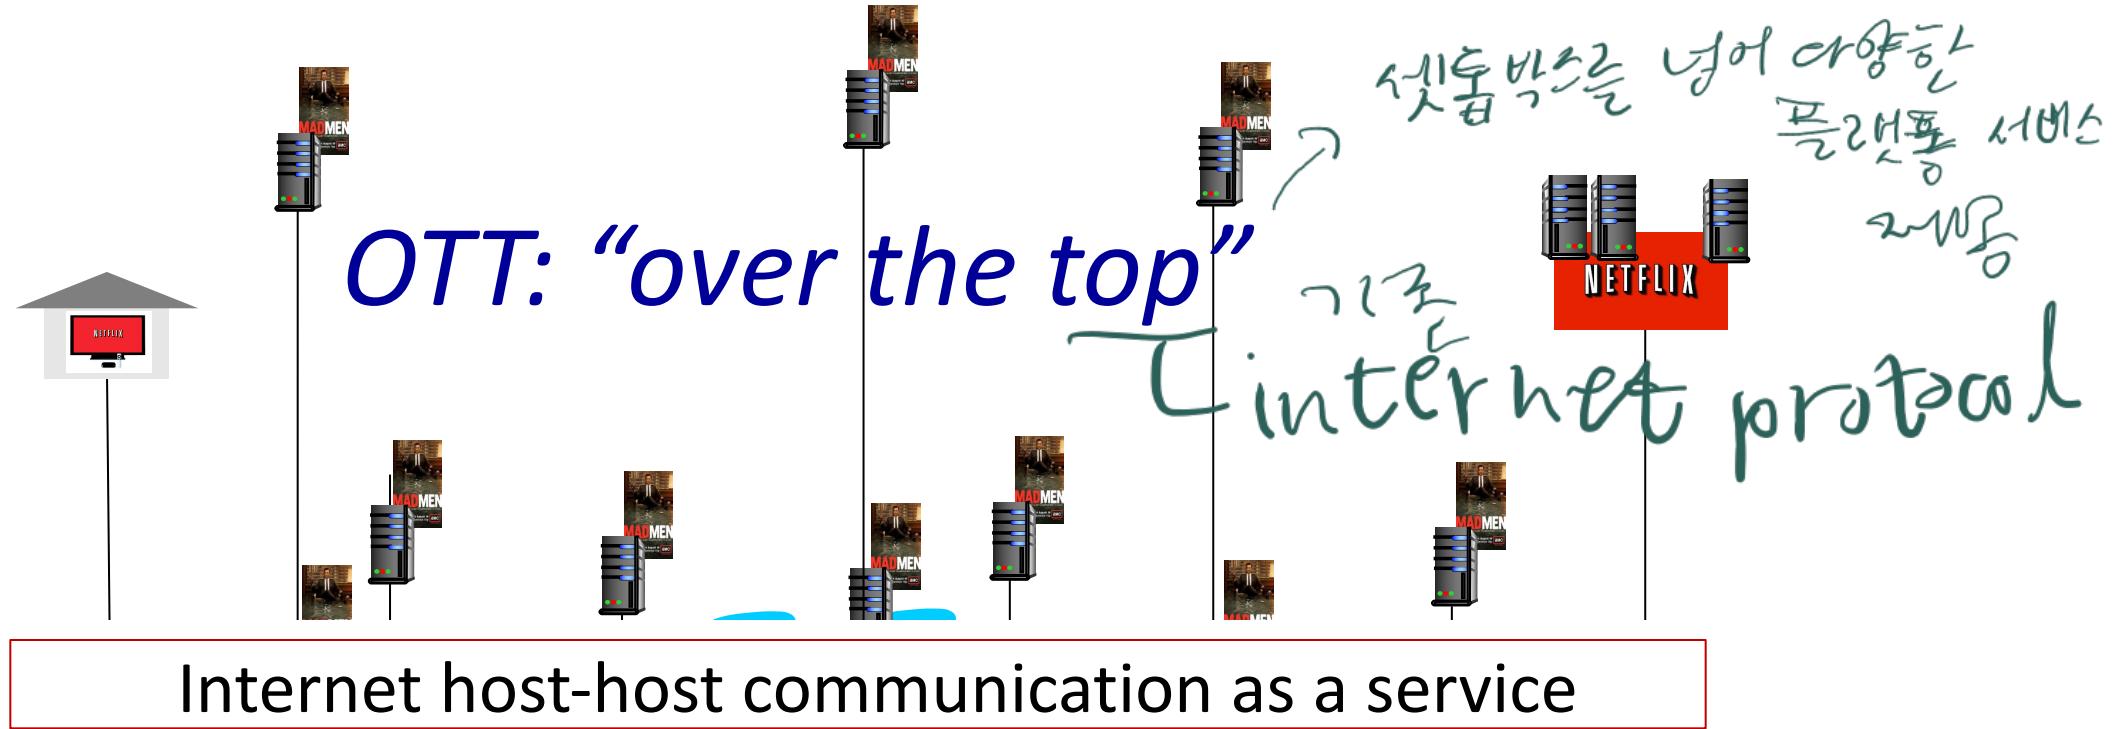
Application Layer: 2 32

Content distribution networks (CDNs)

- CDN: stores copies of content (e.g. MADMEN) at CDN nodes
- subscriber requests content, service provider returns manifest
 - using manifest, client retrieves content at highest supportable rate
 - may choose different rate or copy if network path congested



Content distribution networks (CDNs)



OTT challenges: coping with a congested Internet from the “edge”

- what content to place in which CDN node?
 - from which CDN node to retrieve content? At which rate?

71색하다

Next...

- *Chapter 3.1 Transport-Layer Services*
- *Chapter 3.2 Multiplexing and Demultiplexing*
- *Chapter 3.3 Connectionless Transport: UDP*