

Transport Layer: Reliable Data Transfer

October 8, 2024

Min Suk Kang

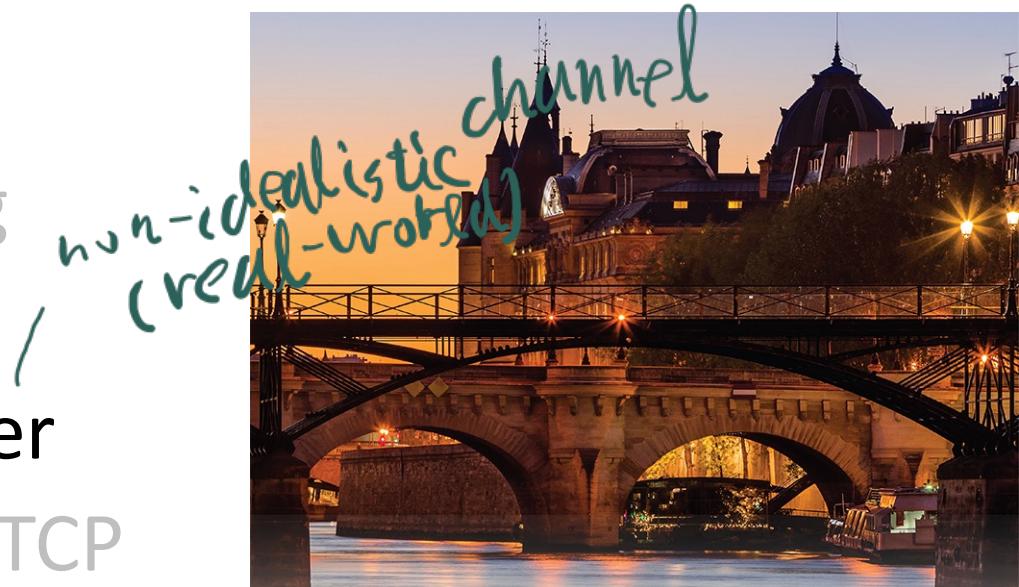
Associate Professor

School of Computing/Graduate School of Information Security



Chapter 3: roadmap

- Transport-layer services
- Multiplexing and demultiplexing
- Connectionless transport: UDP
- **Principles of reliable data transfer**
- Connection-oriented transport: TCP
- Principles of congestion control
- TCP congestion control
- Evolution of transport-layer functionality



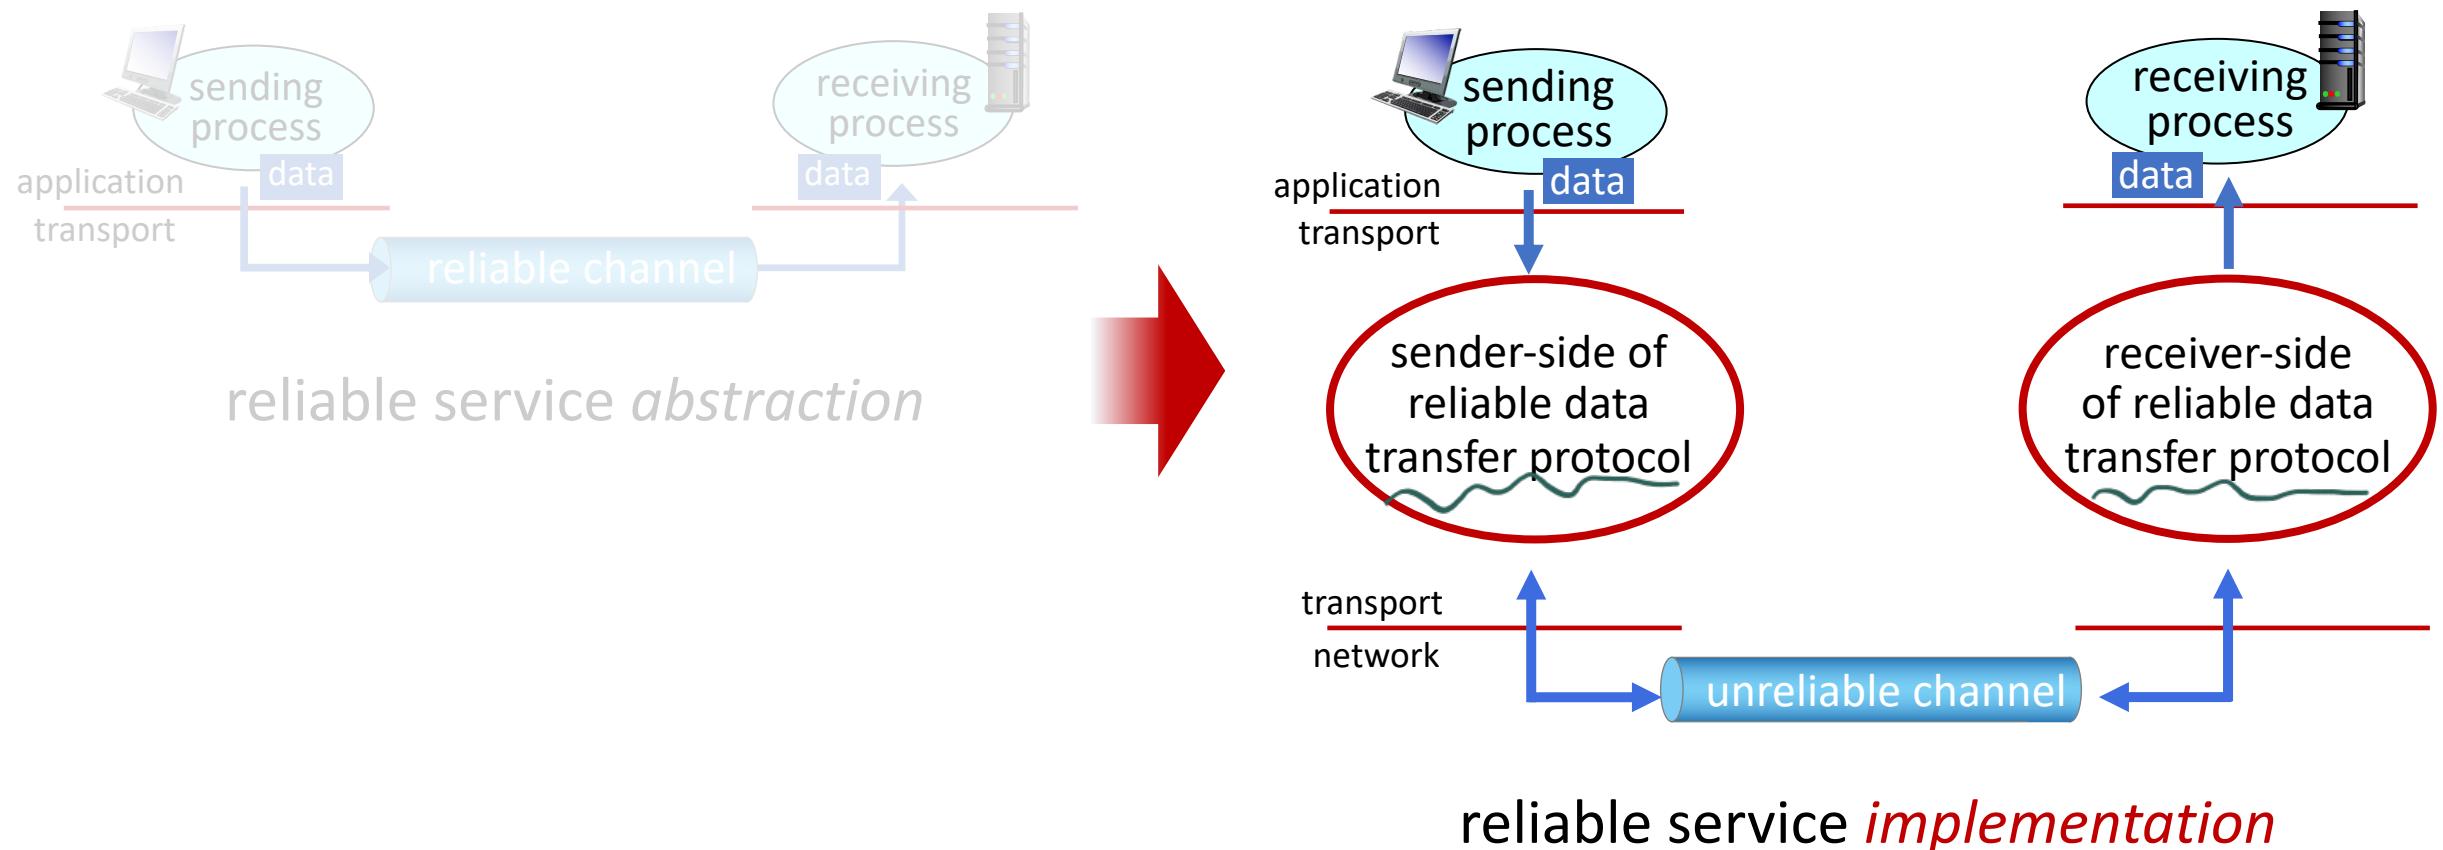
Principles of reliable data transfer



reliable service *abstraction*

즉상호작용

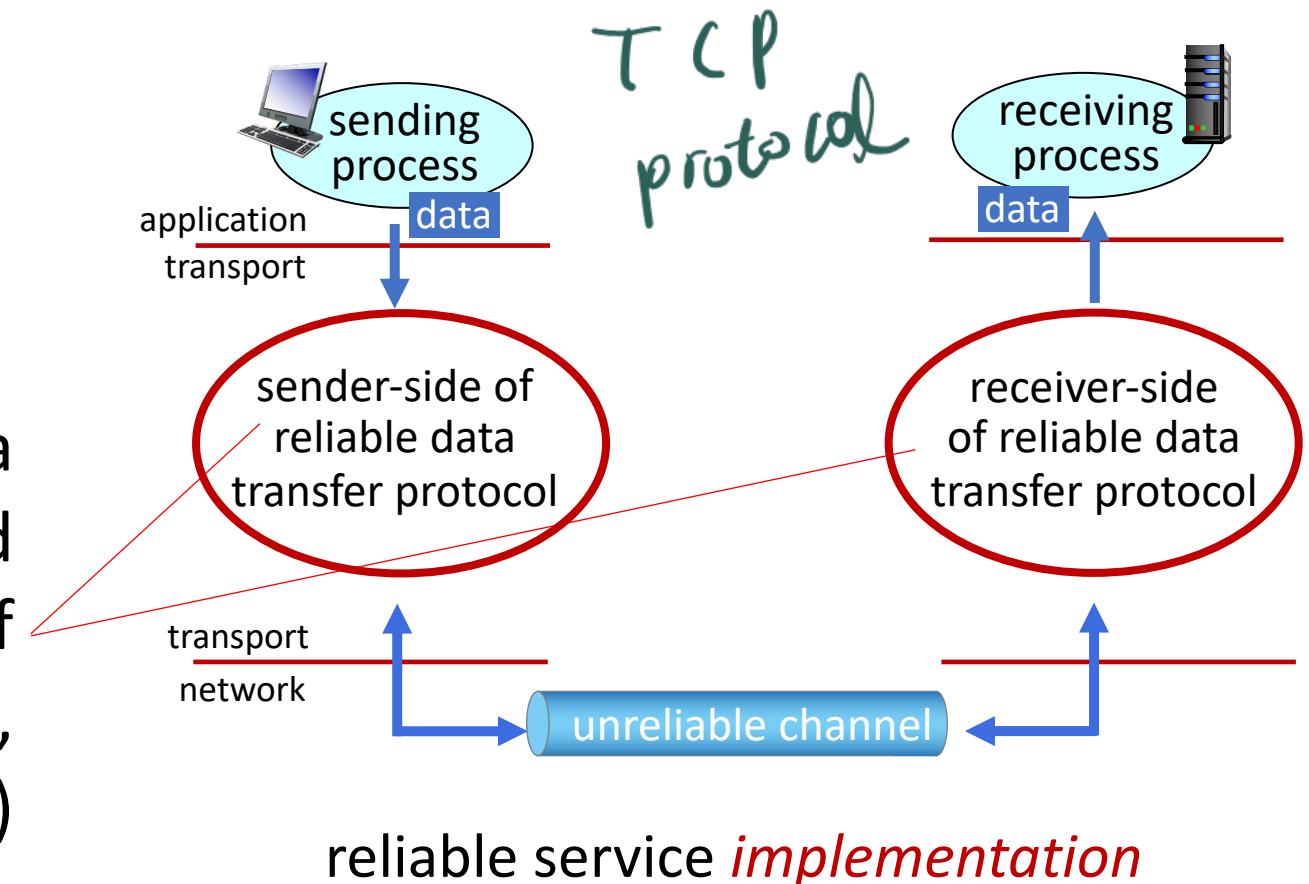
Principles of reliable data transfer



국어

Principles of reliable data transfer

Complexity of reliable data transfer protocol will depend (strongly) on characteristics of unreliable channel (lose, corrupt, reorder data?)

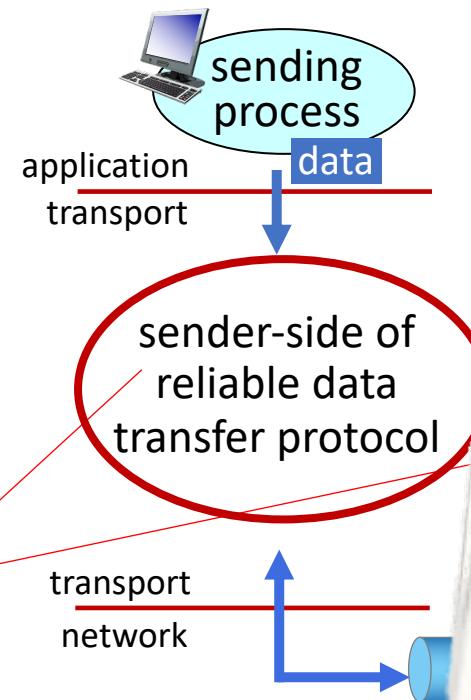


Principles of reliable data transfer

한국어로 된 노트

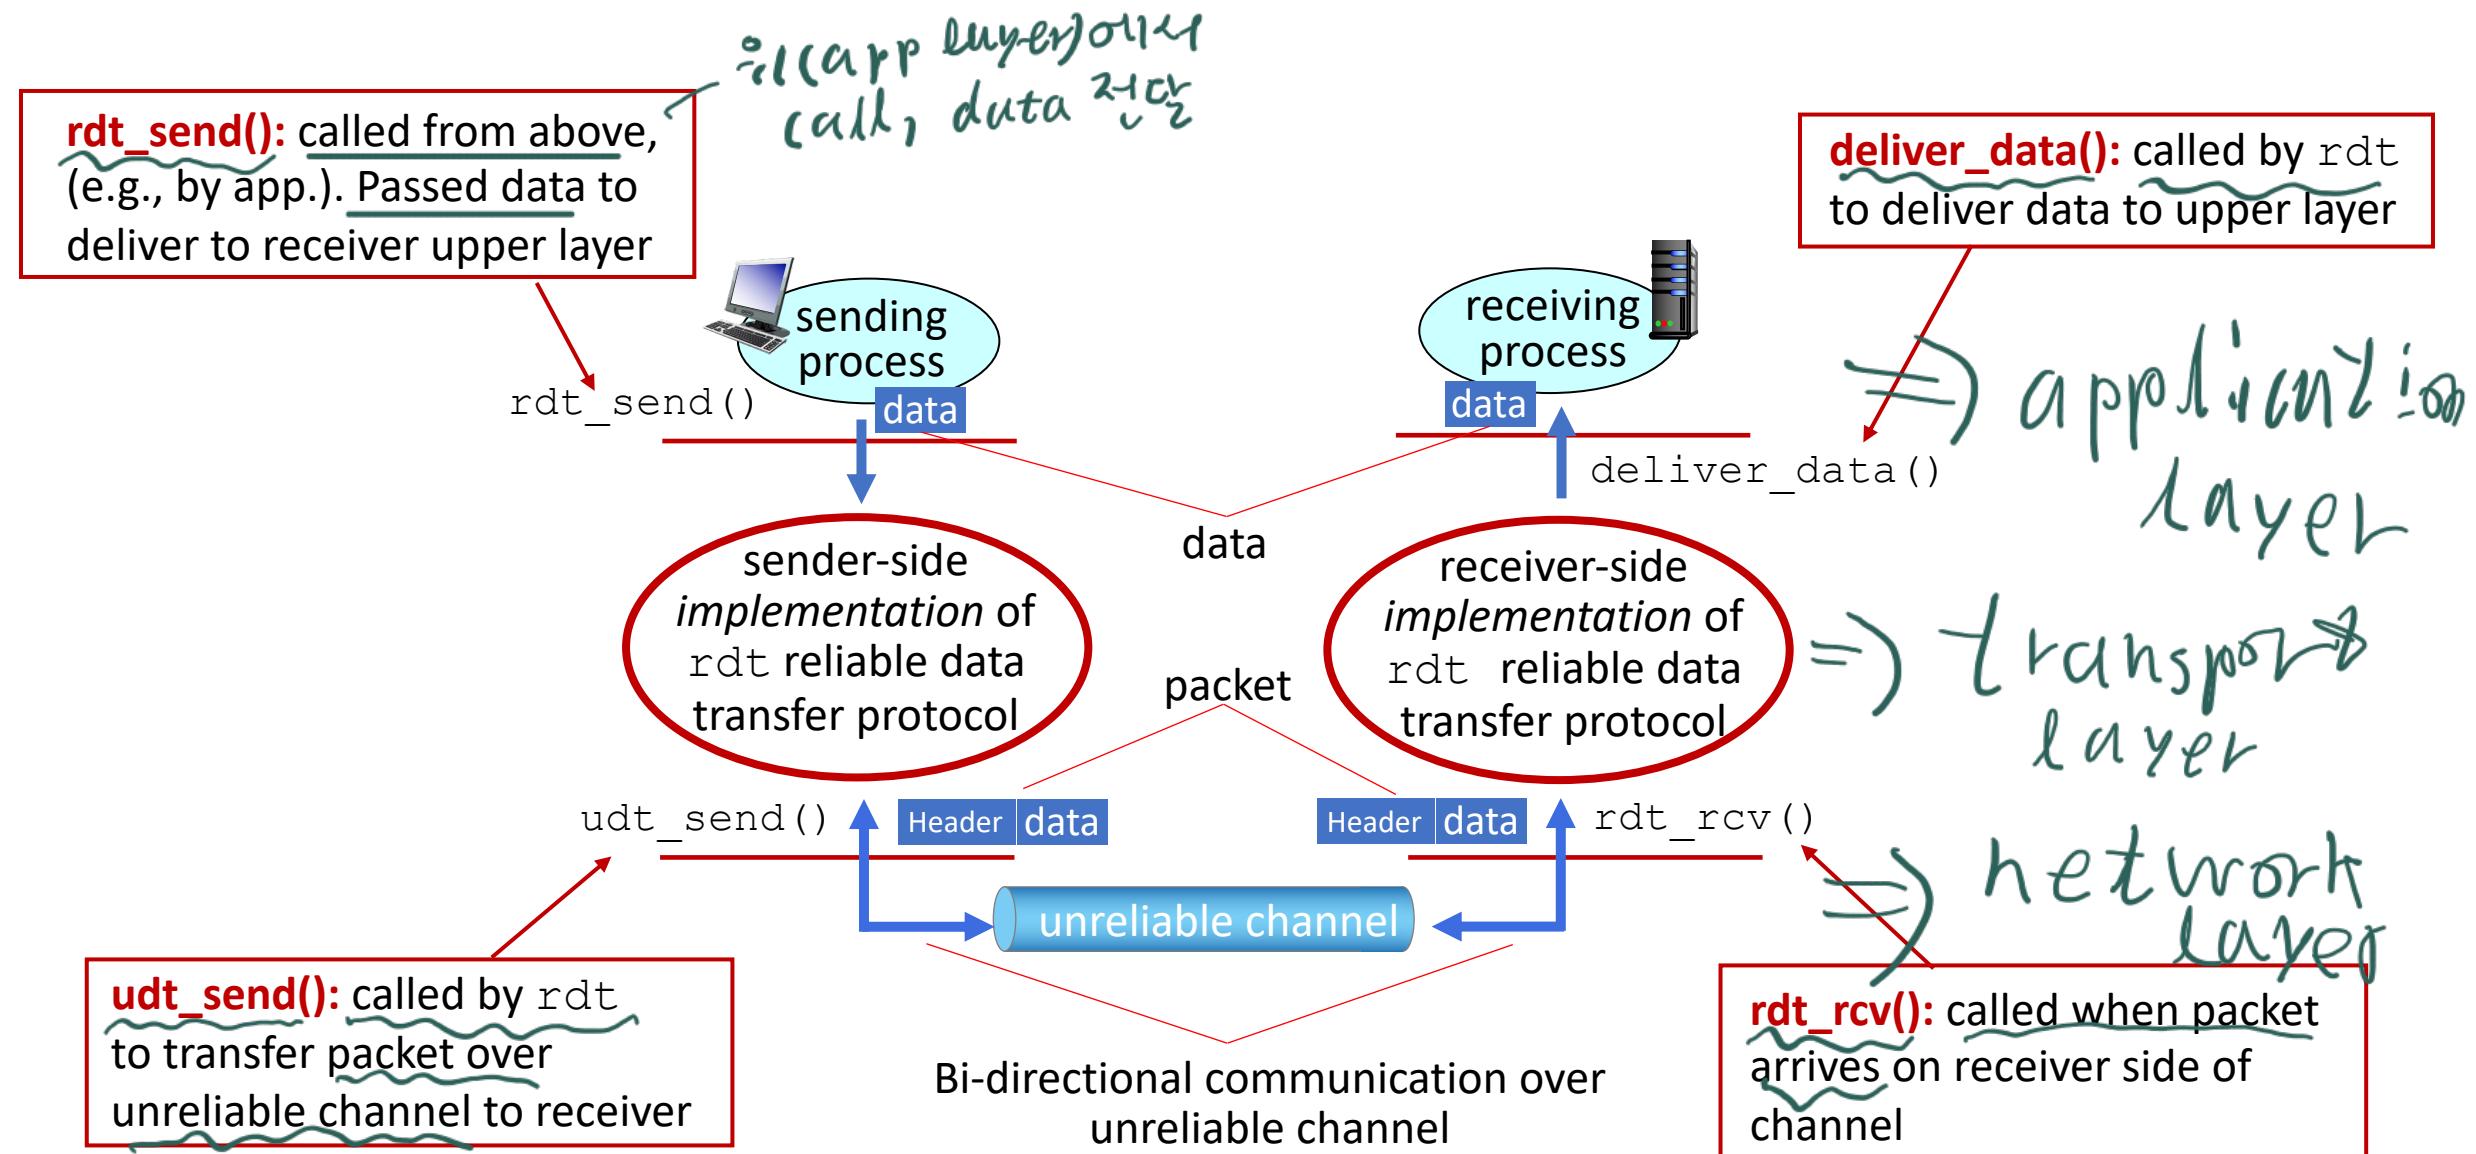
Sender, receiver do *not* know
the “state” of each other, e.g.,
was a message received?

- unless communicated via a message



reliable service *implementation*

Reliable data transfer protocol (rdt): interfaces

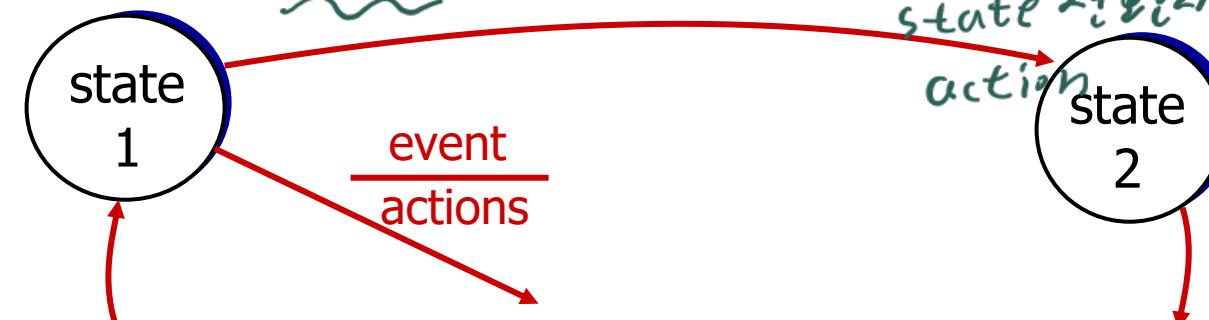


Reliable data transfer: getting started

We will:

- incrementally develop sender, receiver sides of reliable data transfer protocol (rdt)
- consider only unidirectional data transfer \Rightarrow 단방 통신
• but control info will flow in both directions!
- use finite state machines (FSM) to specify sender, receiver

state: when in this “state”
next state uniquely
determined by next
event



rdt1.0: reliable transfer over a reliable channel

↳ simplest

- assumption: underlying channel perfectly reliable

- no bit errors
- no loss of packets

⇒ *perfect channel*

- separate** FSMs for sender, receiver:

- sender sends data into underlying channel
- receiver reads data from underlying channel



rdt2.0: channel with bit errors

↳ bit error 371

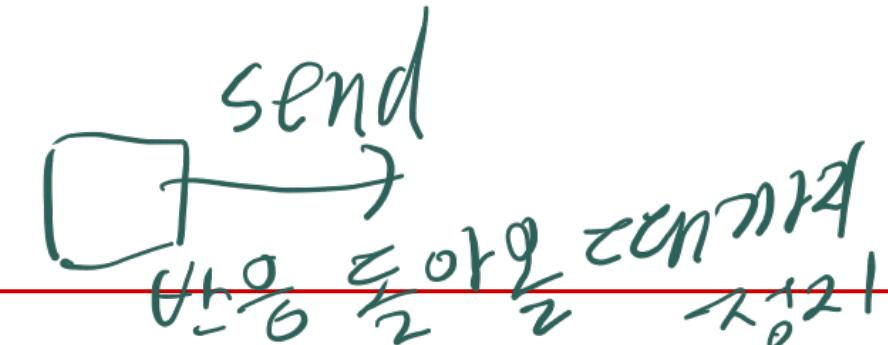
- assumption: underlying channel may flip bits in packet
 - checksum (e.g., Internet checksum) to detect bit errors
 - no loss of packets
 - *the question:* how to recover from errors?

How do humans recover from “errors” during conversation?

↳ recall

rdt2.0: channel with bit errors

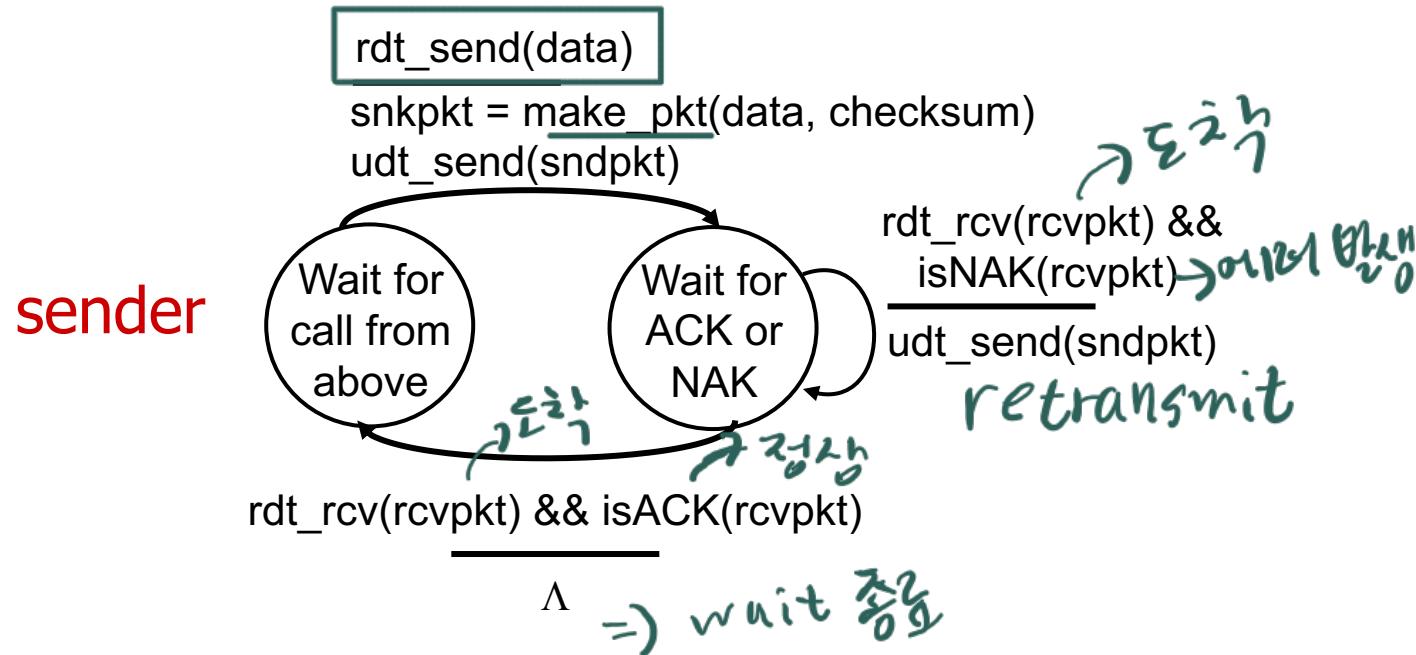
- underlying channel may flip bits in packet
 - checksum to detect bit errors
- the question:* how to recover from errors?
 - acknowledgements (ACKs):* receiver explicitly tells sender that pkt received OK → 성공 시 응답
 - negative acknowledgements (NAKs):* receiver explicitly tells sender that pkt had errors → 실패 시 응답
 - sender *retransmits* pkt on receipt of NAK



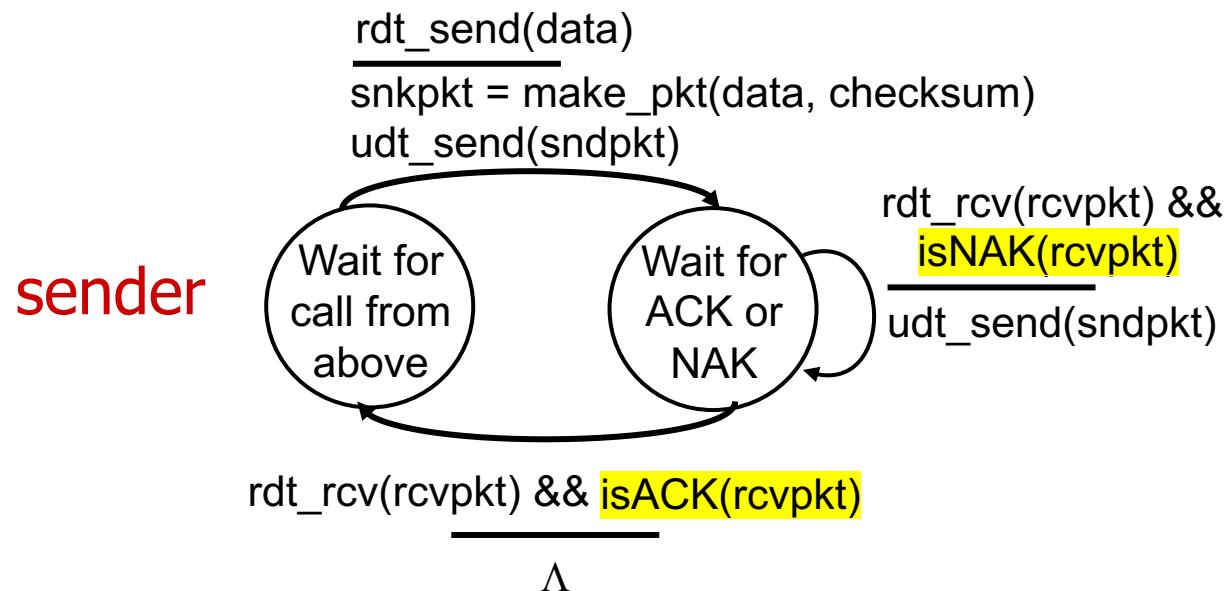
stop and wait

sender sends one packet, then waits for receiver response

rdt2.0: FSM specifications



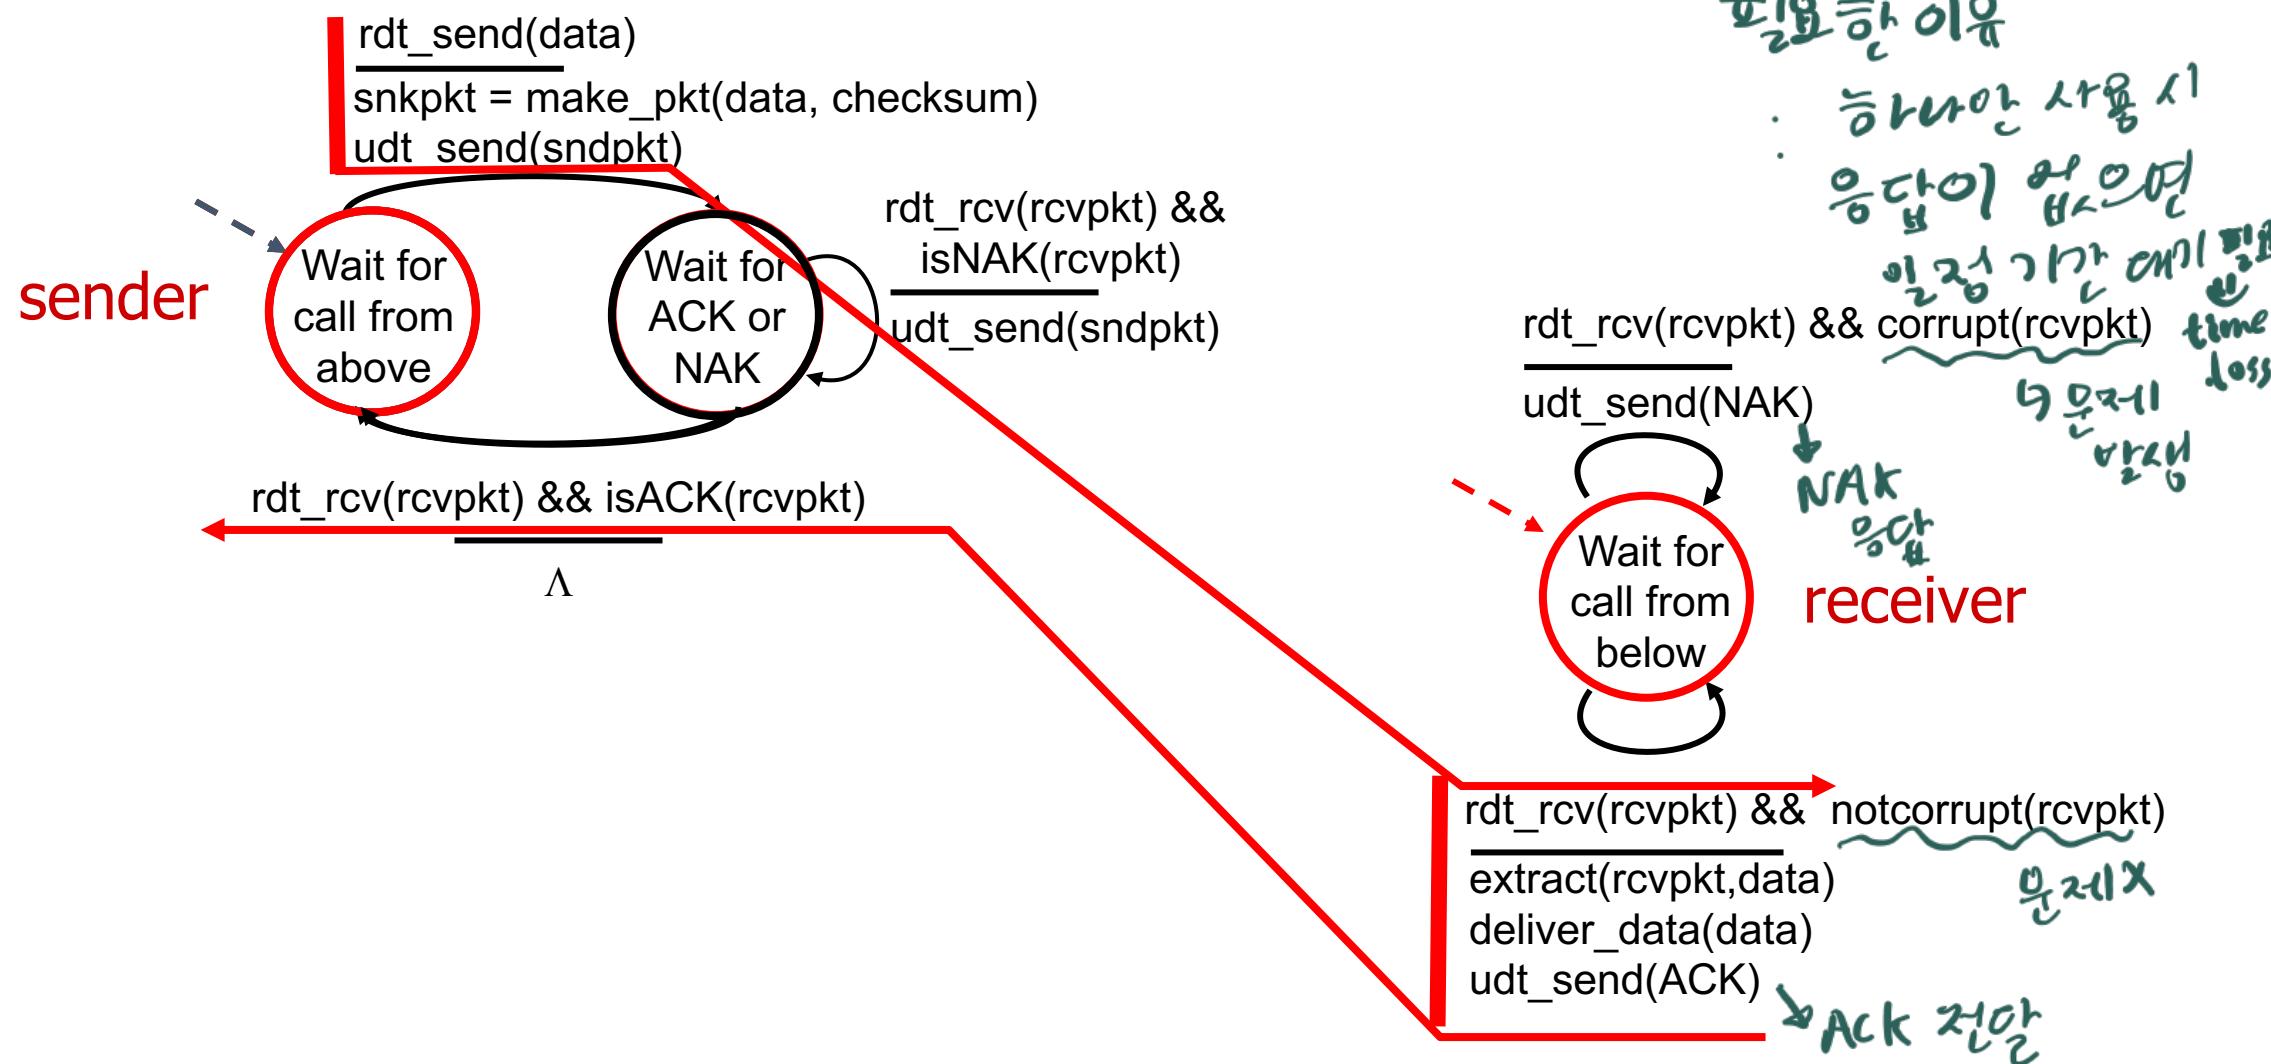
rdt2.0: FSM specification



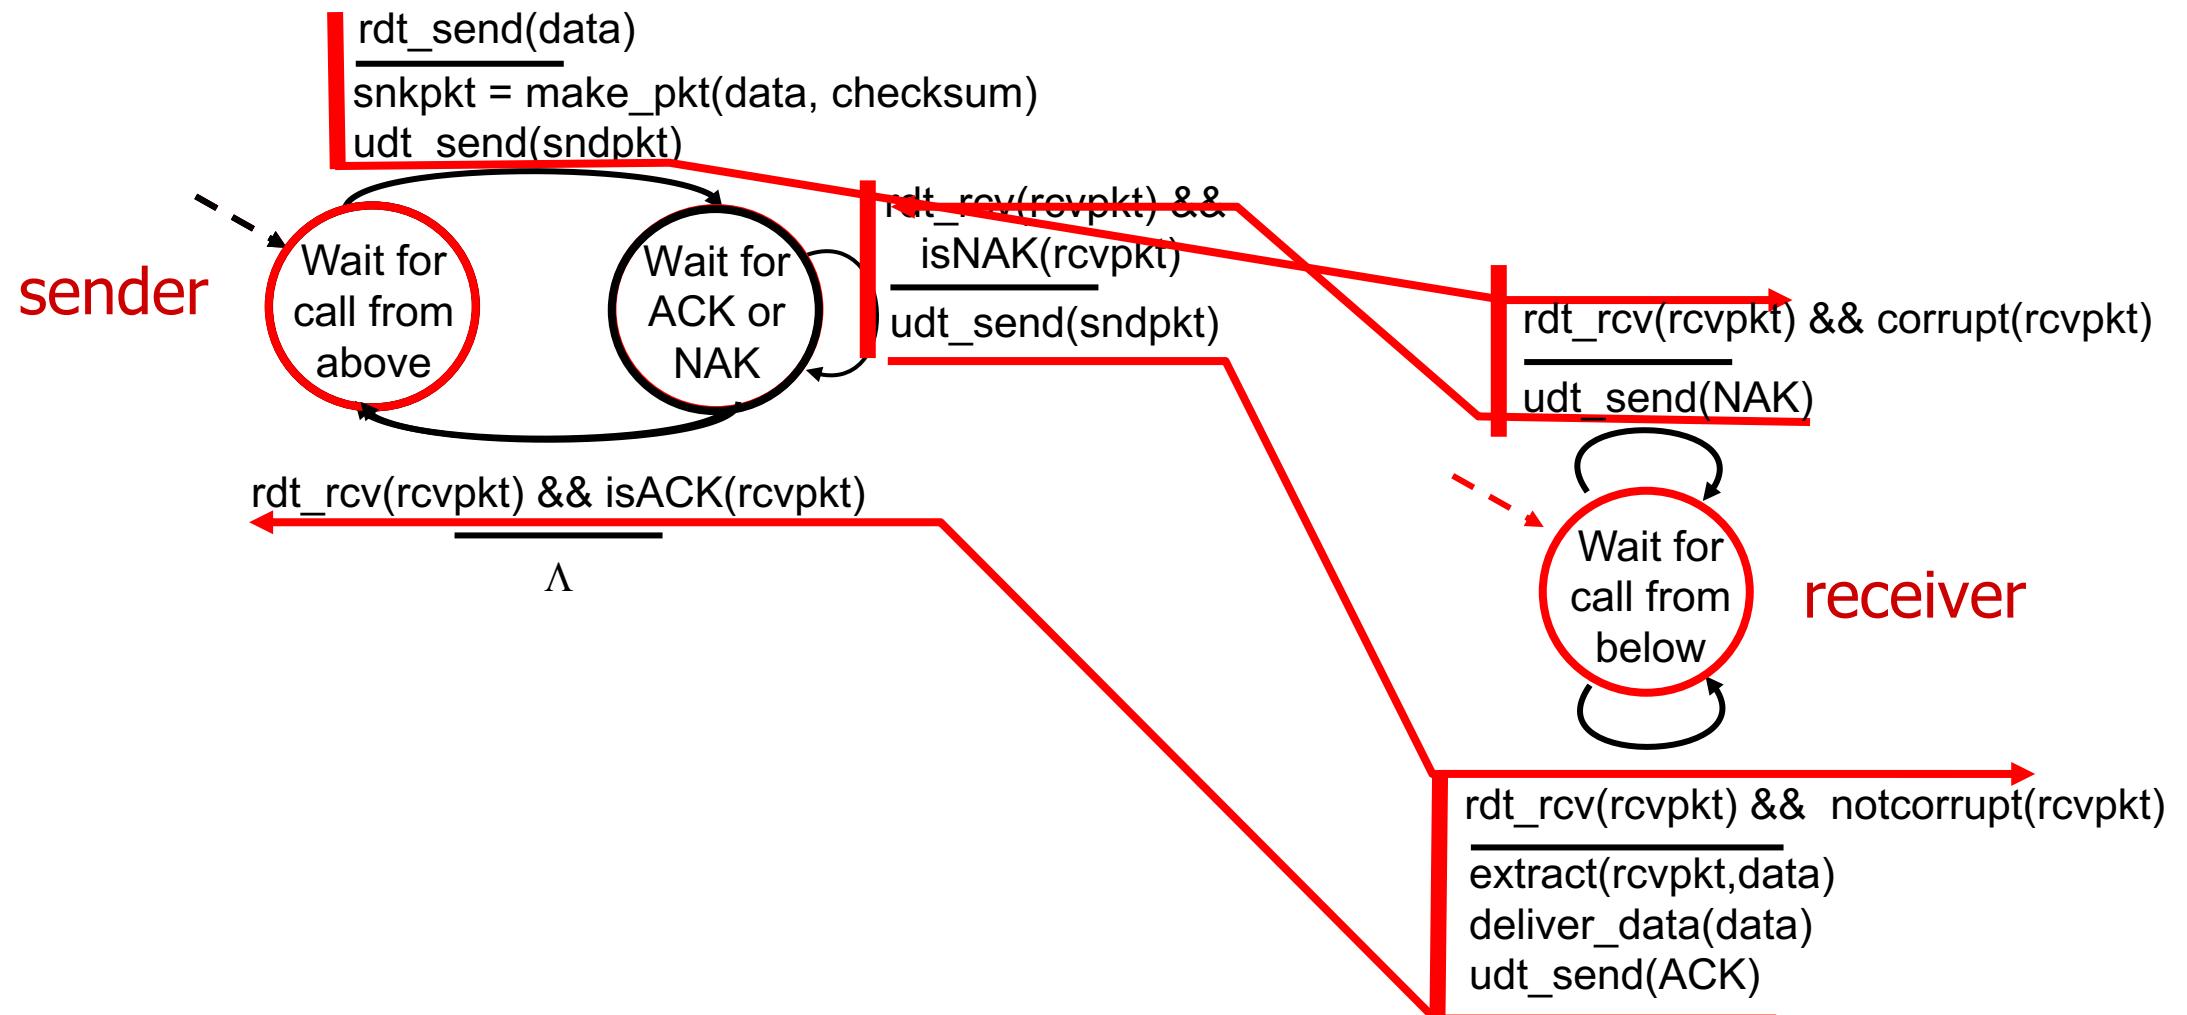
Note: “state” of receiver (did the receiver get my message correctly?) isn’t known to sender unless somehow communicated from receiver to sender

- that’s why we need a protocol!

rdt2.0: operation with no errors



rdt2.0: corrupted packet scenario



rdt2.0 has a fatal flaw!

최종적인 결함

ACK/NAK corrupted?

⇒ sender's state
of the retransmit
will be of the original
duplication packet

retransmit
최종 목적지 가능

handling duplicates:

- sender retransmits current pkt if ACK/NAK corrupted
- sender adds sequence number to each pkt
- receiver discards (doesn't deliver up) duplicate pkt

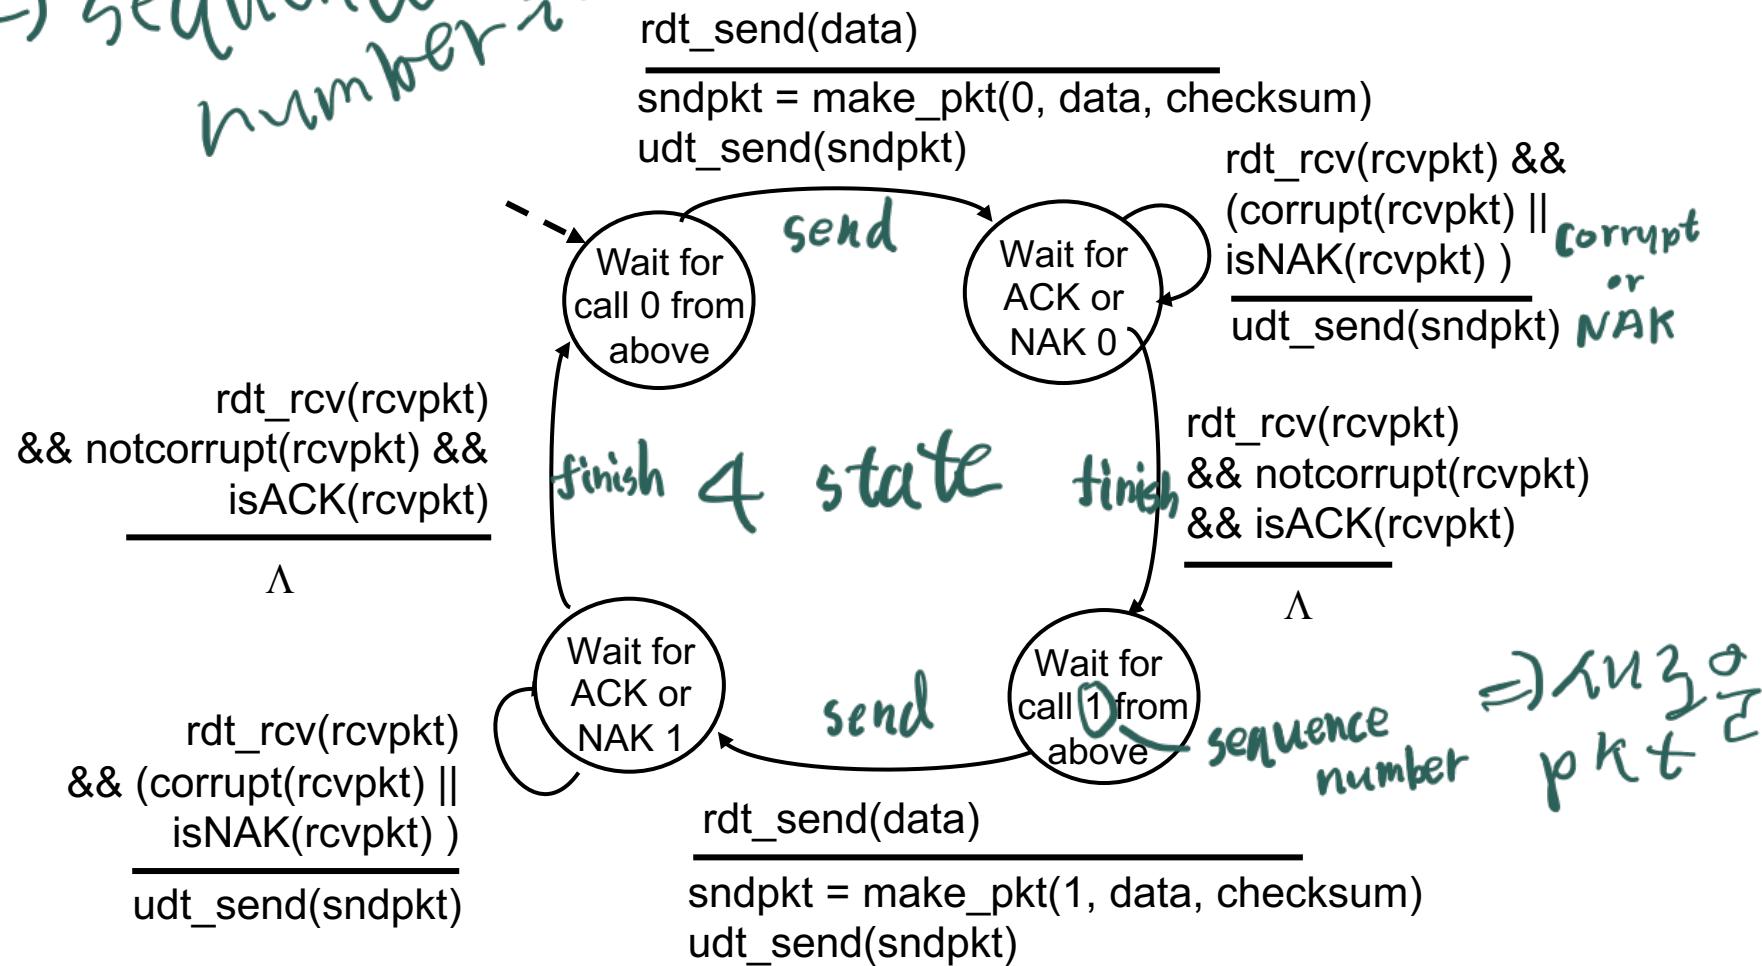
↳ 중복 제거
(retransmit)
하지만 가능

stop and wait

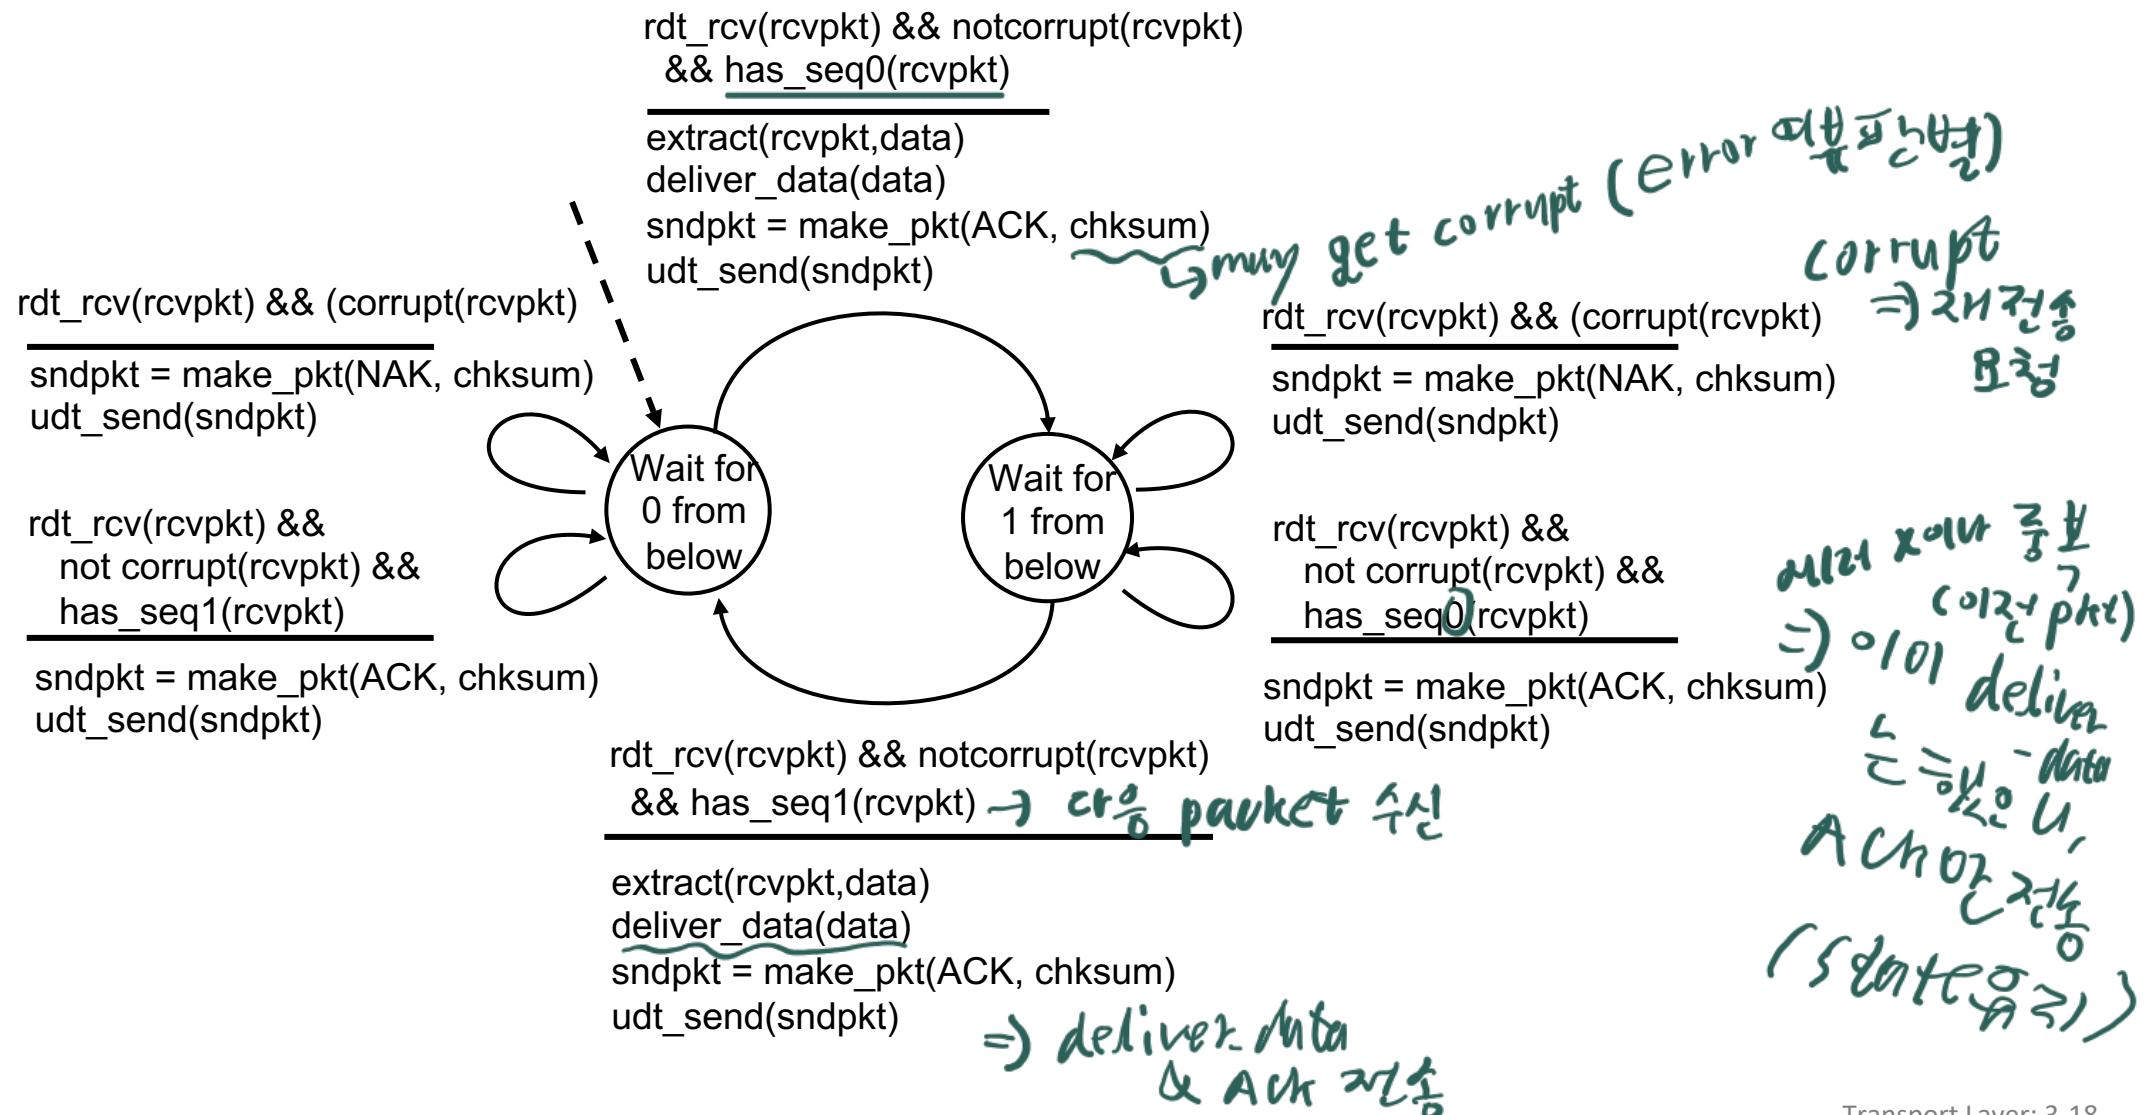
sender sends one packet, then waits for receiver response

rdt2.1: sender, handling garbled ACK/NAKs

↳ sequence number



rdt2.1: receiver, handling garbled ACK/NAKs



rdt2.1: discussion

sender:

- seq # added to pkt
- two seq. #s (0,1) will suffice.
Why? 번갈아가며 재전송 여부(T/F) 표기
- must check if received ACK/NAK corrupted
- twice as many states
 - state must “remember” whether “expected” pkt should have seq # of 0 or 1

증분

receiver:

- must check if received packet is duplicate
 - state indicates whether 0 or 1 is expected pkt seq #
- note: receiver can *not* know if its last ACK/NAK received OK at sender

↳ 자신의 끝에

ACK, NAK가

제대로 도착

되었는지 알

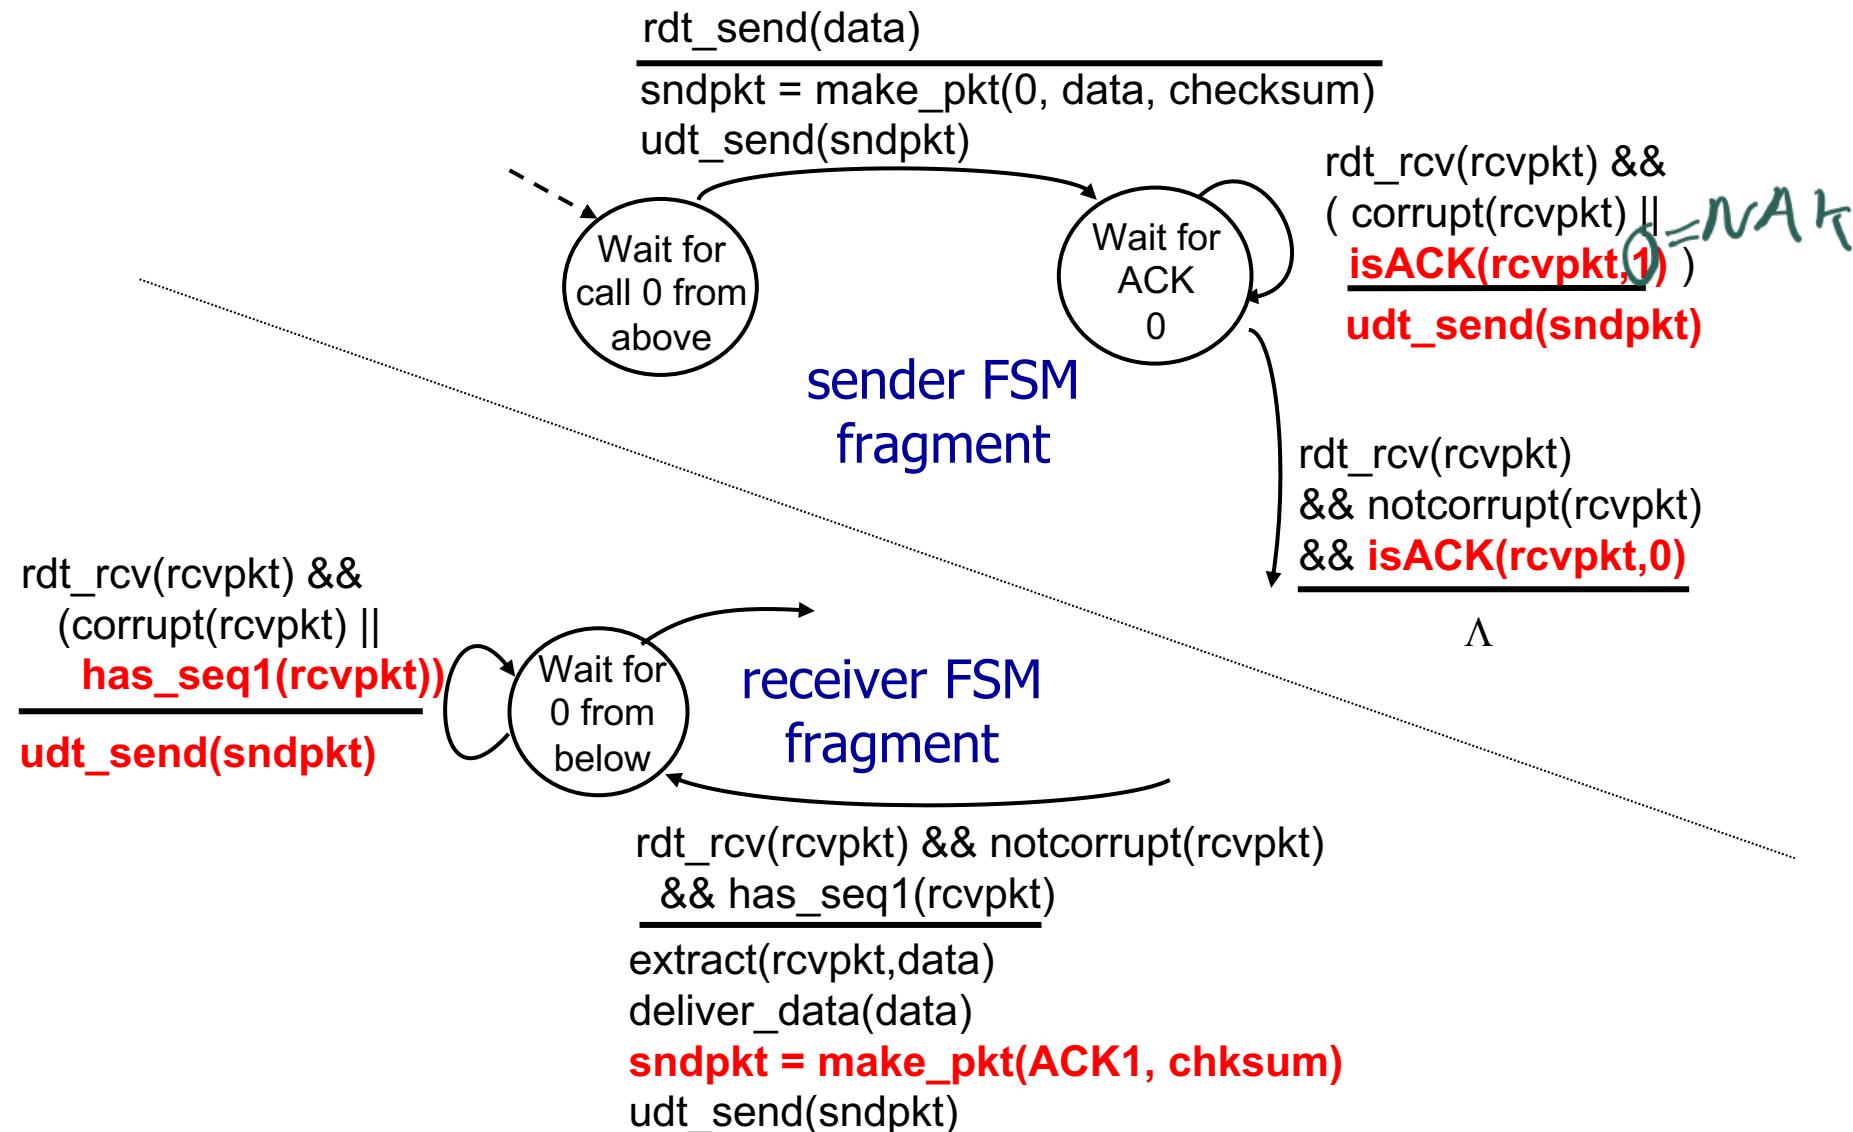
→ 중복 발생을 막아야 한다(ACK)

rdt2.2: a NAK-free protocol

- same functionality as rdt2.1, using ACKs only
 - instead of NAK, receiver sends ACK for last pkt received OK
 - receiver must *explicitly* include seq # of pkt being ACKed → 오른쪽에 그림
 - duplicate ACK at sender results in same action as NAK:
retransmit current pkt 수신된
pkt의
seq# 전달
⇒ 수신여부
확인 가능

rdt2.2: sender, receiver fragments

☞ 3.2 (2/4)



rdt3.0: channels with errors *and* loss

→ loss

New channel assumption: underlying channel can also *lose* packets (data, ACKs)

- checksum, sequence #s, ACKs, retransmissions will be of help ... but not quite enough

Q: How do *humans* handle lost sender-to-receiver words in conversation?

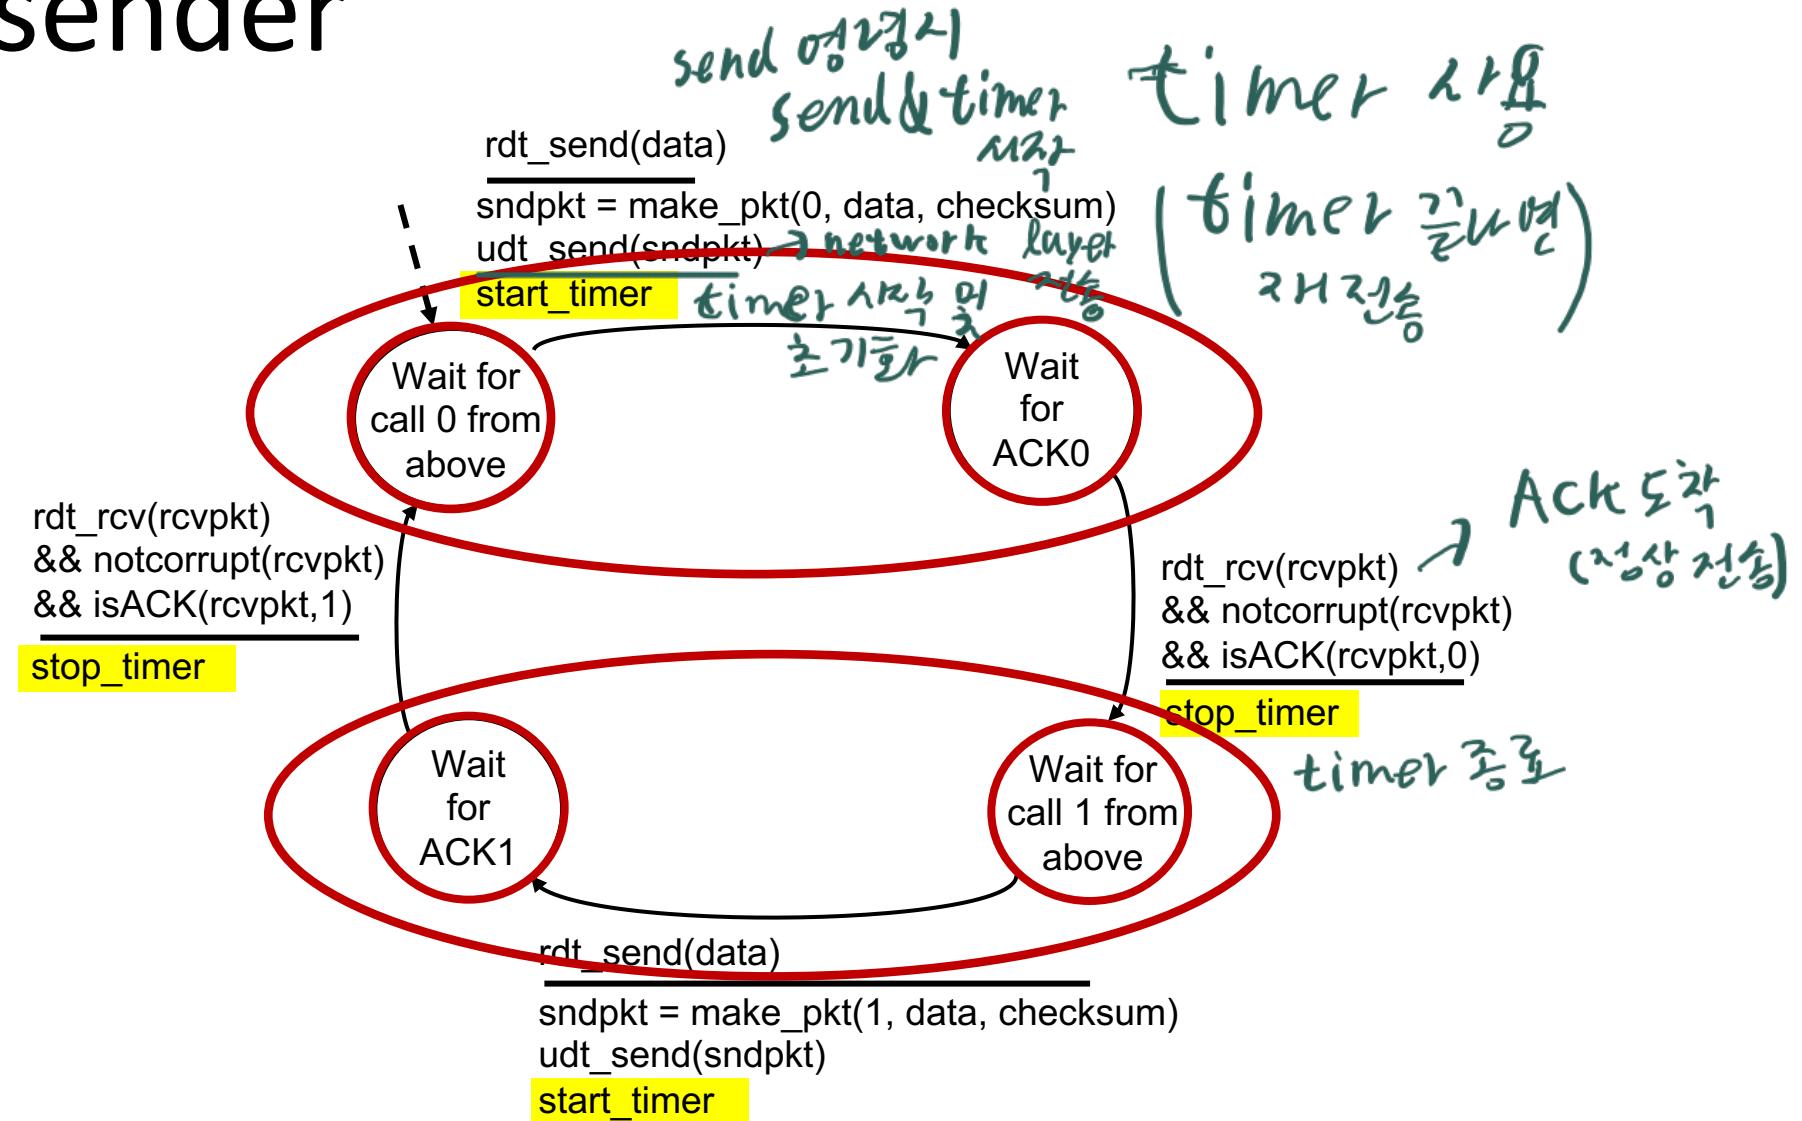
loss ⇒ 충분한 시간 후 질문 반복
(repeat)

rdt3.0: channels with errors *and* loss

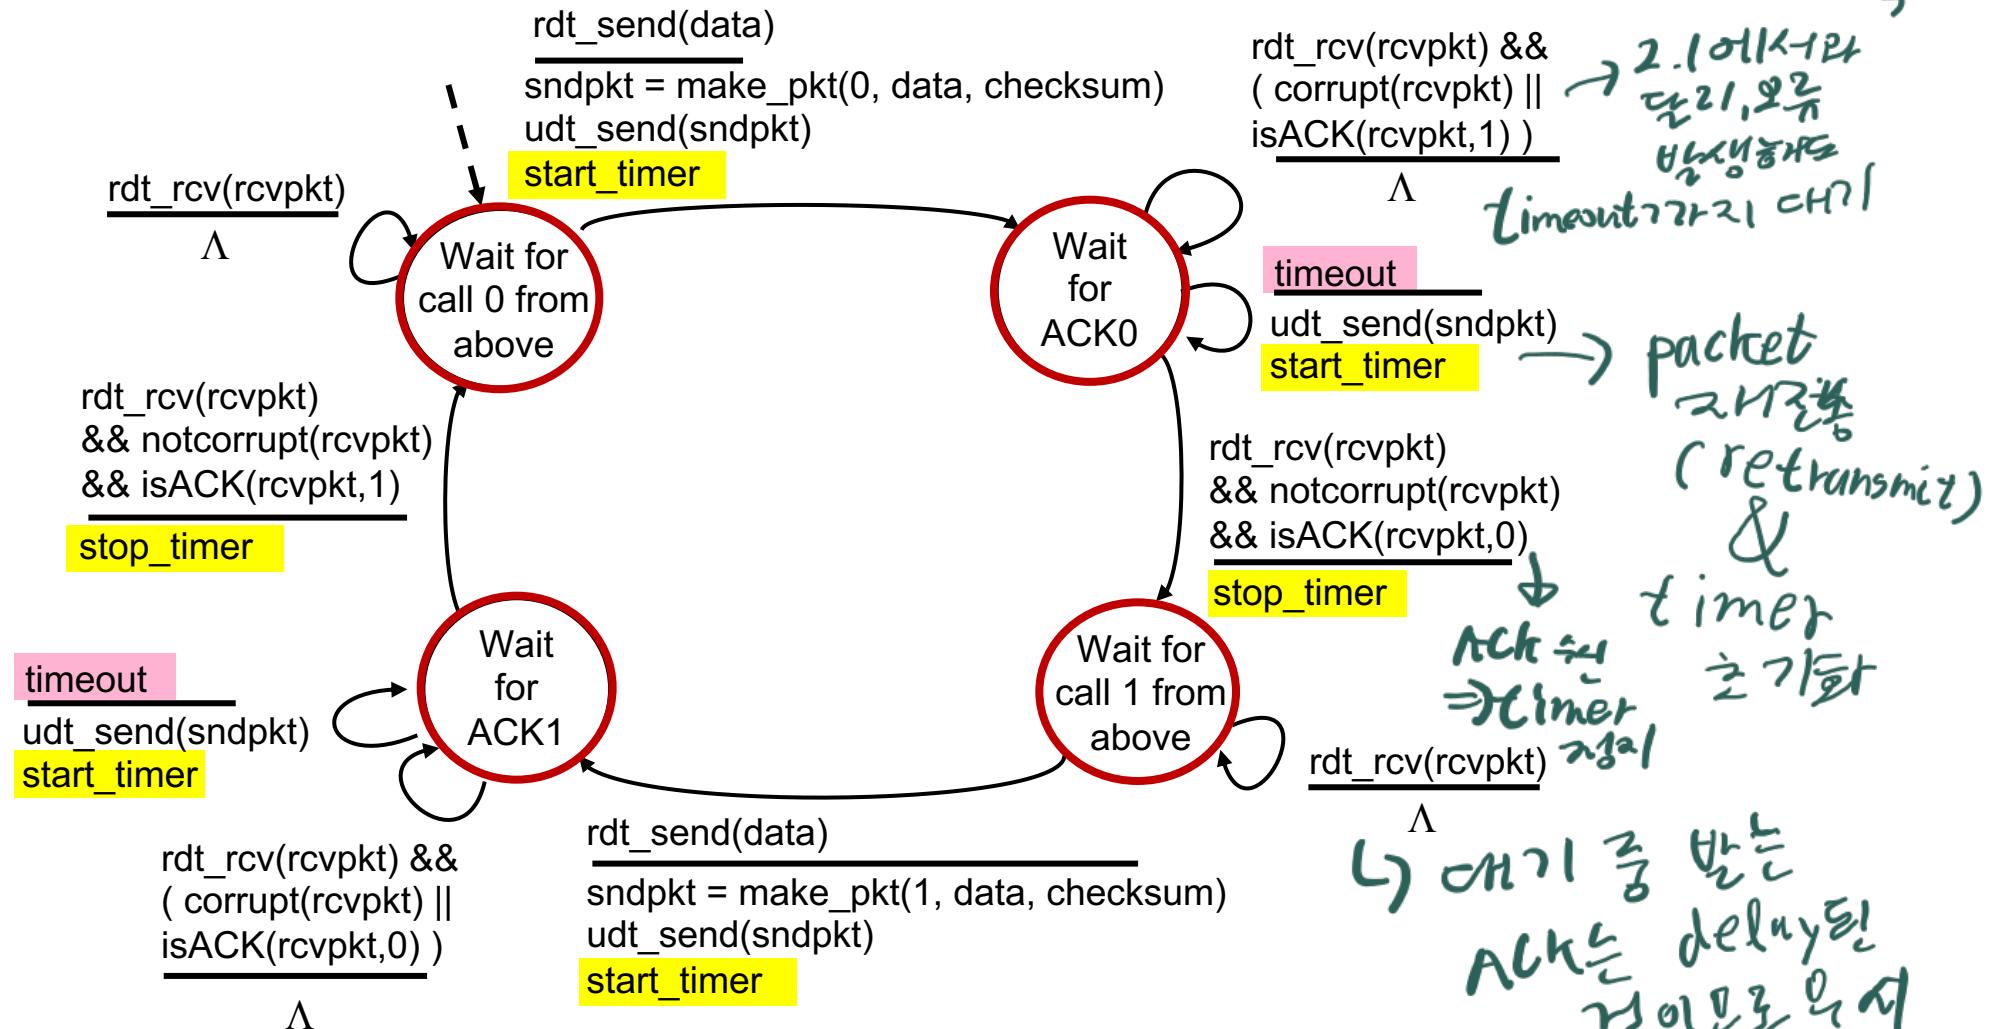
→ ACK $\frac{2}{2}$ ve'll 21nr2/
ch²/

- retransmits if no ACK received in this time
- if pkt (or ACK) just delayed (not lost):
 - retransmission will be duplicate, but seq #s already handles this!
 - receiver must specify seq # of packet being ACKed
- use countdown timer to interrupt after “reasonable” amount of time

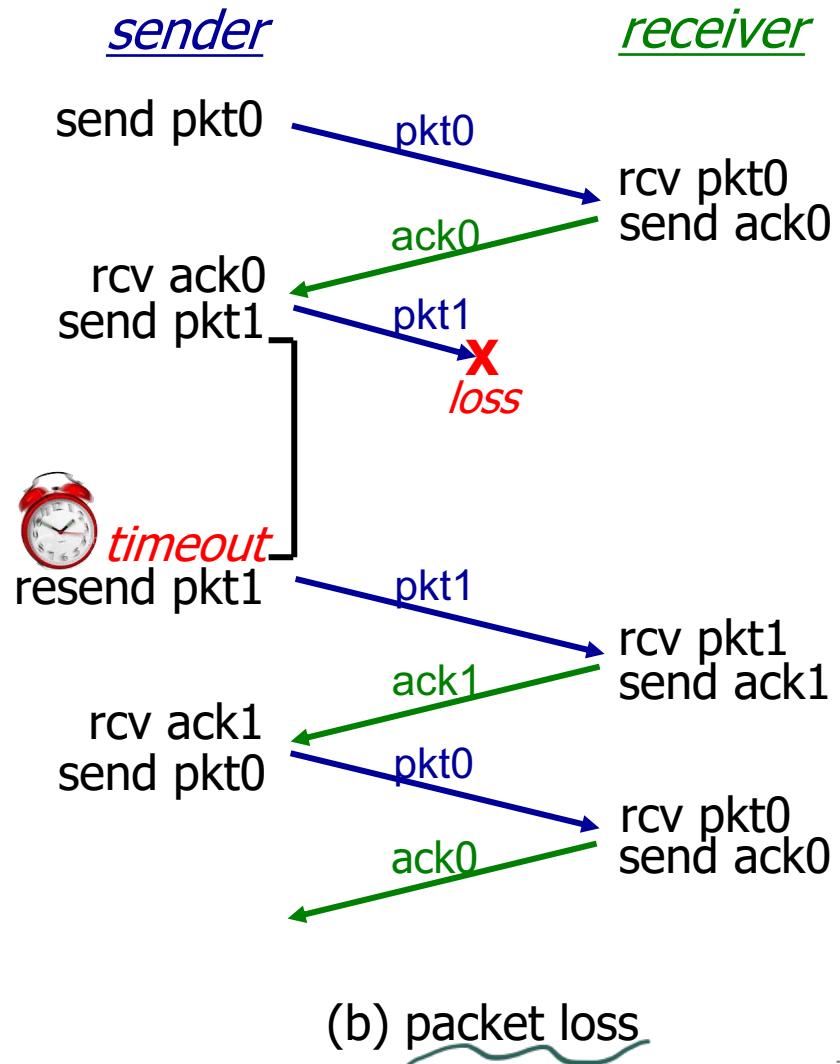
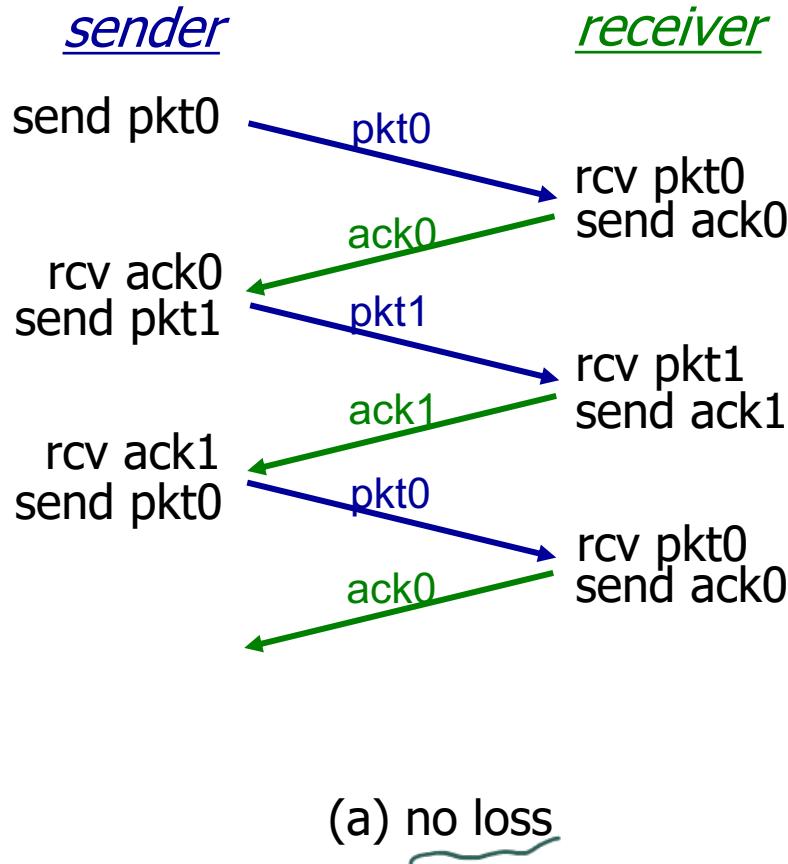
rdt3.0 sender



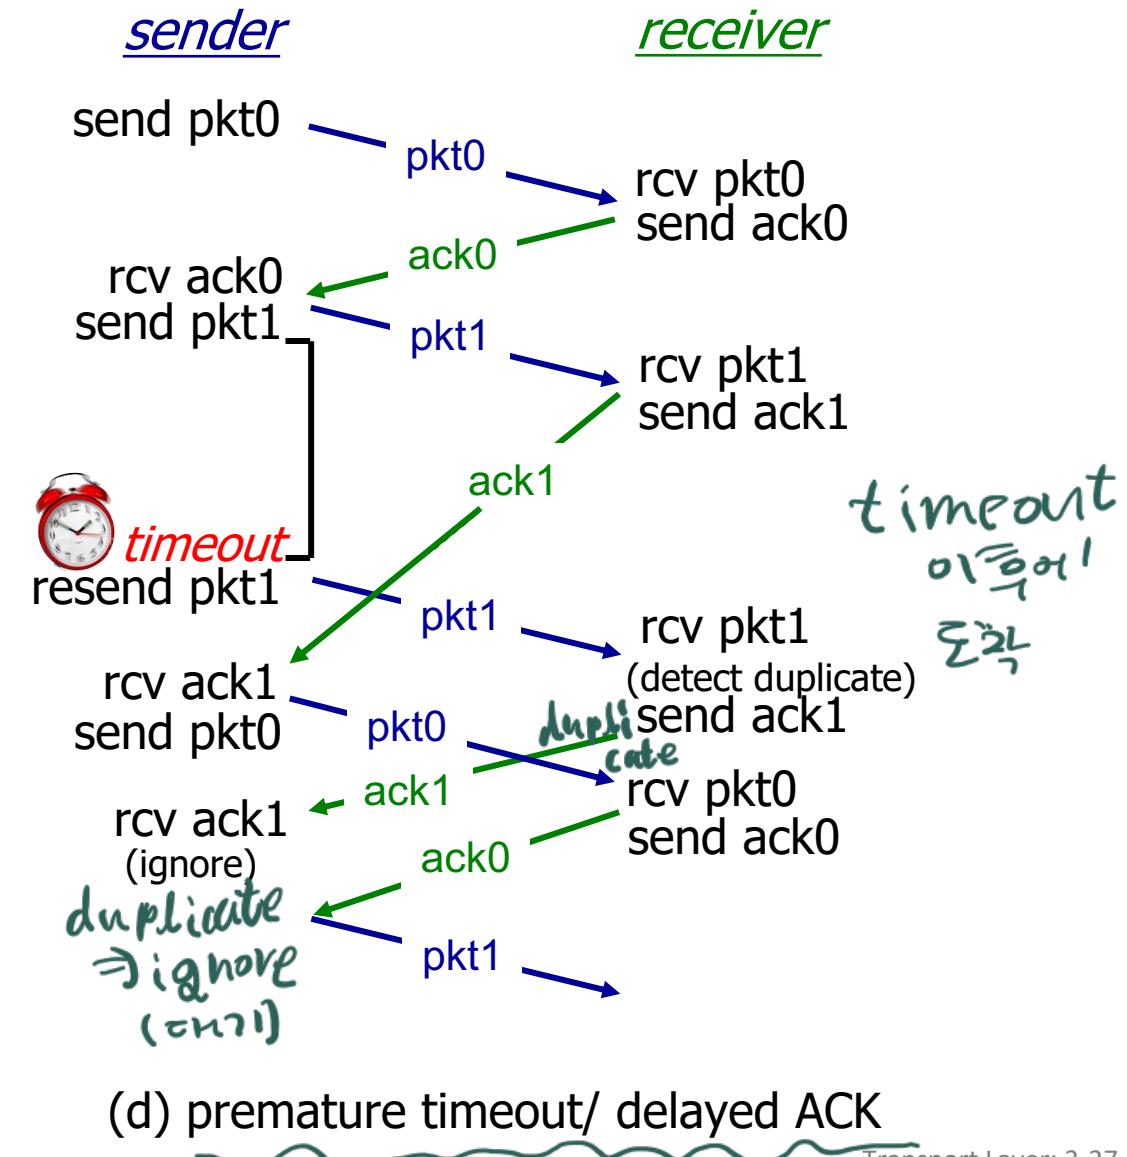
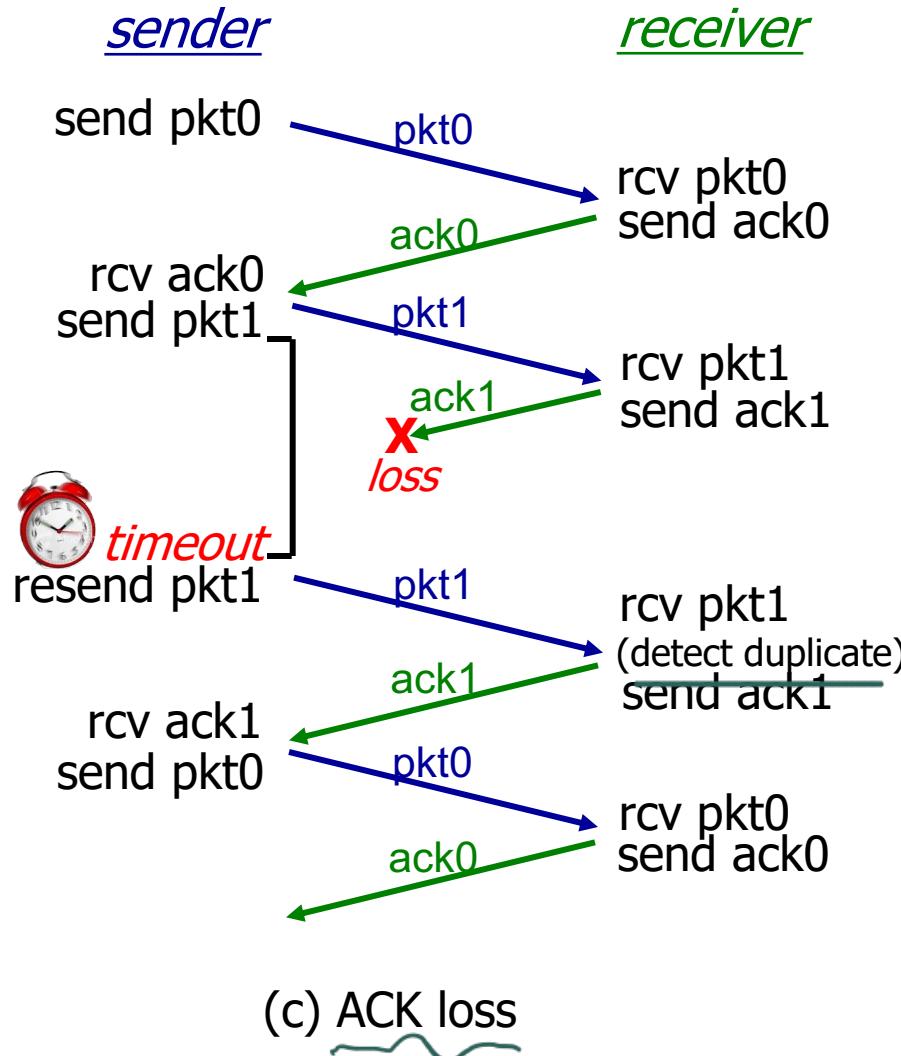
rdt3.0 sender



rdt3.0 in action



rdt3.0 in action



Performance of rdt3.0 (stop-and-wait)

↳ terribly bad (sour 216 307)

- U_{sender} : utilization – fraction of time sender busy sending time
 $\frac{\text{sending time}}{\text{total delay}}$
↳ time loss ↑

- example: 1 Gbps link, 15 ms prop. delay, 8000 bit packet

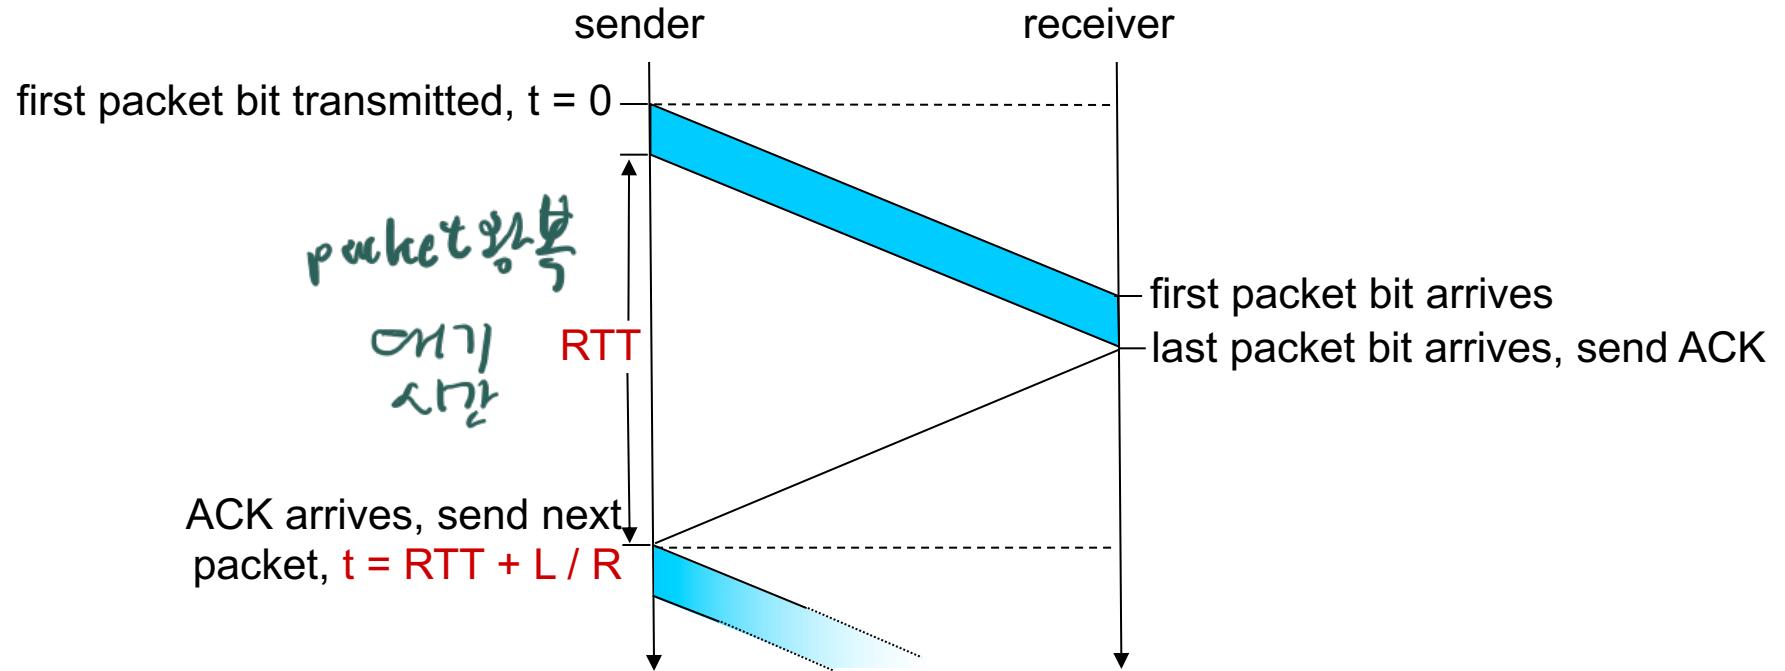
↳ 10^9 bits/sec

- time to transmit packet into channel:

$$D_{\text{trans}} = \frac{L}{R} = \frac{8000 \text{ bits}}{10^9 \text{ bits/sec}} = 8 \text{ microsecs}$$

delay

rdt3.0: stop-and-wait operation



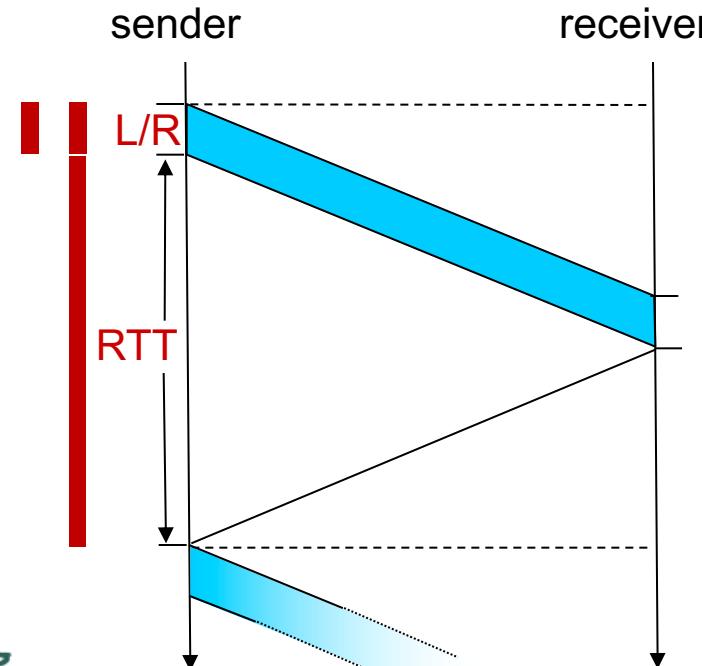
rdt3.0: stop-and-wait operation

$$U_{\text{sender}} = \frac{L/R}{RTT + L/R}$$

$$= \frac{.008}{30.008}$$

$$= 0.00027$$

↳ 這樣慢的
慢到爆

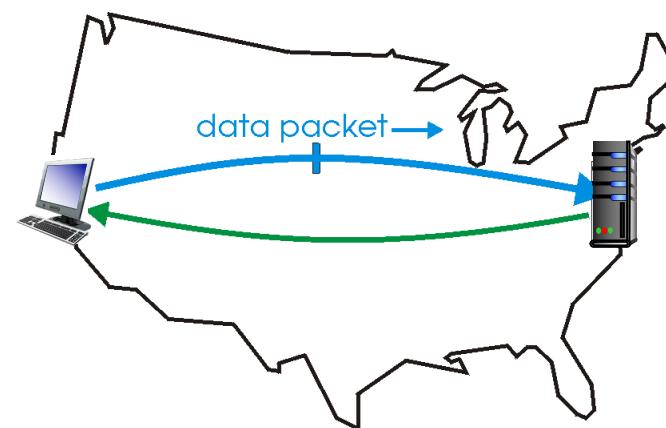


- rdt 3.0 protocol performance stinks!
- Protocol limits performance of underlying infrastructure (channel)

rdt3.0: pipelined protocols operation

pipelining: sender allows multiple, “in-flight”, yet-to-be-acknowledged packets

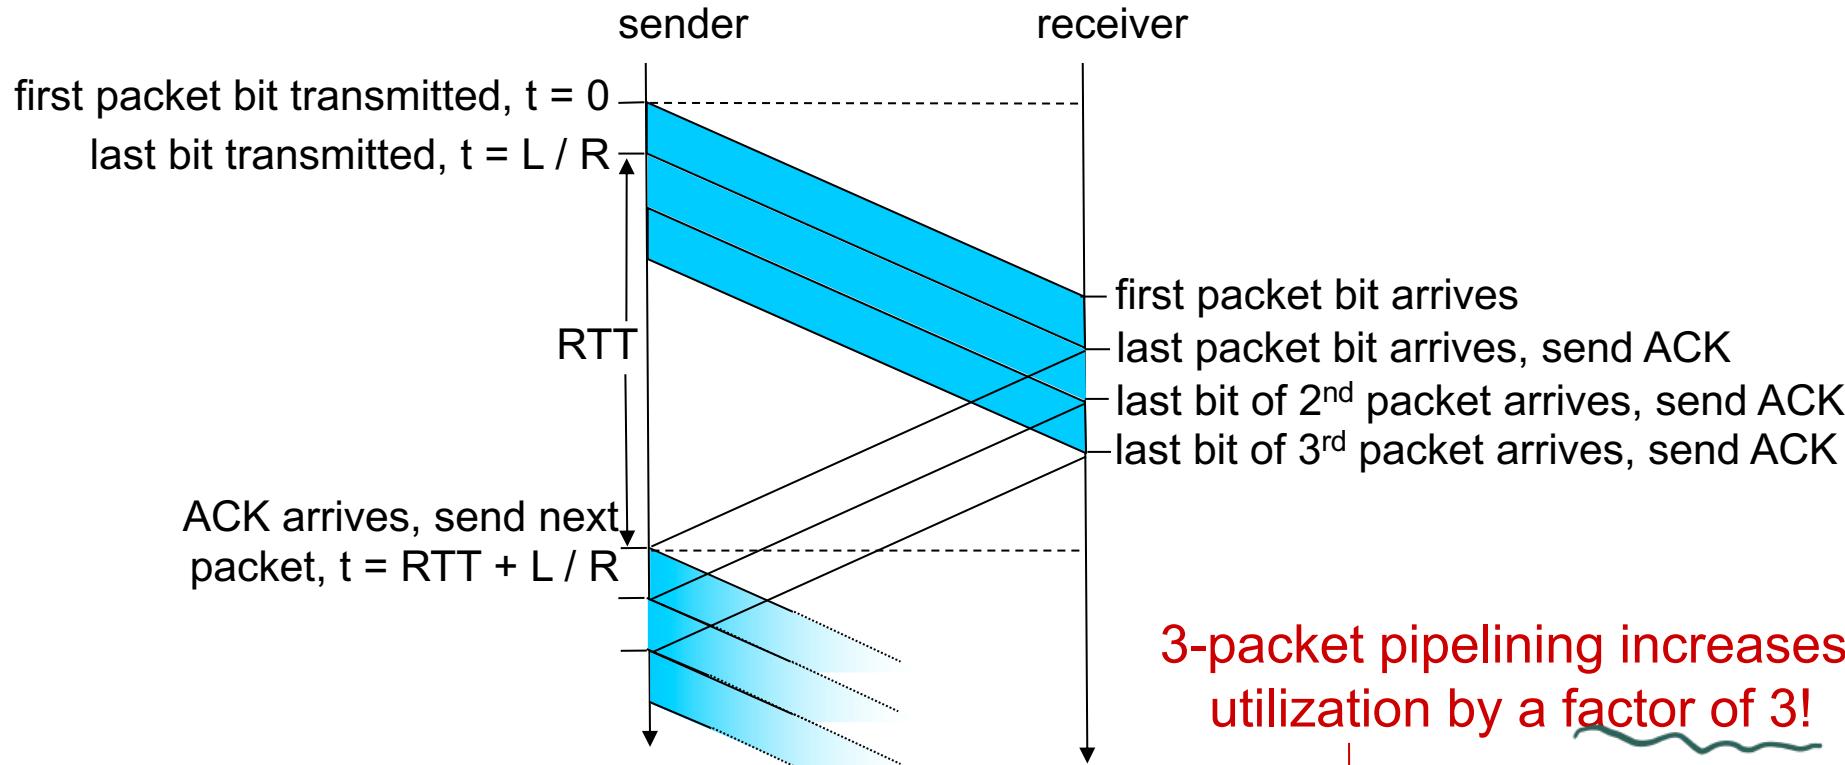
- range of sequence numbers must be increased
- buffering at sender and/or receiver



(a) a stop-and-wait protocol in operation

수신기
수신기
packet 송신

Pipelining: increased utilization



$$U_{\text{sender}} = \frac{3L/R}{RTT + L/R} = \frac{.0024}{30.008} = 0.00081$$

3-packet pipelining increases utilization by a factor of 3!

pipelining
시그널의
utility↑

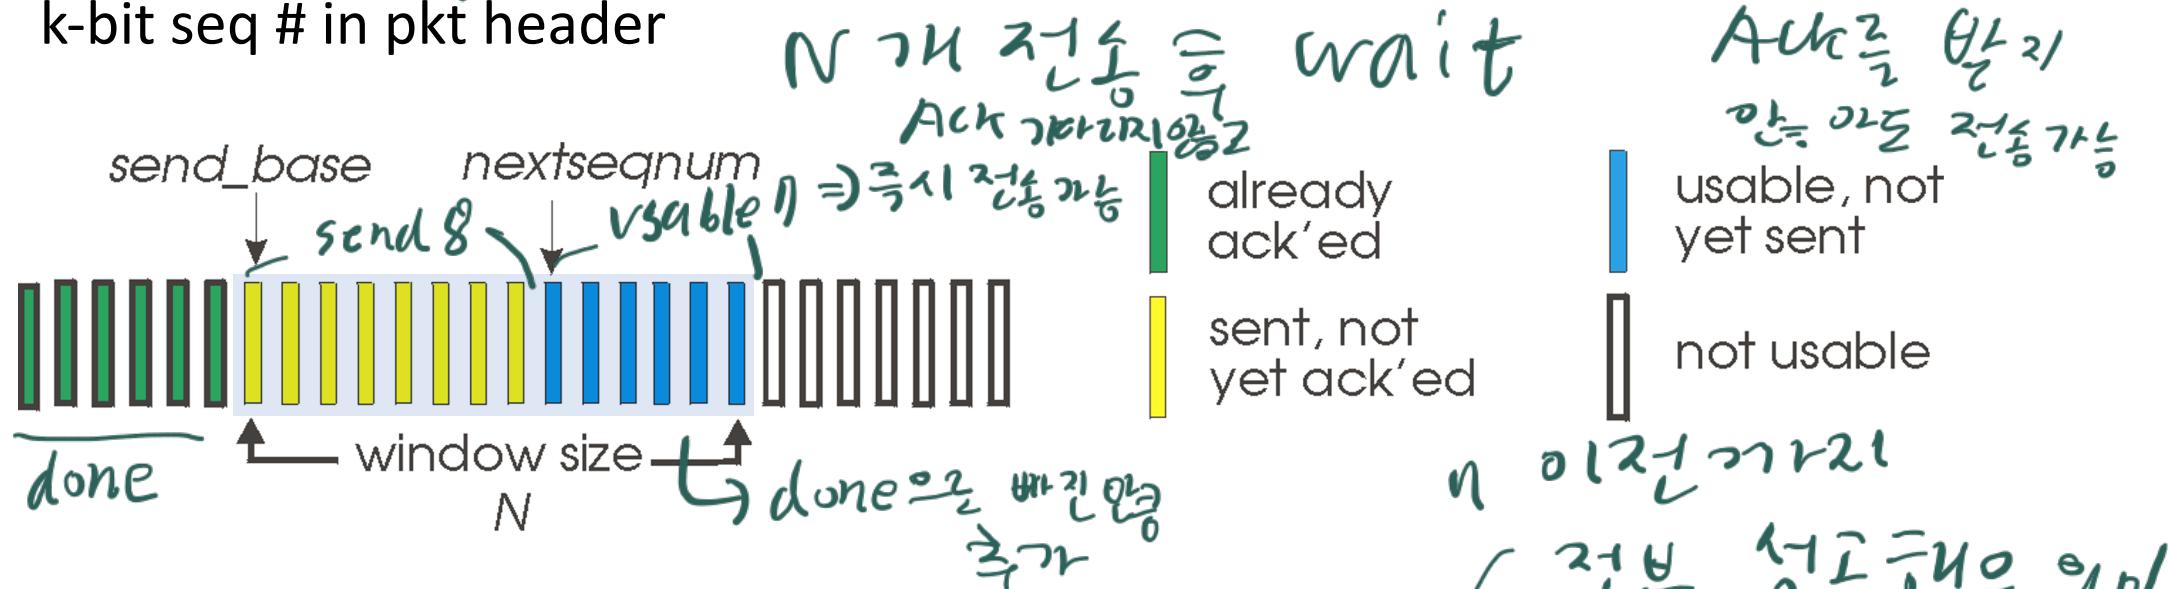
1. sender는 ACK를 받지 않은 상태로 동시에 최대로 보낼 수 있는 패킷의 수인 window size를 설정
2. sender는 send_base, nextseqnum 설정
→ send_base: 송신자가 이미 보냈지만 아직 ACK를 받지 못한 첫번째 패킷의 시퀀스 번호
→ nextseqnum: 송신자가 현재 전송할 수 있는 다음 패킷의 시퀀스 번호
3. 수신자는 패킷들을 연속해서 받고, 각각의 패킷을 받을 때마다 sequence number를 확인하고 ACK를 전송
→ N번까지 잘받았다 → Cumulative Ack
→ 패킷이 유실되어 sequence number에 빈 틈이 생기면, 정상적으로 받은 가장 마지막 패킷 번호를 ACK에 담아 보낸다.
→ 수신자는 정상적으로 받은 마지막 패킷 sequence number를 rcv_base에 기록한다.
4. 송신자는 보낸 패킷의 ACK가 정상적으로 왔다면, 패킷 처리 후 window를 오른쪽으로 한 칸 sliding 한다.

Go-Back-N: sender

한 번의 err이
packet push

- sender: “window” of up to N, consecutive transmitted but unACKed pkts

- k-bit seq # in pkt header



- cumulative ACK**: $\text{ACK}(n)$: ACKs all packets up to, including seq # n
- on receiving $\text{ACK}(n)$: move window forward to begin at $n+1$
- timer for oldest in-flight packet → 가장 오래된 packet ack 알림 → receive 시 window 초기화 알림
- $\text{timeout}(n)$: retransmit packet n and all higher seq # packets in window

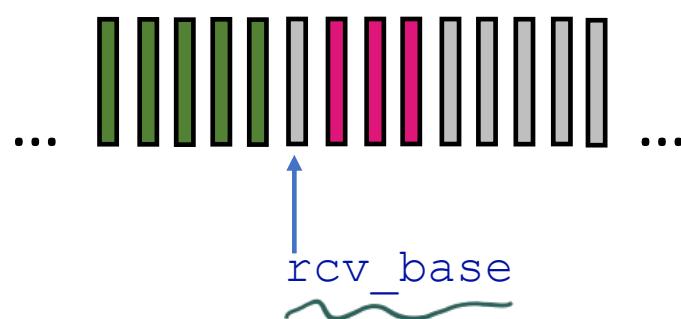
Go-Back-N: receiver

NAK X

- ACK-only: always send ACK for correctly-received packet so far, with highest *in-order* seq #
 - may generate duplicate ACKs
 - need only remember `rcv_base`
 - on receipt of out-of-order packet: → cumulative ACK
(+ BOS) buffer
 - can discard (don't buffer) or buffer: an implementation decision
 - re-ACK pkt with highest *in-order* seq #

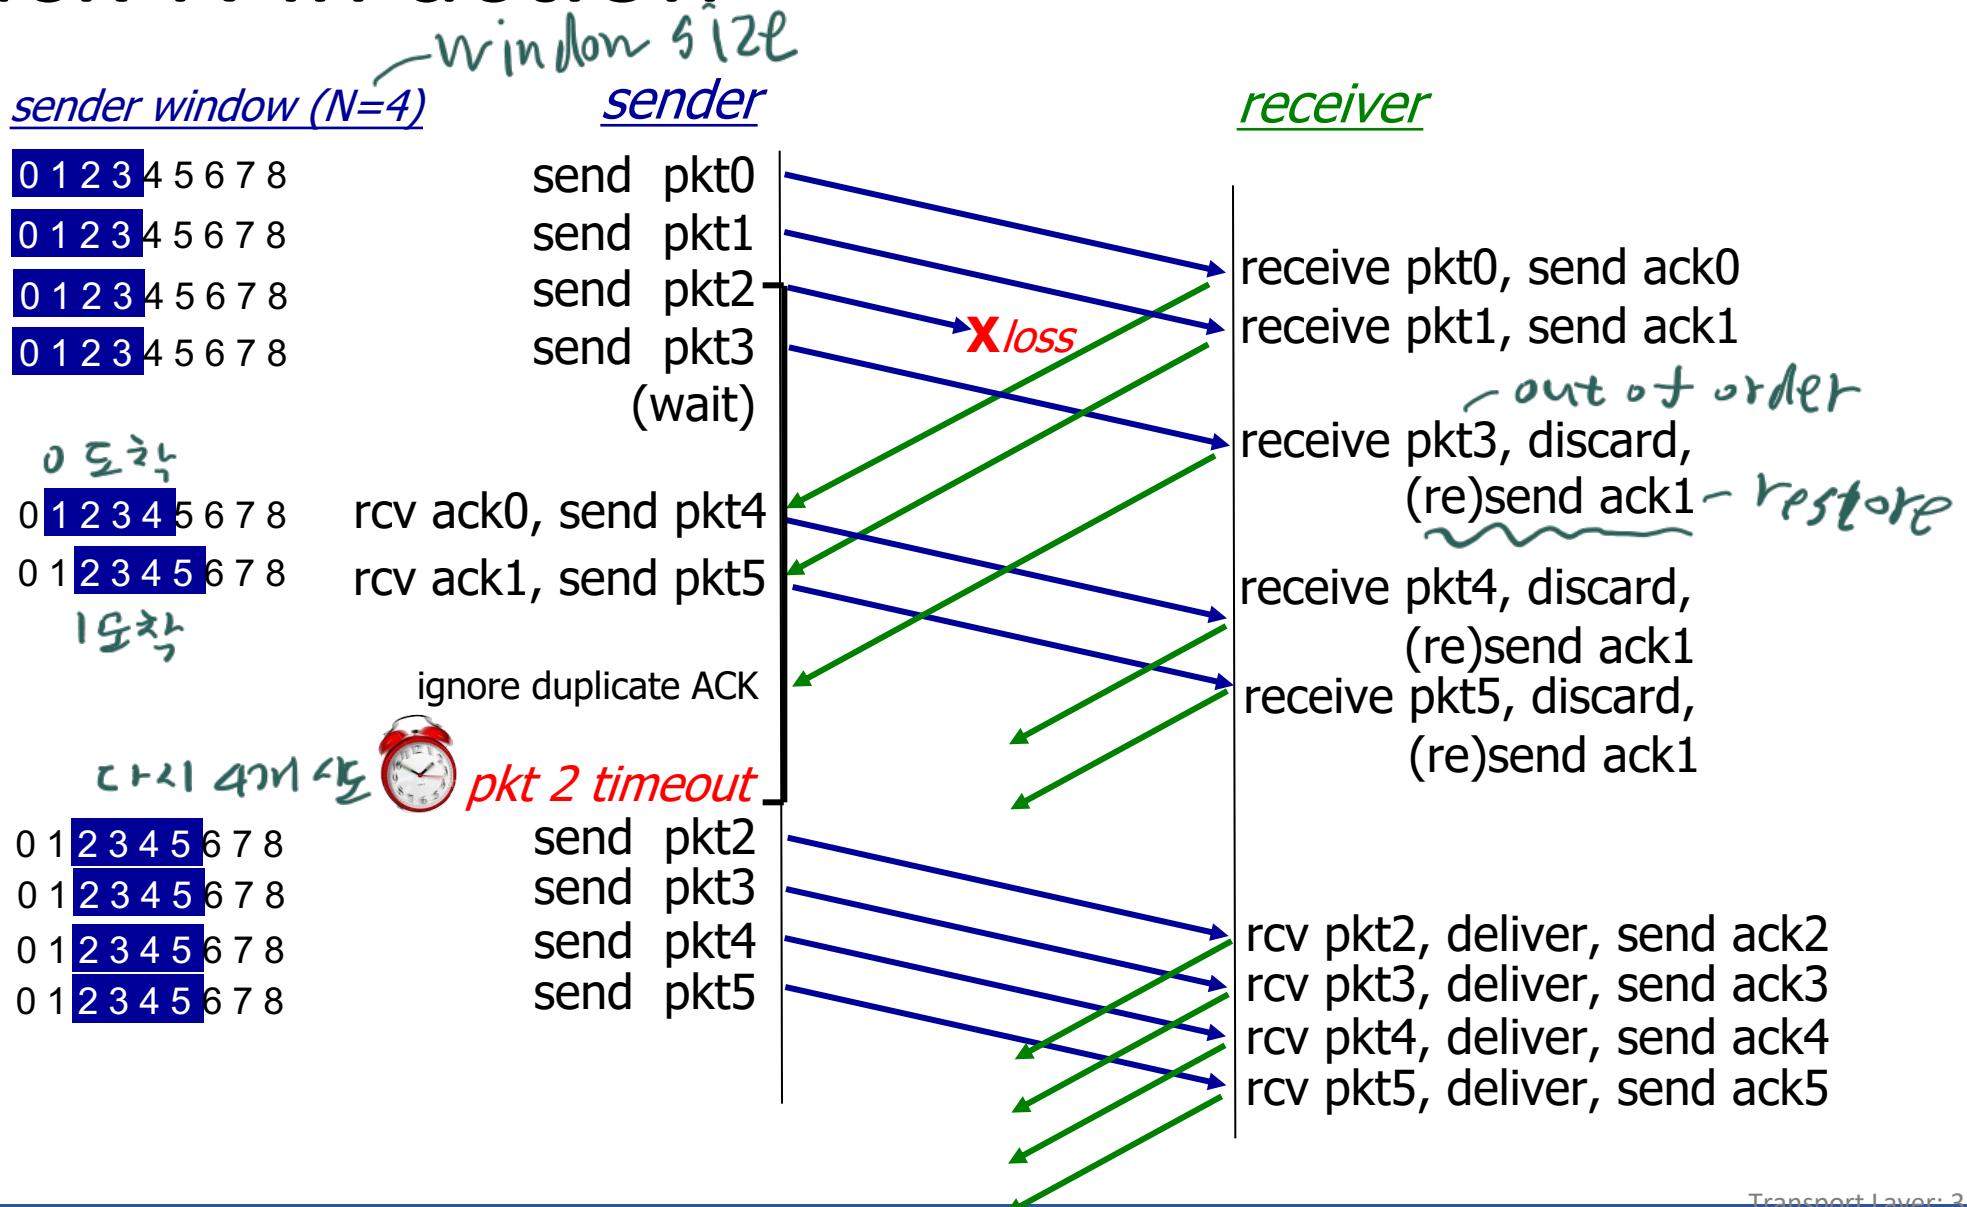
→ cumulative ACK
from buffer to X
implementation decision

Receiver view of sequence number space:



- received and ACKed
 - Out-of-order: received but not ACKed
 - ↳ 앞에서 not received'가
 발생한 경우
 - Not received

Go-Back-N in action



1. sender는 window size 설정
2. window size 내에 들어있는 패킷들을 전부 송신 후 ACK을 기다림
→ 각각의 패킷에 대해 타이머를 가지고 있어 time-out을 활용해 재전송 한다.
3. Selective repeat 방식에서는 수신자가 버퍼를 가지고 있다.
 - 수신 패킷 처리: 수신자는 패킷의 시퀀스 번호를 확인하여 `recv_base`에 저장된 시퀀스 번호와 어긋난다면 패킷들을 임시로 버퍼에 저장해 둔다 (중간에 corrupt 발생).
 - 순서에 맞는 패킷 도착 & ACK 전송: 이후에 순서에 맞는 패킷이 도착하면, 수신자는 버퍼에 저장된 패킷들을 모아 상위 계층으로 전달한다.
 - 윈도우 슬라이딩: ACK 전송 횟수만큼 송신자의 윈도우가 슬라이딩된다.
4. 특정 패킷에 대해 ACK가 들어오면 그 패킷을 완료 기록
→ 송신자가 보낸 패킷들 중 아직 ACK되지 않은 패킷들 중 가장 번호가 작은 패킷이 ACK되면 window를 sliding

Selective repeat — another way

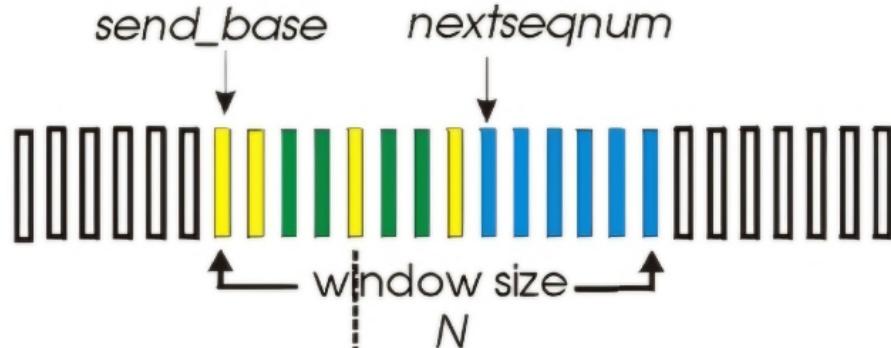
→ 각각의 packet에 대한 정부 저장

- receiver individually acknowledges all correctly received packets
 - buffers packets, as needed, for eventual in-order delivery to upper layer
- sender times-out/retransmits individually for unACKed packets
 - sender maintains timer for each unACKed pkt
- sender window
 - N consecutive seq #s
 - limits seq #s of sent, unACKed packets

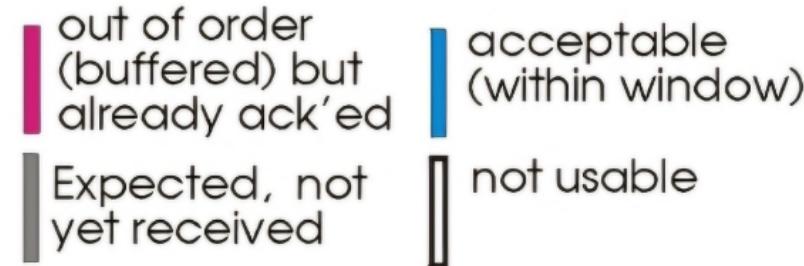
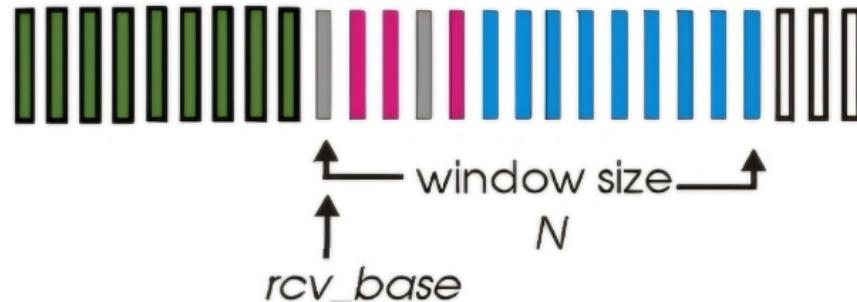
buffer
↑
1 2 3 4 5

선패에 이후를 전부 재전송하는 대신,
ACK 전달 \Rightarrow 선패은 packet만 재전송

Selective repeat: sender, receiver windows



(a) sender view of sequence numbers



(b) receiver view of sequence numbers

Selective repeat: sender and receiver

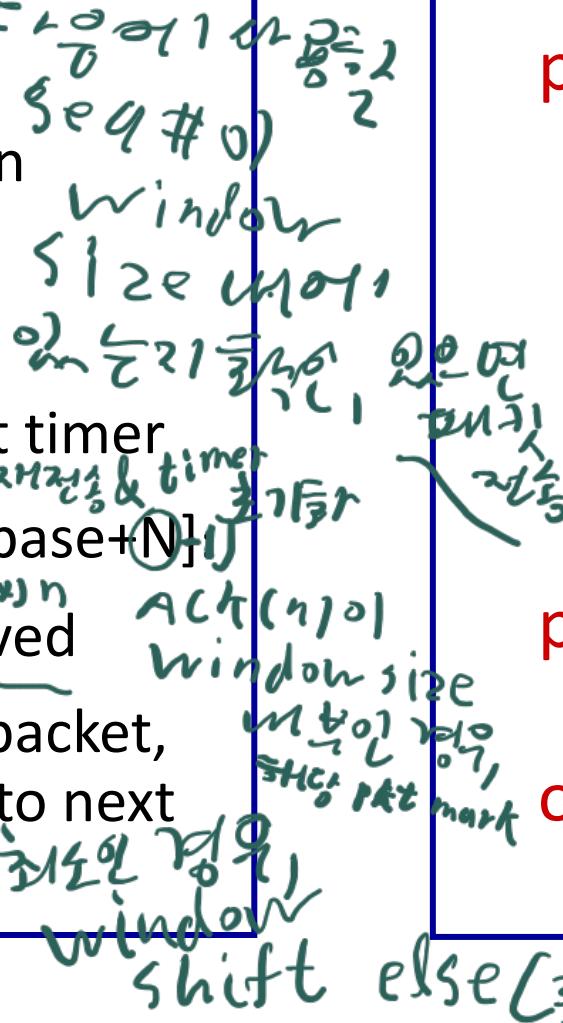
above (application layer)에서
sender

data from above:

- if next available seq # in window, send packet

timeout(n):

- resend packet n , restart timer
 \rightarrow packet timeout & timer
- ACK(n) in $[sendbase, sendbase+N]$
- mark packet n as received
- if n smallest unACKed packet, advance window base to next unACKed seq #



receiver

packet n in $[rcvbase, rcvbase+N-1]$

- send ACK(n)
- out-of-order: buffer = \downarrow \rightarrow buffered, in-order packets, \Rightarrow ACK \rightarrow window
- in-order: deliver (also deliver \Rightarrow buffered, in-order packets), advance window to next not-yet-received packet \Rightarrow window

packet n in $[rcvbase-N, rcvbase-1]$

- ACK(n)

otherwise:

- ignore

수신된 seq # window size

packet n in $[rcvbase, rcvbase+N-1]$

→ 받았던 흔적 확인(Ex) \rightarrow window

수신 알음(중간 missing) \rightarrow window

수신 알음 \rightarrow window

buffered, in-order packets, \Rightarrow window

advance window to next not-yet-received packet \Rightarrow window

packet n in $[rcvbase-N, rcvbase-1]$

↓
수신 알음 (Ex) \rightarrow window

수신 == ACK \rightarrow window

그 다음 window

만나면 \Rightarrow ACK(n)

Selective Repeat in action

sender window ($N=4$)

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8
0	1	2	3	4	5	6	7	8

sender

send pkt0
send pkt1
send pkt2
send pkt3
(wait)

rcv ack0, send pkt4
rcv ack1, send pkt5

record ack3 arrived



pkt 2 timeout

send pkt2
(but not 3,4,5)

Q: what happens when ack2 arrives?

receiver

receive pkt0, send ack0
receive pkt1, send ack1

receive pkt3, buffer,
send ack3

receive pkt4, buffer,
send ack4
receive pkt5, buffer,
send ack5

rcv pkt2; deliver pkt2,
pkt3, pkt4, pkt5; send ack2

Next...

- *Chapter 3.5 Connection-Oriented Transport: TCP*